

# Table of Contents

## Part 1: Electromagnetic Foundations

- Chapter 1: Wave Propagation and Transmission Lines
- Chapter 2: Impedance, Reflections, and Termination
- Chapter 3: Skin Effect and Dielectric Loss
- Chapter 4: Inductance and Magnetic Coupling
- Chapter 5: Return Path Dynamics and Parasitic Effects

## Part 2: Signal Integrity Measurement

- Chapter 6: Time Domain Reflectometry
- Chapter 7: Smith Charts
- Chapter 8: S-Parameters and Vector Network Analysis
- Chapter 9: Fourier Analysis

## Part 3: Jitter Analysis

- Chapter 10: Jitter Decomposition and Measurement
- Chapter 11: Dual-Dirac Model and BER Extrapolation
- Chapter 12: Test Patterns and Jitter Isolation

## Part 4: Equalization and RX Architecture


- Chapter 13: Transmitter Feed-Forward Equalization
- Chapter 14: Continuous Time Linear Equalization and Decision Feedback Equalization
- Chapter 15: PAM4 Signaling and Gray Coding

## Part 5: SerDes Architecture and Clocking

- Chapter 16: SerDes Architecture and Eye Diagrams
- Chapter 17: CDR and PLL Loop Dynamics
- Chapter 18: Line Coding, FEC, and Protocol Framing

## Part 6: Coherent Optics

- Chapter 19: The Physics of the Electromagnetic Wave
- Chapter 20: Electro-Optic Modulation and the Refractive Index
- Chapter 21: The I/Q Modulator Architecture
- Chapter 22: Wideband Signal Analysis and EVM



# Chapter 1: Wave Propagation and Transmission Lines

---

Every high-speed digital system operates on a single physical principle: information travels as electromagnetic waves, not as electrons drifting through copper. The moment a transistor connects a trace to a voltage rail, a localized electric field boundary forms and propagates through the dielectric material at a significant fraction of the speed of light. The electrons themselves barely move. Understanding this distinction is the foundation on which all signal integrity analysis rests, from characteristic impedance and TDR measurements to jitter budgeting and equalization design.

This guide traces the complete physical mechanism of wave propagation on a PCB transmission line: the distributed nature of trace geometry, the CMOS switching event that launches a wavefront, the segment-by-segment chain reaction that constitutes propagation, the relationship between signal voltage and current, and the Telegrapher's Equations that formalize the entire system mathematically.

## Core Concepts

---

### Distributed vs. Lumped Elements

A discrete resistor, capacitor, or inductor is a *lumped* element: a component localized at a single physical coordinate, designed to hold a specific value and small enough that the signal traverses it in effectively zero time. Ohm's law, Kirchhoff's voltage law, and the standard circuit equations all assume lumped behavior.

A transmission line is a *distributed* structure. Its physical dimensions are significant relative to the wavelength of the signal it carries. The capacitance and inductance are not localized hardware elements; they exist continuously along every microscopic fraction of the trace. A bare microstrip trace on a PCB is itself a capacitor: the trace is the top metal plate, the ground plane is the bottom metal plate, and the FR4 fiberglass laminate is the insulating dielectric between them. No soldered ceramic component is required. The capacitance is intrinsic to the geometry and distributed along the entire length.

The inductance is similarly geometric. Any current flowing through a conductor generates a surrounding magnetic field that resists changes in current flow. The signal current travels forward on the trace while the return current travels backward on the ground plane directly below it, forming a current loop. The area enclosed by this loop stores energy in the magnetic field, producing a distributed series inductance per unit length.

These two parameters are geometrically coupled. Widening a trace increases the surface area facing the ground plane, which increases capacitance. That same widening simultaneously provides a broader path for current, which decreases the loop area and lowers inductance. Moving the trace closer to the ground plane increases capacitance and decreases inductance. The ratio of distributed inductance to distributed capacitance per unit length determines the characteristic impedance of the line:

$$Z_0 = \sqrt{\frac{L}{C}}$$

## Why Ohm's Law Fails at High Speed

Ohm's law (  $V = IR$  ) assumes an immediate steady-state condition where electrons drift uniformly across the entire conductor and the voltage is identical at every point. This assumption requires the signal to have already finished propagating and settled into equilibrium.

High-speed serial links operate in a regime where this equilibrium never exists during a signal transition. A CMOS driver transitions from low to high in picoseconds. The voltage at the near end of the trace changes before the wave has reached the far end. The wavefront sees the bare trace as an infinite chain of empty geometric capacitors it must fill with charge. The energy is consumed charging the local geometry, not overcoming the DC resistance of the copper.

The ratio of voltage to current within the traveling wave is determined by the trace geometry alone, expressed as the characteristic impedance  $Z_0 = \sqrt{L/C}$  , not the DC resistance of the conductor. Treating a high-speed transmission line as a lumped resistor produces physically incorrect results.

## Transverse Electromagnetic (TEM) Propagation

A static voltage applied by a disconnected battery creates a pure electric field with no associated magnetic field. A steady DC current creates a constant magnetic field with no changing electric field. Neither of these scenarios produces a propagating wave.

Propagation requires both fields to exist simultaneously and to vary in time. Maxwell's equations mandate that a changing electric field induces a changing magnetic field, and a changing magnetic field induces a changing electric field. These two fields couple into a self-sustaining transverse electromagnetic (TEM) wave that travels through the dielectric material between the trace and the ground plane.

The electromagnetic wave propagates exclusively through the dielectric material, not through the copper. The copper trace and the ground plane serve only as physical guide rails that confine and direct the wave. This is the fundamental reason that the properties of the PCB laminate material (dielectric constant, loss tangent) completely dominate high-speed signal integrity engineering.

## Propagation Velocity and the Dielectric Constant

The speed of the electromagnetic wave in a transmission line is governed by the dielectric constant (  $\epsilon_r$  , also called relative permittivity or  $D_k$  ) of the insulating material surrounding the conductors:

$$v_p = \frac{c}{\sqrt{\epsilon_r}}$$

where  $c$  is the speed of light in a vacuum (  $\approx 3 \times 10^8$  m/s).

Standard FR4 has a dielectric constant of approximately 4.0 to 4.5, which reduces the propagation velocity to roughly half the speed of light. Lower-loss laminates such as Megtron 6 have lower dielectric constants (around 3.6), which allow slightly faster propagation.

The amplitude of the signal has no effect on propagation velocity. A higher voltage produces a stronger electric field (higher "electrical pressure"), but the velocity of an electromagnetic wave in a transmission line is determined entirely by the distributed inductance and capacitance of the physical geometry:

$$v_p = \frac{1}{\sqrt{LC}}$$

Neither  $L$  nor  $C$  changes when the voltage increases from 1V to 5V. The speed remains identical regardless of the signal amplitude.

## Architecture

---

### The CMOS Driver as a High-Speed Valve

The voltage rail (such as +1.2V) and the ground plane (0V) are two reservoirs of constant electrical potential. The power distribution network (PDN) works to keep these reservoirs stable and fully charged, regardless of what the surrounding logic circuits are doing.

At the beginning of a trace sits an output driver (transmitter). At the most basic level, this driver consists of two complementary transistors acting as switches. The pull-up transistor connects the trace to the voltage rail. The pull-down transistor connects the trace to ground. Information is not sent by asking the power supply to generate different voltage levels. Instead, information is encoded in *time*. To send a logic "1," the driver closes the pull-up switch and opens the pull-down switch, connecting the trace to the +1.2V reservoir. To send a logic "0," the driver does the reverse.

The power supply provides the constant potential energy. The transistor provides the timing. The laws of physics dictate that abrupt transitions in the time domain inherently contain broad frequency content in the frequency domain.

### The Wavefront Launch Event

Consider a trace currently sitting at 0V. At the exact picosecond the pull-up transistor switches on, the physical copper at the beginning of the trace is abruptly exposed to the +1.2V potential from the rail.

The entire length of the trace cannot instantly jump to 1.2V. The distributed parasitic capacitance requires time to charge, and the distributed series inductance resists sudden current flow. A sharp boundary forms right at the transmitter between the newly applied 1.2V region and the existing 0V region.

The massive difference in potential across this microscopic physical boundary creates a very intense, localized electric field. This field, accompanied by a magnetic field from the inrush current, is what launches down the trace as an electromagnetic wave.

### Frequency Content of the Transition

The high-frequency content carried by the signal does not come from different voltage levels output by the rail. It comes entirely from the speed of the transition between the two DC states.

Toggling the rail connection at a constant rate creates a square wave in the time domain. Fourier analysis shows that a perfect square wave is constructed by summing a fundamental sine wave with an infinite series of odd harmonics. A fast rise time (for example, 20 picoseconds) requires frequency harmonics extending well past 15 GHz to mathematically construct that steep vertical edge, regardless of how infrequently the transitions occur.

Two distinct frequency scales coexist on the same transmission line:

- **The data rate frequency** describes how often the driver chooses to change states. An alternating 10101010 pattern at 10 Gbps has a fundamental frequency of 5 GHz. A long run of identical bits (like 11111) drops to 0 Hz (DC) for that duration.
- **The edge bandwidth** describes the high-frequency harmonics required to construct the shape of each rising or falling transition. This bandwidth is dictated entirely by the transistor's switching speed, not by the bit pattern.

Both frequency scales propagate simultaneously on the same physical trace. Any mechanism that treats high frequencies differently from low frequencies (such as the skin effect, dielectric absorption, or dispersion) will cause the composite square wave to distort and smear as it propagates.

## Segment-by-Segment Propagation

The physical mechanism of propagation can be understood by mentally dividing the entire trace into millions of microscopic, contiguous segments. Each segment is a tiny piece of copper separated from the ground plane by an insulator, making each segment a tiny capacitor with a tiny series inductance.

**Time  $t = 0$  :** The transistor switches on, abruptly exposing the very first microscopic segment of the trace (at  $d = 0$ ) to the 1.2V rail. To bring this first segment up to 1.2V, charge must be deposited onto it. A localized burst of electrons flows from the power supply into this segment. This physical movement of charge is the conduction current. As charge accumulates, an electric field rapidly builds across the dielectric, pointing from the trace down to the ground plane. Maxwell proved that a rapidly changing electric field behaves exactly like a current. This is the displacement current, and it bridges the physical gap of the insulator, completing the local circuit loop back through the ground plane to the source. At this moment, only the first microscopic segment is affected. The rest of the trace is entirely unaware that anything has happened.

**Time  $t > 0$  :** Segment 1 is now charged to 1.2V, but segment 2 remains at 0V. The voltage difference causes segment 1 to act as the local power source for segment 2. Electrons shift from segment 1 into segment 2. The local series inductance briefly resists this sudden movement, but charge transfers. Segment 2 charges, and an electric field forms beneath it, creating a local displacement current through the dielectric and a local magnetic field around the copper. Segment 2 then acts as the source for segment 3. Segment 3 charges segment 4.

This continuous, sequential charging of microscopic segments is the physical mechanism of electromagnetic wave propagation on a transmission line.

## The Three Zones of a Propagating Wavefront

Freezing time while the wavefront is halfway down the board reveals three distinct zones:

1. **Behind the wavefront:** The trace is fully charged to 1.2V. A static electric field holds steady between the trace and the ground plane, but the voltage is not changing. The conduction current in the trace is zero. The displacement current through the dielectric is zero. The transmission line acts like a charged capacitor holding a DC state.
2. **Ahead of the wavefront:** The trace is still at 0V. The voltage is not changing. The conduction current is zero. The displacement current is zero.
3. **Exactly at the wavefront:** The voltage is actively transitioning from 0V to 1.2V. Electrons are shifting to charge the local segment (conduction current). The electric field is growing, passing energy through the dielectric (displacement current). The current flows in a highly localized loop: forward on the trace to the wavefront, vertically through the dielectric as displacement current, and backward through the ground plane to the source. This entire active loop physically slides down the transmission line at the propagation velocity, carrying the signal with it.

## The Role of the Power Supply During Propagation

A critical physical detail: segment 1 does not simply dump its stored charge into segment 2 and go dead. Segment 1 remains continuously connected to the 1.2V power supply through the closed transistor switch. The power supply acts as a continuous, high-pressure source that forces charge into the trace at  $d = 0$  to feed the propagating wavefront. Every segment behind the wavefront stays firmly pinned at 1.2V.

The forward direction of energy flow is maintained by the permanent voltage gradient. The driver at  $d = 0$  is forcefully held at 1.2V, representing high potential energy. The uncharged trace ahead of the wavefront sits at 0V. Charge always flows from higher potential to lower potential. The electrical pressure behind the wavefront is 1.2V, and the pressure in front is 0V. A forward segment cannot transfer charge backward to a previous segment because the previous segment is held at a higher electrical potential by the continuous supply.

When the transmitter decides to send a logic "0," it disconnects the 1.2V rail and connects the trace to 0V ground. The high-pressure source is removed, a low-resistance path to ground is provided, and the built-up charge rushes out of the trace, discharging it segment by segment until the local electric field collapses.

## Conventional Current vs. Electron Motion

Operating a standard positive voltage rail (+1.2V) means the power supply is performing the exact opposite of "pushing electrons into the trace." Electrons carry negative charge. The copper trace naturally contains an equal number of protons and electrons, giving it a neutral 0V baseline. Creating a +1.2V potential on the trace requires the power supply to *strip electrons away* from the copper atoms and deposit them into the ground plane.

This removal of electrons creates a localized electron deficiency (a net positive charge) at  $d = 0$ . The propagating wavefront requires the power supply to continuously pull more electrons out of the trace at  $d = 0$  so that each newly encountered segment reaches that same +1.2V state of electron deficiency.

Operating a negative voltage rail (-1.2V) reverses this: the power supply literally forces surplus electrons into the trace, creating a state of electron overdensity.

Electrical engineers typically use "conventional current" and state that "the power supply pushes charge into the trace" for a positive voltage. This is a historical mathematical convention established by Benjamin Franklin before the electron was discovered. The physical reality of the subatomic motion is always reversed for positive voltages. The core mechanical principle remains identical regardless of polarity: the power supply actively does physical work at  $d = 0$  to alter the electron density of the copper, and it must continuously perform that work to keep the entire length of the transmission line fully charged behind the wavefront.

## Worked Examples

---

### Electrons, Drift Velocity, and Signal Speed

Electrons move through copper at a drift velocity measured in millimeters per second. The signal arrives at the far end of a trace in under a nanosecond. The resolution to this apparent paradox lies in the distinction between particle transport and wave transport.

The mathematical definition of current is:

$$I = n \cdot A \cdot q \cdot v_d$$

where  $n$  is the electron density (number of free electrons per unit volume),  $A$  is the cross-sectional area of the conductor,  $q$  is the charge of one electron, and  $v_d$  is the drift velocity. The electron density  $n$  is a fixed physical property of the copper lattice (approximately  $8.5 \times 10^{28}$  electrons per cubic meter). To increase current, the drift velocity  $v_d$  increases. While this velocity remains extremely slow, copper contains an astronomical number of free electrons. A microscopic increase in collective drift velocity results in a massive increase in total coulombs passing a physical cross-section per second.

The copper trace is already packed completely full of free electrons. Applying the 1.2V potential at the driver exerts an electromotive force on the electrons at  $d = 0$ . These electrons physically drift into the neighboring segment, creating a cascading compression wave of charge density. An electron at  $d = 0$  only needs to move a fraction of a micrometer to displace its neighbor, which displaces the next neighbor in line.

The signal is the electromagnetic wave traveling through the dielectric. The current is the localized shuffling of electrons facilitating that wave. The electrons leaving the transmitter are not the same electrons arriving at the receiver when the bit flips. The wave propagates through the dielectric at a fraction of the speed of light, locally perturbing the electron sea as it passes.

### Signal, Voltage, and Characteristic Impedance

The "signal" is the electromagnetic wave, and its driving force is the voltage (potential difference). The current is the resulting flow of charge. The relationship linking voltage to current in a traveling wave is the characteristic impedance:

$$Z_0 = \frac{V}{I}$$

A higher voltage pushes more current through a given transmission line geometry. This relationship holds strictly for the traveling wave; it does not describe the DC resistance of the copper.

## How Information Reaches the Receiver

Information is carried by the arrival time of wavefronts, not by the receiver interpreting the frequency content of each edge.

The receiver at the far end of the trace is essentially a voltmeter connected to a very fast, highly synchronized clock recovery circuit. A rising edge wavefront arrives, and the receiver's local voltage jumps from 0V to 1.2V. The clock ticks. The voltage is 1.2V. The receiver records a "1." The clock ticks again. The trace is still fully charged (the transmitter kept the pull-up transistor closed, so no falling edge has arrived). The voltage is still 1.2V. The receiver records a second "1." A falling edge wavefront arrives, discharging the trace back to 0V. The clock ticks. The voltage is 0V. The receiver records a "0."

The data is encoded purely in the time duration between wavefronts. The extreme high-frequency bandwidth discussed earlier is the physical mechanism required to make those wavefront edges steep and sharp. If the edges lacked that high-frequency content, they would be slow and sloped, and the receiver would struggle to identify the exact moment the transition crossed the threshold, leading to jitter and bit errors.

## Edge Cases

---

### The Telegrapher's Equations and the Full Impedance Model

The simplified characteristic impedance formula  $Z_0 = \sqrt{L/C}$  assumes a lossless transmission line. The complete model, derived by Oliver Heaviside from the Telegrapher's Equations, accounts for losses in both the conductor and the dielectric:

$$Z_0 = \sqrt{\frac{R + j\omega L}{G + j\omega C}}$$

In this formulation, each infinitesimal slice of the transmission line contains four distributed parameters:

- **$R$**  (series resistance): the DC and AC resistance of the conductor, including skin effect losses at high frequencies.
- **$L$**  (series inductance): the self-inductance of the current loop formed by the trace and its return path on the ground plane.
- **$G$**  (shunt conductance): the microscopic leakage current that flows through the dielectric insulator from the signal trace to the ground plane.  **$G$**  is the exact mathematical inverse of resistance and accounts for dielectric loss.
- **$C$**  (shunt capacitance): the parallel capacitance between the trace and the ground plane.

The series elements (  $R$  and  $L$  ) are in-line with the current path. They dictate the impedance that the physical current encounters while traveling along the conductor. The shunt elements (  $G$  and  $C$  ) are in parallel between the trace and the ground plane. They determine how much current leaks across the dielectric at each point along the line.

The formula  $Z = R + j\omega L$  describes only the series impedance of a single conductor. It answers the question of which path the physical current will choose when multiple return paths are available (for example, when a signal crosses between reference planes). The ratio of all four distributed parameters (  $R$  ,  $L$  ,  $G$  ,  $C$  ) determines the characteristic impedance the propagating wave experiences as it travels down the line.

## Static Fields Between Transitions

Between edge transitions, no new frequency content propagates along the trace. The section of the trace that the wavefront has already passed acts like a charged capacitor holding a DC state. A static electric field exists between the trace and the ground plane (maintaining the 1.2V potential), but the voltage is not changing, so no displacement current flows, the magnetic field has collapsed, and no propagating electromagnetic wave exists at that location.

The trace remains in this static, charged state until the next wavefront arrives to discharge it. The high-frequency energy content of the signal is concentrated entirely within the rising and falling edges. The flat portions between edges carry only DC energy.

## Chapter 2: Impedance, Reflections, and Termination

---

A transmission line carries electromagnetic energy forward as a propagating wave. The characteristic impedance  $Z_0$  governs the ratio of voltage to current within that wave as long as the geometry remains uniform. The moment the wave encounters a boundary where the impedance changes, physics demands that a portion of the energy reflect backward.

The reflection coefficient  $\rho$  quantifies exactly how much energy reflects and how much transmits through the boundary. It emerges directly from Kirchhoff's voltage and current laws applied at the impedance discontinuity, and it is the single most important quantity in signal integrity engineering. Every TDR measurement, every impedance-controlled PCB stackup, every termination resistor, and every via optimization strategy traces back to controlling  $\rho$ .

This guide derives the reflection coefficient from first principles, traces the physical mechanism of reflections propagating between source and load, explains how termination absorbs reflected energy, and describes the transient process by which a transmission line settles from its initial wave-propagation regime into a DC steady state.

### Core Concepts

---

#### The Reflection Coefficient: Derivation from Kirchhoff's Laws

An incident wave travels down a transmission line with characteristic impedance  $Z_0$  and strikes a boundary where the impedance abruptly changes to  $Z_L$ . At this boundary, the incident wave splits into two components: a reflected wave that bounces backward and a transmitted wave that continues forward into  $Z_L$ .

Three waves exist at the boundary:

1. **Incident wave:** voltage  $V_{inc}$ , current  $I_{inc}$
2. **Reflected wave:** voltage  $V_{ref}$ , current  $I_{ref}$
3. **Transmitted wave:** voltage  $V_{trans}$ , current  $I_{trans}$

Physics imposes two constraints at the exact plane of the boundary.

**Kirchhoff's Voltage Law** requires voltage continuity. The total voltage on the left side of the boundary must equal the total voltage on the right:

$$V_{inc} + V_{ref} = V_{trans}$$

**Kirchhoff's Current Law** requires current continuity. Charge cannot accumulate or vanish at the boundary, so the total current flowing into the junction from the left must equal the current flowing out to the right:

$$I_{inc} + I_{ref} = I_{trans}$$

The relationship between voltage, current, and impedance (  $I = V/Z$  ) links these three waves to the impedances on each side:

- Incident current:  $I_{inc} = V_{inc}/Z_0$
- Transmitted current:  $I_{trans} = V_{trans}/Z_L$
- Reflected current:  $I_{ref} = -V_{ref}/Z_0$

The reflected current carries a negative sign. Voltage is a scalar quantity, but current is a vector with a physical direction. The reflected wave travels backward, so its current vector is opposite to the incident wave's current vector.

Substituting these current expressions into Kirchhoff's Current Law:

$$\frac{V_{inc}}{Z_0} - \frac{V_{ref}}{Z_0} = \frac{V_{trans}}{Z_L}$$

Replacing  $V_{trans}$  with the voltage continuity equation (  $V_{inc} + V_{ref}$  ):

$$\frac{V_{inc}}{Z_0} - \frac{V_{ref}}{Z_0} = \frac{V_{inc} + V_{ref}}{Z_L}$$

Cross-multiplying to clear the denominators:

$$Z_L(V_{inc} - V_{ref}) = Z_0(V_{inc} + V_{ref})$$

Distributing the impedances:

$$Z_L V_{inc} - Z_L V_{ref} = Z_0 V_{inc} + Z_0 V_{ref}$$

Collecting  $V_{inc}$  terms on the left and  $V_{ref}$  terms on the right:

$$V_{inc}(Z_L - Z_0) = V_{ref}(Z_L + Z_0)$$

Isolating the ratio of reflected voltage to incident voltage:

$$\frac{V_{ref}}{V_{inc}} = \frac{Z_L - Z_0}{Z_L + Z_0}$$

Defining this ratio as the reflection coefficient  $\rho$  :

$$\rho = \frac{Z_L - Z_0}{Z_L + Z_0}$$

This equation is the absolute foundation of signal integrity measurement. Every impedance profile on a TDR screen, every S-parameter return loss plot, and every termination design decision derives from it.

## Three Canonical Terminations

The reflection coefficient formula produces three boundary cases that define the extremes of termination behavior.

**Matched termination (  $Z_L = Z_0$  ):** The numerator becomes zero.  $\rho = 0$ . No energy reflects. The entire wave transmits into the load and is absorbed. This is the ideal condition for signal integrity. A 50-ohm trace terminated by a 50-ohm resistor produces zero reflections.

**Open circuit (  $Z_L = \infty$  ):** The reflection coefficient approaches  $\rho = +1$ . The entire wave reflects with the same polarity as the incident wave. The reflected voltage adds constructively to the incident voltage, doubling it at the boundary. No current flows into the open circuit, and the reflected wave carries all of the incident energy backward.

**Short circuit (  $Z_L = 0$  ):** The reflection coefficient equals  $\rho = -1$ . The entire wave reflects with inverted polarity. The reflected voltage subtracts from the incident voltage, forcing the boundary voltage to zero. The current doubles at the boundary, and the reflected wave again carries all incident energy backward.

## Impedance as a Property of the Medium

Electromagnetic waves do not possess their own independent impedance. Impedance is a strictly physical property of the copper geometry and the surrounding dielectric material.

The characteristic impedance of a trace dictates the strict mathematical ratio between voltage and current for any wave currently traveling on that specific geometry. A 50-ohm trace forces any propagating wave to maintain a ratio where voltage divided by current equals exactly 50 ohms. A reflected wave created by a localized defect immediately steps onto the trace heading backward. The trace takes complete physical control of that wave, locking its voltage and current into the 50-to-1 ratio demanded by the geometry.

The defect determines only the initial amplitude of the reflected wave. The reflection coefficient at the discontinuity dictates exactly how much voltage the defect rejects backward. Once launched, the reflected wave behaves exactly like a new signal traveling in the reverse direction. The physical geometry of the trace it currently occupies entirely defines its voltage-to-current relationship.

A reflected wave with a smaller voltage amplitude necessarily carries a proportionally smaller current. If a transmitter launches a 1-volt forward wave on a 50-ohm trace (producing 20 milliamps of forward current), and a parasitic via rejects 0.1 volts backward, the reflected wave's current is forced by the 50-ohm geometry to equal 2 milliamps (  $0.1\text{V}/50\Omega$  ). The reflected wave carries significantly less energy than the original forward wave.

## Architecture

---

### The Reflection as a Separate Electromagnetic Wave

A reflection is not the return current of the forward wave. A reflection is an entirely separate electromagnetic wave propagating in the backward direction.

The forward wavefront carries its own localized current loop: conduction current flows forward on the trace, displacement current bridges the dielectric at the wavefront location, and return current flows backward on the ground plane. This loop structure, described in the [Wave Propagation guide](#), closes at all times during forward propagation.

When the forward wavefront reaches an impedance discontinuity, the mismatch launches a new backward-traveling wave. This reflected wave possesses its own independent current loop: conduction current on the trace, displacement current through the dielectric at the reflected wavefront location, and return current on the ground plane. The reflected wave propagates backward toward the source, completely independent of the DC state established by the forward wave behind it.

## Ping-Pong Reflections and Energy Absorption

The reflected wave travels backward along the trace until it reaches the source. The driver transistor at the source possesses an internal resistance, typically between 10 and 30 ohms. The backward-traveling reflection encounters this internal source resistance as a physical boundary.

The reflection coefficient at the source determines how much of the returning wave is absorbed and how much re-reflects forward. If the source impedance matches the trace impedance, the source absorbs the returning wave completely and no re-reflection occurs. If the source impedance differs from  $Z_0$ , a portion of the wave re-reflects forward toward the load, and the process repeats.

The wave literally bounces back and forth between the source termination and the load termination. Each time the wave strikes a termination, the resistive element absorbs a fraction of the wave's electromagnetic energy and converts it permanently into thermal heat. The current of the arriving wave flows directly into the physical resistor material, and the energy dissipation destroys a portion of the wave amplitude.

The reflections become progressively smaller with each successive bounce. The exponentially decaying amplitude eventually drops below any measurable threshold, and the ping-ponging ceases.

## Settling to Steady State

The transient ping-pong process is the physical mechanism that bridges the gap between wave-propagation physics and DC circuit analysis.

When a voltage is first applied, the wavefront travels down the line, strikes the termination, and reflects backward. These reflections superimpose on one another as they bounce back and forth. The transient voltages ring down as the termination resistors absorb energy at each boundary.

Only after this infinite series of reflections converges does the entire line reach a uniform equipotential state. At that exact moment, the transient  $Z_0$  behavior ceases. The voltage is identical at every point along the trace. No voltage gradient exists to drive a propagating wave. The system has reached DC equilibrium, and Ohm's law (  $V = IR$  ) now perfectly describes the circuit.

The DC source acts as a constant-pressure reservoir holding the line at its intended voltage. The initial wavefront establishes this baseline. The reflections act as transient pressure ripples bouncing on top of that baseline. Once the ripples dissipate entirely into heat, the transmission line reaches a quiet steady state, held firmly at the source voltage.

Energy is lost to the DC resistance of the copper at all times, producing a voltage drop along the trace and generating heat. During the picosecond transient phase, however, the vast majority of the source energy is actively consumed by building the electric and magnetic fields in the surrounding dielectric, not by resistive heating.

## Worked Examples

---

### Open Circuit Response on a TDR

A time domain reflectometer contains an internal voltage source and an internal series resistor. The industry standard for this internal source resistance is exactly 50 ohms.

Connecting the TDR to a 50-ohm trace creates a simple voltage divider at the connector. The internal source generates a 1V step. The 1V source sits in series with the internal 50-ohm resistor, followed by the 50-ohm characteristic impedance of the trace. Equal resistances split the voltage evenly: 0.5V drops across the internal resistor as heat, and the remaining 0.5V launches onto the trace as an electromagnetic wave.

The TDR screen displays the voltage measured at the connector. It shows an instantaneous jump to 0.5V and remains flat at that level while the wavefront is actively traveling down a uniform 50-ohm trace. The ratio of voltage to current (  $V/I = Z_0$  ) remains constant across the uniform geometry, so no reflections are generated during propagation.

If the far end of the trace is completely disconnected (an open circuit), the 0.5V wavefront hits the open boundary and reflects completely (  $\rho = +1$  ). A +0.5V reflection travels backward toward the source. The TDR screen continues to display the 0.5V baseline until the exact moment this reflection arrives back at the connector. The returning +0.5V wave superimposes on the existing 0.5V baseline, and the screen jumps to 1.0V.

The returning wave then encounters the internal 50-ohm source termination. The perfect impedance match ( **50Ω** source to **50Ω** trace) completely absorbs the returning wave energy. No further bouncing occurs. The system achieves steady state, holding the full 1V potential across the entire trace.

Turning the TDR on without connecting any trace at all causes the connector itself to act as an open circuit with infinite impedance. Zero current flows, so zero voltage drops across the internal 50-ohm resistor. The full 1V potential appears immediately at the connector with no intermediate 0.5V stair. The round-trip flight time is zero, and the TDR screen shows a flat line at 0V followed by a single sharp edge jumping directly to 1V.

The steepness of that edge is limited by the instrument's bandwidth. A perfectly vertical edge would require zero rise time, which mathematically demands summing sine waves extending to infinite frequencies. A high-end TDR with 20 GHz bandwidth can only generate harmonics up to that limit, so the edge takes a finite time (perhaps 20 picoseconds) to climb from 0V to 1V. The slight slope of the transition is a direct consequence of the instrument's finite frequency content.

## Impedance Bump and Return to Baseline

A localized impedance discontinuity, such as a narrow trace segment measuring only a few millimeters, demonstrates how the TDR maps physical structure to reflections.

A positive impedance bump (impedance higher than  $Z_0$ ) produces a positive reflection coefficient. The reflected wave adds to the baseline, creating a voltage bump on the screen. A negative impedance dip (impedance lower than  $Z_0$ ) produces a negative reflection coefficient. The reflected wave subtracts from the baseline, creating a voltage dip.

The voltage returns to the baseline after the defect for a straightforward reason: the defect has a finite physical length. The forward wavefront hits the start of the narrow section and begins generating a positive reflection backward. The wavefront requires only a few picoseconds to traverse the two-millimeter defect. Exiting the narrow section returns the wavefront to the properly designed 50-ohm trace, which generates zero reflections. The defect physically stopped producing reflected energy the exact moment the wavefront cleared the bottleneck. The TDR screen, which plots reflected energy arriving at the instrument against time, draws a bump and then returns to the flat 0.5V baseline.

## Edge Cases

---

### Masking and Spreading of Downstream Discontinuities

Multiple impedance discontinuities along a trace do not appear independently on a TDR measurement. The first discontinuity physically alters the interrogating wavefront, degrading the instrument's ability to characterize everything downstream.

A parasitic inductor early in the trace severely strips the high-frequency content from the forward-traveling wave. The inductor generates a back-EMF proportional to the rate of current change ( $V = L \cdot dI/dt$ ). A fast, steep edge produces a large rate of change, triggering a strong inductive response that removes energy from the highest-frequency harmonics. The surviving wavefront continuing past the inductor possesses a much slower, rounded rising edge.

This degraded wavefront becomes the new interrogating signal for the remainder of the trace. Striking a second parasitic inductor further down the line with a slow, rounded edge produces a drastically different response. The slow edge possesses a small rate of change, so the second inductor generates only a weak opposing voltage. The TDR records this weak response as a small bump on the screen. The first discontinuity physically masks the true severity of all subsequent discontinuities.

The spatial resolution also degrades. A perfectly sharp voltage step hits a parasitic inductor instantaneously, and the inductor reacts with a massive, concentrated back-EMF spike. The TDR records this rapid event as a sharp, narrow feature. A degraded, rounded wavefront hits the same inductor gradually, producing a weak voltage spread out over a longer time interval. The TDR plots a wide, flattened mound on the screen instead of a sharp needle. The degraded rise time stretches out the physical duration of the interaction, smearing the feature and reducing the instrument's ability to pinpoint the defect's exact location.

This masking and spreading effect is one of the fundamental limitations of time domain reflectometry. A massive impedance error at the far end of a board might appear as a tiny perturbation on the TDR screen simply because the interrogating wavefront lost its high-frequency energy to discontinuities encountered earlier in the path.

## Parasitic Inductance and Capacitance Cancellation

Every millimeter of a trace possesses both capacitance and inductance simultaneously. Their continuous ratio defines the baseline characteristic impedance of the line. At a structural discontinuity, one parameter typically dominates the local geometry: a via introduces both extra metal area (capacitance) and vertical travel distance (inductance), but the proportions may not match the baseline ratio.

If a discontinuity introduces additional inductance and additional capacitance in the exact proportions required to maintain  $Z_0 = \sqrt{L/C}$ , the impedance remains matched. The TDR trace displays a flat line with no reflection, even though the physical geometry changed. Engineers exploit this phenomenon to optimize signal links. Deliberately routing a narrow, highly inductive trace segment adjacent to a wide, highly capacitive connector pad can maintain a uniform 50-ohm impedance profile through a region where either element alone would produce a severe mismatch.

## Chapter 3: Skin Effect and Dielectric Loss

---

A copper transmission line is not a perfect conductor, and the dielectric material surrounding it is not a perfect insulator. Both the metal and the insulator introduce frequency-dependent losses that attenuate the signal as it propagates. The two dominant loss mechanisms in high-speed PCB design are the skin effect (conductor loss) and dielectric absorption (insulator loss). Together, they cause the transmission line to behave as a physical low-pass filter, selectively attenuating the high-frequency harmonics that construct fast signal edges while passing lower-frequency content with relatively little distortion.

Understanding the physics of each mechanism separately, and then understanding how they combine, is essential for predicting signal degradation, selecting laminate materials, specifying surface roughness requirements, and designing equalization strategies that compensate for the damage.

### Core Concepts

---

#### Resistivity vs. Resistance

The distinction between resistivity and resistance is the foundation for understanding why the skin effect increases loss at high frequencies.

Resistivity (  $\rho$  ) is a constant chemical property of the copper atomic lattice. It describes the innate opposition to electron flow provided by the material itself and does not change with geometry. Resistance (  $R$  ) describes the total opposition to current flow for a specific physical piece of copper, and it depends directly on the geometry of the path:

$$R = \rho \cdot \frac{L}{A}$$

where  $L$  is the length of the current path and  $A$  is the cross-sectional area available to carry that current.

A wide ground plane carrying DC current has an enormous cross-sectional area. Every available parallel path through the copper is utilized simultaneously, driving  $A$  in the denominator to its maximum value and minimizing the total resistance. Any mechanism that reduces the usable cross-sectional area increases the resistance proportionally. The skin effect is precisely such a mechanism: it forces high-frequency current out of the interior of the conductor and into a thin surface layer, drastically reducing the effective  $A$  and increasing the AC resistance.

#### The Skin Effect: Eddy Current Cancellation

The skin effect is driven entirely by internal magnetic fields and induced eddy currents within the conductor.

A primary signal current flowing through a copper trace generates a surrounding magnetic field. The Right-Hand Rule dictates that this field forms concentric circles around the current axis. A longitudinal cross-section through a cylindrical conductor reveals the magnetic field pointing in one direction on the top half of the wire and the opposite direction on the bottom half.

When the primary current changes rapidly (as with a high-frequency signal), these internal magnetic fields strengthen rapidly. Faraday's Law of Induction dictates that a changing magnetic field cutting through a solid conductor induces a localized voltage. Inside the bulk copper, this changing magnetic field induces microscopic, circular currents known as eddy currents.

Lenz's Law mandates that these induced eddy currents must flow in a direction that opposes the original change that created them. The physical geometry of these circular loops is where the cancellation mechanism originates.

Consider the top half of the wire, where the primary magnetic field is strengthening. Lenz's Law requires the copper to generate an opposing field in that region. Applying the Right-Hand Rule determines the direction of the induced eddy current loop. Tracing the physical path of this loop through the conductor reveals a critical asymmetry:

- The portion of the eddy current loop located deep inside the center of the wire flows in the **opposite direction** of the primary signal current, causing cancellation.
- The portion of the loop located at the outer surface flows in the **same direction** as the primary signal current, causing reinforcement.

Examining the bottom half of the wire produces the identical result through the same Faraday-Lenz-Right-Hand-Rule chain, with appropriately reversed field directions and loop orientations. The net effect is universal throughout the conductor volume: eddy currents subtract from the primary current at the core and add to it at the surface.

Each eddy current forms a perfectly closed circular loop. A closed loop traveling in a circle contributes exactly zero net current from one end of the wire to the other. The side flowing backward cancels the side flowing forward. The eddy currents do not create new net current, nor do they destroy existing net current. They act strictly as a physical redistribution mechanism, moving the active current density from the interior to the surface.

DC signals produce zero rate of change in the magnetic field, generate zero eddy currents, and utilize the entire copper cross-section. High-frequency signals generate intense eddy currents that banish the active current flow to the outermost microscopic skin of the metal.

## The Exponential Gradient

The cancellation is not binary. It manifests as a smooth, continuous exponential gradient across the conductor cross-section.

The internal magnetic field strength inside a solid conductor is not uniform. The field is essentially zero at the exact dead center and grows steadily stronger moving outward toward the surface. A stronger magnetic field at the outer radii generates proportionally stronger eddy currents. Trillions of these overlapping circular eddy currents exist simultaneously throughout the copper volume.

Summing the effects of all overlapping loops produces an exponential decay function for the current density:

$$J = J_0 e^{-d/\delta}$$

where  $J_0$  is the maximum current density at the absolute outer surface,  $d$  is the physical depth into the conductor, and  $\delta$  is the skin depth constant. At one skin depth ( $d = \delta$ ), the current density has fallen to  $1/e$  (approximately 37%) of its surface value. At five skin depths, the current is effectively zero.

The skin depth is inversely proportional to the square root of the frequency:

$$\delta \propto \frac{1}{\sqrt{f}}$$

A 10 GHz signal is forced into a microscopic outer ring of the copper trace. A 1 GHz signal penetrates deeper, utilizing more copper. A DC signal uses the entire cross-section with no restriction. The practical consequence is that high-frequency signal components encounter dramatically higher AC resistance ( $R_{AC}$ ) than low-frequency components traveling on the same physical trace.

## The Low-Pass Filter Effect

The frequency-dependent resistance created by the skin effect causes the transmission line to function as a physical low-pass filter. Applying a voltage step to the trace launches a wavefront containing frequency harmonics spanning from DC to many gigahertz (the bandwidth determined by the rise time). The low-frequency components of that step propagate with minimal attenuation because they penetrate deep into the copper and encounter low resistance. The high-frequency components are squeezed into the thin surface layer, encounter massive AC resistance, and lose amplitude over distance.

The result is progressive degradation of the signal edge. The steep, sharp transition that left the transmitter arrives at the receiver with its high-frequency harmonics attenuated, producing a slower, rounded edge. The data rate and pattern have not changed, but the edge quality has been physically damaged by the conductor itself.

## Skin Effect and Loop Inductance

The skin effect also influences the return path geometry and the total loop inductance of the transmission line.

At high frequencies, the return current on the ground plane concentrates on the top skin of the metal, riding the surface closest to the signal trace. This bunching minimizes the vertical dielectric gap between the outbound and return currents, tightening the current loop. Inductance is proportional to the magnetic flux enclosed by the loop ( $L = \Phi/I$ ), so a smaller geometric loop captures fewer field lines and reduces the total loop inductance.

Narrowing the active current distribution does increase the self-inductance of the return path slightly (the surrounding magnetic field becomes denser when current is concentrated in a smaller area). This penalty is real but small compared to the massive reduction in mutual loop area achieved by the tight bunching. The current distribution organically settles at the exact profile that minimizes total mathematical impedance, balancing the competing self-inductance and loop-inductance terms.

## Surface Roughness

The exponential current density profile places the highest current density at the absolute outer surface of the conductor. Any physical irregularity at that surface forces the current to travel a longer, more tortuous path around microscopic peaks and valleys of the copper grain structure. This additional path length increases the effective AC resistance beyond what the smooth-conductor skin depth formula predicts.

The Hammerstad-Jensen correction factor quantifies this roughness penalty by comparing the RMS surface roughness to the skin depth. When the roughness features are much smaller than the skin depth (at lower frequencies), the correction is negligible. When the roughness features approach or exceed the skin depth (at high frequencies, where the skin depth shrinks to micrometers), the correction becomes substantial. Low-loss PCB designs intended for multi-gigahertz operation specify smoother copper foil profiles (such as very-low-profile or hyper-low-profile copper) specifically to minimize this roughness-induced loss.

## Architecture

---

### Dielectric Loss: Molecular Rotation and Energy Absorption

The dielectric insulator surrounding the conductor introduces a second, independent loss mechanism. The dielectric material (FR4 fiberglass, or more advanced laminates) contains polar molecules with asymmetric charge distributions. These molecules respond physically to the electric field of the propagating electromagnetic wave.

When the wavefront passes through a section of dielectric, the intense electric field forces the polar molecules to rotate and align with the field direction. This alignment absorbs kinetic energy from the wave. The molecules do not simply snap into position; the rotation involves physical work against intermolecular forces, and a portion of the wave's energy is dissipated as heat in the process.

The loss tangent (  $\tan \delta$  , also called the dissipation factor or  $D_f$  ) quantifies this energy absorption as the ratio of energy dissipated per cycle to energy stored per cycle. A higher loss tangent means the dielectric absorbs more of the wave's energy during each oscillation of the electric field. Standard FR4 has a relatively high loss tangent (approximately 0.02), making it unsuitable for multi-gigahertz serial links. Low-loss laminates such as Megtron 6 achieve loss tangents below 0.005 by using molecular structures that resist rotation more effectively.

The dielectric loss scales directly with frequency. Each cycle of the electric field forces the molecules through one full rotation-and-relaxation event, so higher-frequency signals force more absorption events per unit time. This frequency dependence makes dielectric loss, like the skin effect, a contributor to the transmission line's low-pass filter behavior.

### The Relaxation Asymmetry

A subtle but physically significant asymmetry exists in the dielectric response. The rising edge of a signal brings a massive, highly organized external force. The intense electric field of the wavefront acts on the polar molecules with the full driving power of the transmitter voltage, forcing them into strict alignment rapidly.

The falling edge simply removes this external force. The transmitter does not actively push the molecules back into their original random arrangement; it stops holding them in place. The molecules must rely entirely on ambient thermal energy to return to their neutral state. Random thermal jostling is an extremely weak and slow physical process compared to a directed electromagnetic force.

The polarization (alignment) is fast because it is actively driven. The relaxation (return to random orientation) is slow because it is passively thermal. This speed difference creates a temporal asymmetry in the dielectric response: the molecules release their stored energy gradually over a period that extends well beyond the duration of the original pulse.

## The Trailing Wake and Inter-Symbol Interference

The slow relaxation creates a physical effect analogous to the trailing wake behind a moving boat. The leading edge of a pulse creates a sharp, immediate disturbance as the molecules snap into alignment. After the pulse passes and the transmitter drops the voltage, the molecules slowly drift back to their resting state, gradually releasing their stored energy back into the copper trace.

This gradual energy release produces a lingering voltage tail that trails behind the pulse. The next data pulse arriving at the receiver rides on top of this decaying tail from the previous pulse. The receiver measures the total combined voltage at the sampling instant, and the residual energy from the old pulse physically adds to (or subtracts from) the new pulse voltage. This corruption of the receiver threshold voltage is trailing inter-symbol interference (ISI).

The asymmetry between the fast leading disturbance and the slow trailing wake is the physical reason why transmitter pre-emphasis (FFE) typically requires much stronger post-cursor correction taps than pre-cursor taps. The pre-cursor tap only needs to sharpen the immediate leading edge. The post-cursor taps must subtract voltage aggressively to flatten out the long, persistent trailing wake created by the dielectric relaxation.

## Worked Examples

---

### Frequency-Dependent Attenuation on a PCB Trace

Consider a 10 Gbps NRZ signal with a 30-picosecond rise time propagating on a 12-inch FR4 microstrip trace. The alternating 1010 pattern has a fundamental frequency of 5 GHz, but the 30 ps rise time requires harmonic content extending past 10 GHz to construct the edge shape.

The 5 GHz fundamental experiences moderate skin effect loss and moderate dielectric absorption. It arrives at the receiver attenuated but still with significant amplitude. The 10 GHz third harmonic encounters a skin depth that has shrunk by a factor of  $\sqrt{2}$  compared to the 5 GHz case, roughly doubling the AC resistance contribution. The dielectric absorption at 10 GHz is also double that at 5 GHz, since dielectric loss scales linearly with frequency.

The combined effect is that the high-frequency harmonics responsible for the steep edges are attenuated far more heavily than the fundamental. The received signal retains its 5 GHz repetition rate but arrives with rounded, slower transitions. The eye diagram closes vertically (reduced voltage margin) and horizontally (increased timing uncertainty at the decision threshold).

## DC vs. High-Frequency Return Path Geometry

A DC test current injected into a trace spreads broadly across the ground plane, utilizing the maximum available copper cross-section. The current takes the shortest geometric path from source to sink, ignoring the trace routing entirely, because the inductance term in the impedance equation ( $j\omega L$ ) vanishes at zero frequency. Only DC resistance matters, and spreading maximizes  $A$  to minimize  $R$ .

The same trace carrying a 10 GHz signal produces a dramatically different return path. The massive frequency multiplier ( $\omega$ ) causes the inductive reactance ( $j\omega L$ ) to dominate the impedance equation completely. The return current ignores the vast expanse of available copper and crowds into a tight, narrow band directly underneath the signal trace, perfectly shadowing its physical routing. The current voluntarily accepts a higher DC resistance penalty to achieve the lowest possible loop inductance. The skin effect further confines this return current to the top surface of the ground plane (the side facing the trace), minimizing the vertical separation and tightening the loop to its physical minimum.

## Edge Cases

---

### Combined Loss Budget

In practice, the total insertion loss of a transmission line channel is the sum of conductor loss (skin effect plus surface roughness) and dielectric loss. At lower frequencies (below approximately 1 GHz on standard FR4), conductor loss dominates. At higher frequencies, dielectric loss grows linearly while conductor loss grows as the square root of frequency, so dielectric loss eventually overtakes conductor loss and dominates the total budget.

This crossover frequency depends on the specific laminate material and copper roughness profile. Upgrading from standard FR4 to a low-loss laminate dramatically reduces the dielectric loss contribution but does nothing for conductor loss. Specifying smoother copper foil reduces the roughness contribution but does nothing for dielectric loss. Achieving acceptable total loss at multi-gigahertz data rates typically requires addressing both mechanisms simultaneously: low-loss laminate paired with low-roughness copper.

### The Skin Depth Floor

The skin depth formula predicts that the active current layer becomes arbitrarily thin at sufficiently high frequencies. In practice, the copper surface roughness imposes a floor. Once the skin depth shrinks to the same scale as the RMS surface roughness (typically a few micrometers), the current can no longer be treated as flowing in a smooth, well-defined layer. The rough surface topology scatters the current, and the simple exponential decay model breaks down. Electromagnetic simulation tools that account for the three-dimensional roughness profile (rather than applying a scalar correction factor) provide more accurate loss predictions in this regime.

## Chapter 4: Inductance and Magnetic Coupling

---

Inductance governs the magnetic dimension of every high-speed transmission line. The electric field stores energy in the dielectric between the trace and the ground plane, defining capacitance. The magnetic field stores energy in the current loop formed by the outbound trace and the return path, defining inductance. The ratio of these two distributed parameters sets the characteristic impedance, and the product of them sets the propagation velocity. Controlling inductance is therefore inseparable from controlling signal integrity.

The concept divides naturally into two distinct mechanisms. Self-inductance describes the opposition that a single conductor's own magnetic field generates against changes in its own current. Mutual inductance describes the magnetic coupling between two separate conductors, where the changing field of one induces a voltage in the other. In a transmission line, both mechanisms act simultaneously. The total loop inductance of the system depends on the self-inductance of the outbound trace, the self-inductance of the return plane, and the mutual inductance bridging the gap between them. Engineers who understand how physical geometry controls each of these three quantities can predict and manage the inductive behavior of any PCB structure.

This guide develops the physics from the microscopic scale upward: the origin of macroscopic fields from individual electron contributions, the internal field gradient inside a solid conductor, the mathematical structure of self-inductance and mutual inductance, the Reciprocity Theorem that enforces symmetric coupling, and the total loop inductance equation that ties the entire system together.

### Core Concepts

---

#### The Origin of Macroscopic Fields: Superposition

Every moving electron in a conductor generates its own microscopic magnetic field. These individual fields are tiny, perfectly circular loops (often visualized as "hula hoops") centered on each electron's trajectory. The macroscopic magnetic field that engineers measure with instruments is the vector sum of trillions of these microscopic contributions. This is the Principle of Superposition: the total field at any point in space is the mathematical addition of every individual electron's field at that point.

The shape of the physical conductor dictates the geometric arrangement of the moving electrons, and that arrangement determines how the microscopic fields combine.

Packing all the electrons into a tight cylindrical wire forces their individual hula hoops into close proximity. The fields add constructively, producing tight, dense concentric circles immediately outside the wire. Spreading the same electrons across a massive flat copper plane changes the superposition pattern. The horizontal components of adjacent hula hoops cancel each other, while the vertical components add together, transforming the macroscopic field into a wide, diffuse oval rather than a compact circle.

The microscopic physics of each electron remains identical in both cases. The geometry of the copper simply determines which components of the individual fields reinforce and which cancel. Compressing a conductor does not alter the fundamental behavior of any single electron; it packs the charge carriers into a tighter bundle, forcing their superimposed macroscopic field to occupy a much denser volume.

## Self-Inductance and the Effect of Trace Width

Self-inductance measures the magnetic opposition that a conductor's own field exerts against changes in its current. The physical mechanism is direct: a changing current generates a changing magnetic field, which (by Faraday's Law) induces a voltage that opposes the original current change. The magnitude of this opposition depends on how much magnetic flux the conductor's geometry allows it to generate per ampere of current.

Trace width has a pronounced effect on self-inductance. Ampere's Law dictates that magnetic field strength is inversely proportional to distance from the current source. In a wide trace, the current spreads across a large cross-section, producing a broad, diffuse magnetic field with significant internal cancellation between the left and right edges. A narrow trace forces the same total current through a small physical bottleneck. The magnetic field lines crowd extremely close to the concentrated charge flow, creating an intense, spatially compressed field that stores more magnetic energy per unit length. The result is a higher self-inductance for the narrow trace.

This relationship has a practical interpretation: a narrow trace acts like a series inductor. The dense magnetic field surrounding the bottleneck physically resists rapid changes in current. High-frequency signals, which carry an astronomically high rate of change (  $dI/dt$  ), encounter a violent opposing voltage (back-EMF) that chokes off the fast edge content. Low-frequency signals, with a near-zero rate of change, pass through unaffected. The narrow trace therefore behaves as a geometric low-pass filter, rounding the edges of the propagating wavefront and degrading the signal at the receiver.

## The Internal B-Field Gradient

The relationship between field strength and distance from the current takes on a distinctive character inside a solid conductor. Outside the wire, the field decreases with distance, matching the common intuition that fields weaken as you move away from the source. Inside the wire, the gradient runs in the opposite direction.

Ampere's Law states that the magnetic field strength along a circular boundary is proportional to the total current physically enclosed within that boundary. Applying this rule at different radii inside the conductor reveals a systematic gradient:

At the exact center, a measuring circle of zero radius encloses zero current. The magnetic field is exactly zero. Moving outward, the measuring circle captures progressively more cross-sectional area and therefore more current. The field strength grows steadily. At the surface, the measuring circle encompasses the entire cross-section and encloses 100 percent of the current, producing the maximum internal field strength.

The field increases linearly from zero at the center to its maximum at the surface. This linear behavior arises from a geometric ratio: the enclosed current grows with the area of the measuring circle (  $\propto r^2$  ), while the boundary length grows only with its circumference (  $\propto r$  ). Dividing the enclosed current by the circumference leaves a field strength proportional to the radius.

Stepping outside the wire, the enclosed current stops growing (all of it is already inside the boundary), and the field begins to fall off with distance.

## The Shell Theorem Analog: Why Outer Current Cancels

The zero-contribution of current outside the measuring radius is not an approximation; it is a rigorous consequence of cylindrical symmetry, analogous to the gravitational shell theorem in mechanics.

A solid conductor can be conceptually decomposed into a solid inner core of radius  $r$  and a hollow outer pipe surrounding it. The question is: what magnetic field does the current in the outer pipe generate at a point inside its cavity?

At any point inside the hollow pipe, competing magnetic forces act from every direction. The current in the pipe wall closest to the point generates a strong, localized field. The current on the far side of the pipe, being more distant, generates a weaker field pointing in the opposite direction. Geometry dictates that there is more copper (and more total current) spanning the wide arc on the far side than on the narrow arc nearby. The larger magnitude of distant current perfectly offsets the stronger influence of the close current. Integrating the field contributions from every degree of the circular pipe wall reveals exact, uniform cancellation. The net magnetic field anywhere inside the cavity of a hollow current-carrying cylinder is precisely zero.

A solid wire is an infinite series of concentric hollow tubes nested together. At any internal radius  $r$ , all the copper outside that radius behaves as a collection of these hollow tubes, each contributing zero net field to the interior. The only current capable of generating a measurable magnetic field at radius  $r$  is the current flowing physically inside radius  $r$ .

The copper medium does not change across the boundary, and the electrons do not behave differently depending on where an imaginary line is drawn. The boundary simply divides the uniform electron population into two geometric groups: electrons located farther from the center than the measurement point surround it symmetrically, forcing their individual field vectors to cancel at that coordinate; electrons located closer to the center fail to surround the point, so their field vectors add constructively and produce the measurable field.

## Mutual Inductance and Geometry

Mutual inductance ( $M$ ) measures the efficiency of magnetic coupling between two separate current paths. Current flowing through the outbound trace generates a magnetic field that expands into three-dimensional space. A portion of those field lines physically intersect the return plane. Mutual inductance quantifies exactly how much of that flux successfully couples to the second conductor.

This coupling efficiency is entirely dictated by physical geometry. A thick dielectric separating the trace from the ground plane allows a large fraction of the magnetic field to dissipate into the surrounding material before reaching the return path. The resulting poor magnetic coupling yields a low mutual inductance. Compressing the dielectric gap moves the ground plane closer to the trace, allowing it to intercept a much larger fraction of the expanding field lines and drastically increasing the mutual inductance.

## The Reciprocity Theorem ( $M_{12} = M_{21}$ )

A narrow trace and a wide ground plane are physically very different shapes. The trace generates a dense, localized magnetic field; the plane generates a broad, diffuse one. It seems counter-intuitive that these two conductors would share the same mutual inductance value.

The Reciprocity Theorem resolves this apparent asymmetry. Mutual inductance is a shared geometric property of the entire two-conductor system, not an attribute of either conductor individually. The theorem guarantees that the magnetic coupling between any two arbitrary shapes is perfectly symmetric, regardless of their individual physical profiles.

The underlying mechanism is a precise geometric balance. A highly concentrated magnetic source (the narrow trace) projects its field onto a massive receiving area (the wide plane). The field weakens over distance, but the enormous surface area of the plane acts as a giant net, intercepting the weaker, expanding field lines. The total intercepted flux reaches some definite value.

Reversing the experiment, the wide plane generates a weak, diffuse field spread across its entire surface. The narrow trace presents a tiny cross-sectional area, acting as a very small net. Intercepting a weak field with a small net seems destined to produce negligible flux. However, every microscopic slice of current distributed across the vast width of the plane contributes a tiny field vector that reaches up and touches the trace. Integrating all of these contributions from the full width of the plane reveals an exact mathematical balance: the total flux coupled to the narrow trace equals the total flux coupled to the wide plane. The physical shapes dictate the shape of each conductor's magnetic field, but the integrated coupling ratio remains a perfectly symmetric geometric constant.

This symmetry is why the total loop inductance equation contains the term  $2M$  rather than  $M_{12} + M_{21}$  : the two values are identical by physical law.

## Architecture

---

### The Total Loop Inductance Equation

The magnetic behavior of a complete transmission line circuit is captured by the total loop inductance:

$$L_{loop} = L_1 + L_2 - 2M$$

where  $L_1$  is the self-inductance of the outbound trace,  $L_2$  is the self-inductance of the return plane, and  $M$  is the mutual inductance between them.

This formula derives directly from Faraday's Law and the voltage drops within the current loop. Inductance translates a changing current into a voltage:  $V = L(dI/dt)$  . The total voltage drop around the loop is the sum of four contributions:

1. The outbound trace produces a self-induced voltage of  $L_1(dI/dt)$  .
2. The return plane produces a self-induced voltage of  $L_2(dI/dt)$  .
3. The outbound trace's magnetic field cuts through the return plane. The return current flows in the opposite direction (  $-dI/dt$  ), so this mutually induced voltage is  $-M(dI/dt)$  .
4. Symmetrically, the return plane's magnetic field cuts through the outbound trace, inducing another  $-M(dI/dt)$  .

Adding all four components and factoring out the current change yields the total loop inductance. The engineering consequence is immediate: physically compressing the dielectric gap maximizes  $M$ , which mathematically cancels the self-inductance contributions and drives the total loop inductance toward zero.

The quantities  $L_1$  and  $L_2$  are not assigned to the entire length of the trace and plane at once. They describe an infinitesimally small geometric slice of the transmission line (typically modeled as one millimeter long). The per-unit-length parameters are integrated along the full trace length to model the complete board.

## Loop Inductance vs. Self-Inductance

Self-inductance in isolation measures the magnetic opposition generated by a single conductor suspended in an infinite vacuum, disconnected from any return path. It answers a purely theoretical question about one wire's intrinsic magnetic storage.

Loop inductance measures the total magnetic opposition of the complete, closed circuit that exists in reality. It incorporates the outbound conductor, the return conductor, and the mutual magnetic interaction bridging the space between them. The distinction matters because no real signal exists without a return path. The self-inductance of an isolated trace is a mathematical abstraction; the loop inductance is the physical quantity that determines signal behavior.

High-frequency current crowds tightly underneath the trace precisely to minimize loop inductance. The self-inductance of the return path increases slightly as the current concentrates into a narrow band, but this penalty is overwhelmed by the massive reduction in loop area and the corresponding increase in mutual inductance. The current organically settles into the exact spatial distribution that minimizes total impedance at each frequency.

## Loop Resistance and Why It Lacks Spatial Interaction

Resistance and inductance respond to conductor spacing in fundamentally different ways.

The geometric resistance of a conductor is  $R = \rho \cdot (L_{trace}/A)$ , where  $\rho$  is the copper resistivity,  $L_{trace}$  is the trace length, and  $A$  is the cross-sectional area of the current path. Maximizing  $A$  by spreading the return current across the full width of the ground plane places a large number in the denominator, driving resistance to its minimum.

Magnetic fields project outward into three-dimensional space. The physical distance separating the outbound and return paths drastically alters the mutual inductance. Resistance, by contrast, remains entirely confined within the physical boundaries of the copper. Moving a return trace closer to an outbound trace produces zero change in total series resistance. The spacing does not affect the pure addition of the two resistive values. The concept of an interacting "loop" is physically irrelevant for resistance calculations, which is why the term "loop resistance" rarely appears in engineering literature, while "loop inductance" is a core design parameter.

This asymmetry explains the frequency-dependent behavior of return current distribution. At low frequencies, inductive reactance (  $X_L = 2\pi fL$  ) vanishes, leaving DC resistance as the dominant impediment. The current spreads across the full ground plane to maximize cross-sectional area and minimize resistance. At high frequencies, the inductive reactance grows to dominate the impedance budget. The current crowds under the trace to minimize loop area and loop inductance, accepting a higher resistance penalty as the cost of reducing the overwhelmingly larger inductive term.

## Worked Examples

---

### Proving Reciprocity: Narrow Trace Driving a Wide Plane

The mutual inductance formula  $M = \Phi/I$  provides a concrete framework for verifying that  $M_{AB} = M_{BA}$  .

**Setup:** Trace A is a narrow, focused copper wire. Trace B is a massive, wide copper plane. They are positioned parallel to each other, separated by a vertical distance  $D$  .

**$M_{AB}$  (Trace drives Plane):** Injecting exactly 1 Ampere into the narrow Trace A generates a dense, concentrated magnetic field radiating outward in tight concentric circles. The field weakens over the distance  $D$  , but the vast surface area of the wide plane below acts as a large net, intercepting the weakening field lines across its full extent. Summing all intercepted flux yields  $\Phi_B = 50$  Webers (hypothetical value). The mutual inductance is  $M_{AB} = 50/1 = 50$  Henrys.

**$M_{BA}$  (Plane drives Trace):** Reversing the experiment, 1 Ampere flows through the wide plane, generating a weak, diffuse field at any localized point above it. The narrow trace presents a tiny cross-sectional area. Intercepting a weak field with a small net appears destined to produce negligible flux. However, every microscopic current element across the full width of the plane contributes a tiny field vector reaching up to the trace. Integrating these contributions from the entire width yields  $\Phi_A = 50$  Webers. The mutual inductance is  $M_{BA} = 50/1 = 50$  Henrys.

The geometric balance is exact. A concentrated source projecting onto a large receiver produces the same total flux as a diffuse source projecting onto a small receiver. The coupling ratio is a symmetric geometric constant.

### Narrow Trace Bottleneck as a Series Inductor

Consider a 50-ohm microstrip trace with a manufacturing defect that narrows a short section of the trace. The properly designed section spreads the signal current across a calculated width, producing a diffuse magnetic field with significant internal cancellation between the left and right edges.

At the narrow bottleneck, the same current is forced through a reduced cross-section. The magnetic field lines compress into a dense ring surrounding the bottleneck, storing more magnetic energy per unit length and locally increasing the self-inductance.

The effect is frequency-dependent. A high-frequency wavefront carries a rapid  $dI/dt$ , which, when combined with the elevated local inductance, generates a large opposing back-EMF. This back-EMF partially blocks the high-frequency content of the edge, reflecting energy backward. A low-frequency signal has a near-zero  $dI/dt$ ; the dense magnetic field remains static and generates no opposition, so the low-frequency content passes through encountering only a trivial increase in DC resistance.

The net effect at the receiver is a signal with its high-frequency building blocks stripped away: a rounded, sloped transition instead of a sharp edge. On a TDR display, the instrument records the localized impedance increase as a positive bump, pinpointing the exact physical location of the bottleneck.

## Edge Cases

---

### Two Geometries, One Inductance Increase

Trace width and loop area both affect inductance, but through different physical mechanisms. Confusing the two is a common source of error.

A *narrow trace width* increases self-inductance by concentrating the magnetic field. All of the current flows through a small cross-section, eliminating the parallel cancellation that exists in a wide trace. The field density rises, and more magnetic energy is stored per unit length.

A *wide loop area* (large vertical separation between the trace and the return plane) increases loop inductance by allowing the magnetic field to occupy a larger physical volume. More field lines are enclosed by the loop geometry, yielding more total flux per ampere.

Both mechanisms increase inductance, but they are geometrically independent. A designer can have a wide trace (low self-inductance) on a thick dielectric (high loop inductance), or a narrow trace (high self-inductance) on a thin dielectric (low loop inductance). Optimal design minimizes both: a correctly calculated trace width avoids unnecessary field concentration, and a thin dielectric maximizes mutual coupling to drive the loop inductance term toward zero.

### The Parasitic Via Capacitor vs. the Parasitic Narrow Trace

Both parasitic capacitance and parasitic inductance degrade high-frequency edge content, but through opposite mechanisms. A parasitic via capacitor provides a parallel escape route, shunting high-frequency energy into the ground plane and draining it from the signal path. A parasitic narrow trace acts as a series barrier, generating inductive back-EMF that blocks high-frequency energy and reflects it backward toward the source.

The measured result at the receiver is identical in both cases: a transition missing its high-frequency harmonics, producing a slower, more rounded edge. The distinction matters for diagnosis. On a TDR display, the capacitive parasitic appears as a negative dip (local impedance decrease), while the inductive parasitic appears as a positive bump (local impedance increase).

## Chapter 5: Return Path Dynamics and Parasitic Effects

---

Every high-speed signal exists as an electromagnetic wave propagating through the dielectric between a trace and a reference plane. The signal current on the trace and the return current on the ground plane travel together as two halves of a single closed loop, connected at the wavefront location by displacement current bridging the insulator. The physical geometry of this current loop determines the distributed inductance and capacitance of the transmission line, which in turn set the characteristic impedance and propagation velocity.

The return current does not follow a fixed path. Its spatial distribution shifts with frequency. Low-frequency currents spread across the entire ground plane to minimize resistance. High-frequency currents crowd into a narrow band directly beneath the signal trace to minimize loop inductance. Any physical structure that disrupts the return path or introduces unintended geometry into the loop creates a parasitic element: a capacitor, an inductor, or both. These parasitics produce frequency-dependent impedance mismatches that reflect energy, degrade edge quality, and corrupt signal integrity at multi-gigahertz data rates.

This guide traces the physics of return path behavior from its electromagnetic origin through the frequency-dependent transition, then examines the specific parasitic structures that arise from PCB geometry: vias, anti-pads, neck-downs, and pad structures. The treatment connects the abstract concept of reactive impedance to the concrete physical mechanisms that produce it.

### Core Concepts

---

#### The Current Loop and Its Electromagnetic Origin

Current requires a closed loop to flow. A transmitter launching a voltage step into a trace pushes current into the top copper while simultaneously pulling an equal amount of current from the ground plane. The active current loop travels outward along the trace, drops vertically through the dielectric as displacement current at the exact location of the propagating wavefront, and returns immediately along the ground plane to the source.

The return current does not scatter randomly across the ground plane. The electrons injected into the ground plane at the source are strongly attracted to the positively charged trace sitting directly above them across the dielectric. This electromagnetic attraction forces them to flow directly underneath the trace, perfectly shadowing its physical routing. An equal and opposite wavefront of charge travels in lockstep with the expanding electric field on the trace above.

The electrons in the ground plane stop moving the moment the wavefront passes their specific physical location. They settle into a static position directly beneath the newly charged segment. The strong DC electric field spanning the dielectric holds them in place: the trace possesses a localized deficit of electrons, the ground plane directly below it possesses a matching surplus, and the insulator prevents recombination. They remain locked in this position until the transmitter sends a logic "0," which closes the pull-down transistor, creates a direct path between the trace and ground, and allows the trapped surplus electrons to rush back and neutralize the charge disparity.

## Geometric Origins of Distributed Capacitance and Inductance

The capacitance and inductance of a transmission line are not soldered components. They emerge from the physical geometry of the copper and the dielectric.

Capacitance requires two conductors separated by a dielectric. The signal trace serves as one plate, the adjacent ground plane serves as the second, and the PCB laminate acts as the insulating layer between them. This geometry stores energy in the electric field. Widening the trace increases the surface area facing the ground plane, increasing the capacitance. Moving the trace closer to the ground plane achieves the same effect.

Inductance requires current flowing in a loop. The signal current traveling down the trace and the return current traveling back along the ground plane form a geometric loop. The area enclosed by this loop stores energy in the magnetic field. Narrowing the trace concentrates the current into a tighter channel, producing a denser magnetic field and higher self-inductance per unit length. Moving the trace closer to the ground plane reduces the loop area, increasing mutual inductance and reducing the total loop inductance.

These two parameters are geometrically coupled. Widening a trace increases the plate area (raising capacitance) while simultaneously providing a broader current path (lowering self-inductance). Moving the trace closer to the ground plane increases capacitance and decreases the loop area. The continuous ratio of distributed inductance to distributed capacitance per unit length determines the characteristic impedance:  $Z_0 = \sqrt{L/C}$ .

## The Impedance Cancellation Principle

Every millimeter of a properly designed trace possesses both capacitance and inductance simultaneously. Their ratio defines the baseline characteristic impedance. At a structural discontinuity, one parameter typically dominates the local geometry. A via introduces both extra metal area (capacitance from the barrel-to-plane gap) and vertical travel distance (inductance from the narrow vertical current path).

A component that introduces additional inductance and additional capacitance in the exact proportions required to maintain the  $L/C$  ratio keeps the impedance matched. The TDR trace displays a flat line with no reflection, even though the physical geometry changed. Engineers exploit this phenomenon to optimize signal links, deliberately routing a narrow, highly inductive trace segment adjacent to a wide, highly capacitive connector pad to maintain a uniform 50-ohm impedance profile through a region where either element alone would produce a severe mismatch.

## Frequency-Dependent Return Path Selection

The return current's choice of path is governed by the total impedance equation:  $Z = R + j\omega L$ . The behavior transitions smoothly as a function of frequency.

At DC and low frequencies, the angular frequency  $\omega$  approaches zero. The inductive reactance term ( $j\omega L$ ) vanishes, leaving DC resistance ( $R$ ) as the sole impediment. The current minimizes this resistance by spreading across the full width of the ground plane, taking a direct straight-line path from source to sink. Maximizing the cross-sectional area of the current flow places the largest possible number in the denominator of the resistance equation ( $R = \rho \cdot L_{trace}/A$ ), driving total resistance to its minimum.

At high frequencies,  $\omega$  grows to dominate the impedance budget. The inductive reactance ( $j\omega L$ ) completely overwhelms the DC resistance, rendering  $R$  mathematically irrelevant. The current must now minimize the physical inductance ( $L$ ). Loop inductance depends on the geometric area separating the outbound signal current from the return current. A wider spatial loop captures more magnetic flux and produces more inductance. The high-frequency return current crowds into a tight, narrow band directly underneath the signal trace to close this spatial gap, perfectly shadowing the trace routing regardless of how narrow or physically long that path becomes. It voluntarily accepts a higher DC resistance penalty to achieve the absolute lowest possible loop inductance.

This behavior is not a binary switch but a smooth transition. At intermediate frequencies, both the resistance and inductance terms contribute meaningfully to the impedance, and the current distribution reflects a balance between the two optimization pressures. The skin effect further confines the return current to the top surface of the ground plane (the side facing the trace), minimizing vertical separation and tightening the loop to its physical minimum.

## Loop Area and the Ballooning Effect

The geometric loop area has a simple physical definition: the cross-sectional region bounded by the outbound signal path and the return current path. Visualizing a cross-section of the PCB, the current travels forward along the top trace, drops vertically through the termination resistor at the far end, travels back along the ground plane, and rises vertically at the source. This forms a very long, extremely thin rectangular loop. The enclosed area equals the trace length multiplied by the dielectric thickness.

If the return current encounters a slot or hole in the ground plane, it cannot jump the gap. It must detour horizontally around the obstacle. The return path is no longer directly underneath the trace. The physical loop bows outward, changing from a tight vertical rectangle into a wide, sprawling polygon. This drastically increases the enclosed surface area, which increases the total loop inductance and degrades the impedance match. Engineers call this effect "loop ballooning," and it is one of the most damaging layout errors in high-speed PCB design.

The magnetic field geometry inside the loop is not uniform. Around the narrow trace, the field lines form tight concentric circles. Around the flat ground plane, the field lines stretch into wide, diffuse ovals. At high frequencies, the skin effect prevents magnetic fields from penetrating deep into the copper, and the field lines wrap around the exterior surfaces. The high-frequency return current crowds beneath the trace not to minimize its own width, but to get as physically close to the outbound trace as possible, minimizing the enclosed geometric area and the total flux linkage.

## Architecture

---

### Parasitic Structures: Unintentional Circuit Elements

Physical structures on a PCB present reactive impedances because their specific geometries force electrical energy to be temporarily stored in electric or magnetic fields. A via is a vertical copper barrel passing through a clearance hole in a solid ground plane. Two metal conductors (the barrel and the surrounding plane) separated by a dielectric insulator (FR4 filling the clearance gap) form the exact physical definition of a capacitor. Simultaneously, the signal current flowing vertically through the narrow barrel generates a concentrated magnetic field, forming the exact physical definition of an inductor.

These unintended physical structures are called parasitics. They are not soldered components; they are geometric artifacts that behave exactly like circuit elements. A parasitic inductor acts as an unwanted series component in the signal path. A parasitic capacitor acts as an unwanted shunt component bridging the signal to ground.

PCB elements also present resistive impedance. The copper making up a via possesses a measurable DC resistance. Signal integrity engineers largely ignore this resistance in high-speed analysis because it is exceptionally small (often less than a milliohm) and constant across frequency. A one-milliohm resistor in series with a 50-ohm transmission line creates an undetectable reflection. The reactive impedance is the primary concern because it scales dramatically with frequency.

### Reactive Impedance and Frequency Scaling

The frequency dependence of capacitive and inductive impedance originates directly from the fundamental defining equations of these components.

The capacitor equation is  $I = C \cdot (dV/dt)$  : current equals capacitance multiplied by the rate of change of voltage. The inductor equation is  $V = L \cdot (dI/dt)$  : voltage equals inductance multiplied by the rate of change of current. Applying a sinusoidal signal to these equations requires taking the mathematical derivative of the sine function. The chain rule extracts the frequency variable from inside the sine wave and places it as a multiplier in front, directly producing the impedance formulas:

$$Z_C = \frac{1}{2\pi fC} \qquad Z_L = 2\pi fL$$

The capacitive impedance is inversely proportional to frequency. A parasitic via capacitor might present 10,000 ohms at 1 kHz but plunge to 5 ohms at 10 GHz. This drastic, frequency-dependent drop causes a massive impedance mismatch against the 50-ohm trace. The inductive impedance scales in the opposite direction: low at low frequencies, massive at high frequencies.

The physical mechanics behind these mathematical relationships are grounded in energy storage and release.

**Capacitors:** Driving a low-frequency signal changes the voltage slowly. The metal plates have time to fill completely with charge. As electrons crowd onto the plate, electrostatic repulsion (like charges repel) builds a back-pressure that opposes further current flow. When the repulsive voltage from the crowded plate exactly equals the driving voltage, current drops to zero and the capacitor acts as an open circuit. This represents the high-impedance behavior at low frequencies. Driving a high-frequency signal reverses polarity before the plates can fill. The electrons shuttle back and forth continuously, never building enough charge to create significant back-pressure. Current flows freely, defining a low impedance.

**Inductors:** An inductor opposes changes in current by generating a back-EMF from its magnetic field. Changing the current slowly generates a weak opposing field, presenting low impedance. Attempting to change the current rapidly at high frequencies forces the inductor to generate a massive opposing field, presenting enormous impedance.

## The Via Anti-Pad: Anatomy of a Parasitic Capacitor

A via is a solid vertical copper barrel connecting traces on different PCB layers. The barrel must pass through internal ground planes without touching them, which would create a short circuit. PCB manufacturers drill a larger clearance hole through each ground plane to let the via pass safely. This empty circular void is called an anti-pad.

The via barrel plunges through the center of the anti-pad. The physical gap between the outer wall of the barrel and the inner circular rim of the ground plane is filled with FR4 fiberglass resin. This geometry creates an unintentional cylindrical capacitor: the barrel wall is one plate, the circular ground plane rim is the other, and the FR4 in the clearance gap is the dielectric.

The signal traveling vertically down the barrel sees two available paths:

**Path A (the copper highway):** The intended path requires the signal to flow straight down the solid barrel, exit the bottom, and continue along the destination trace. This provides continuous metal with near-zero resistance.

**Path B (the dielectric gap):** The high-speed voltage transition generates a rapidly expanding electric field that reaches horizontally across the FR4 clearance gap to the surrounding ground plane. Physical electrons never cross the insulator. The changing electric field projects across the gap and pushes electrons in the ground plane away, creating the illusion of current flow. This is displacement current: no electron crosses the gap, but the electric field reaches across to drive charge motion on the other side.

A low-frequency signal sees the parasitic capacitor as an impenetrable wall (millions of ohms of capacitive impedance). The current remains entirely on Path A and passes through the via undisturbed. A high-frequency signal sees the impedance of Path B plunge to a value comparable to or lower than the 50-ohm trace, creating a significant parallel shunt. During the brief picoseconds of a voltage transition, when  $dV/dt$  is astronomically high, a burst of displacement current leaks through Path B into the ground plane.

## Displacement Current and Capacitor Charging Physics

The displacement current mechanism deserves detailed examination because it governs the parasitic leakage at every via, pad, and plane transition in a high-speed design.

Electrons physically crowd onto one plate of the capacitor. Their combined negative charge projects a strong electric field across the physical gap of the insulator. This field penetrates the ground plane on the opposite side and repels the free electrons resting in the copper, forcing them to flow away. Electrons physically arriving on one plate, combined with completely different electrons being pushed away from the opposite plate, creates a continuous circuit without any particle crossing the gap.

The fully charged state occurs when the repulsive back-pressure of the crowded plate exactly equals the forward driving voltage. At this stalemate, net electron flow drops to zero. The displacement current ceases because the electric field stops changing. The capacitor acts as an open circuit, and the signal continues exclusively down the solid copper of Path A.

The practical consequence is that the parasitic via capacitor acts as a low-impedance shunt exclusively during voltage transitions. While the trace holds a steady 0V or 1V state, the rate of change ( $dV/dt$ ) is zero, the displacement current is zero, and the via is electrically invisible. The exact moment a transition begins,  $dV/dt$  spikes, and the parasitic capacitor temporarily steals high-frequency energy from the forward-traveling wave.

## Energy Conservation in Via Parasitic Leakage

The via does not sort the signal into two isolated frequency buckets. The parasitic capacitance of a typical via is extremely small (often less than a picofarad). A small capacitor only achieves low impedance at extremely high frequencies. The vast majority of the signal successfully travels down the copper barrel. The parasitic capacitor steals only a tiny fraction of the highest-frequency components, rounding off the sharp corners of the square wave without stripping the signal to pure low frequencies.

The leaked energy does not vanish. Current must always flow in a completely closed loop. The high-frequency energy that jumps the clearance gap via displacement current lands directly in the surrounding ground plane, where it merges with the main return current and flows backward to the transmitter.

Tracing the energy accounting: the transmitter launches a wavefront carrying a total current. At the via, a small fraction of that current takes the parasitic shortcut directly into the ground plane. The remainder continues forward along the trace. The forward-moving wavefront creates its own return current loop through the ground plane. Both return currents travel backward to the transmitter, where their sum equals the original total. The conservation of charge is maintained perfectly; the via simply provided a premature exit for a fraction of the high-frequency energy.

The high-frequency return current follows the same path-of-least-inductance physics that governs all return current at high frequencies. The leaked current immediately forms a tight band flowing directly under the input trace leading back to the source. It merges with the main return current in the same narrow spatial channel, summing via superposition as the two flows occupy the same conductive region of the ground plane.

The net effect on the forward wavefront is a subtle loss of high-frequency edge content. The signal arriving at the receiver is missing the sharpest frequency components, producing a slightly rounded transition visible on an oscilloscope as a slower rise time.

## Excess Capacitance and Inductance on the TDR

A time domain reflectometer provides direct physical measurement of parasitic structures.

**Capacitive parasitic (impedance dip):** Impedance is the ratio of voltage to current (  $Z = V/I$  ). When the propagating wavefront encounters a region of excess metal (a wide pad, a via barrel), the excess geometric capacitance demands a sudden surge of current to charge the extra surface area. A surge in current for a given voltage means the instantaneous impedance drops. The TDR records this as a downward dip. The physical mechanism of the negative reflection follows from the pressure analogy: opening a new parallel path to ground is equivalent to puncturing a pressurized pipe. The local electrical pressure drops, and this drop propagates backward as a negative voltage ripple that the TDR instrument records.

**Inductive parasitic (impedance bump):** A narrow bottleneck in the trace forces all the current through a reduced cross-section, concentrating the magnetic field into a dense ring. The elevated local inductance generates a violent opposing back-EMF against the high-frequency wavefront. This opposition acts like a partial open circuit, establishing a local impedance above the 50-ohm baseline. A positive reflection propagates backward, and the TDR displays an upward bump at the corresponding location.

The distinction matters for diagnosis. Both parasitics degrade high-frequency edge content and produce a rounded transition at the receiver, but their TDR signatures are opposite: dip for capacitive, bump for inductive. The physical difference lies in the mechanism: the capacitive parasitic shunts energy out of the signal path through a parallel escape route, while the inductive parasitic blocks energy from passing through the signal path with a series barrier. Both extract the high-frequency content that constructs sharp edges, leaving the low-frequency foundation intact.

## Worked Examples

---

### DC vs. High-Frequency Return Path Geometry

Injecting a DC test current into a trace produces a return current that spreads broadly across the entire ground plane. The current takes a straight-line path from source to sink, ignoring the physical routing of the trace. The inductance term in the impedance equation (  $j\omega L$  ) vanishes at zero frequency, leaving only DC resistance, which is minimized by maximizing the cross-sectional area of the return path.

The same trace carrying a 10 GHz signal produces a dramatically different return path. The massive frequency multiplier  $\omega$  causes the inductive reactance to completely dominate the impedance equation. The return current ignores the vast expanse of available copper and crowds into a tight, narrow band directly underneath the signal trace, shadowing its physical routing with high precision. The skin effect further confines this current to the top surface of the ground plane.

An intermediate frequency of 100 MHz produces an intermediate distribution: the return current concentrates partially beneath the trace but retains some lateral spreading. The transition between the resistance-dominated and inductance-dominated regimes is gradual, governed continuously by the relative magnitudes of  $R$  and  $\omega L$  .

## Via Energy Split: A Quantitative Example

Consider a transmitter launching a wavefront carrying 10 milliamps of total current. The wavefront reaches a via with a parasitic capacitance of 0.3 pF.

At the via, the parasitic capacitor shunts approximately 1 milliamp of the highest-frequency current directly into the ground plane. The remaining 9 milliamps continue forward along the copper barrel to the destination trace.

The 1 milliamp of leaked current immediately turns around in the ground plane and flows backward toward the transmitter, forming a tight band under the input trace. The 9 milliamp forward wavefront creates its own return current loop through the ground plane. When the two return currents physically merge (at the via location, where the 9 milliamp band transitions from following the output trace to following the input trace), they sum to exactly 10 milliamps. The total return current arriving at the transmitter matches the original source current precisely.

The forward wavefront departing the via carries only 9 milliamps. The missing 1 milliamp of high-frequency energy visibly rounds off the leading edge of the voltage wave. On a TDR display, the parasitic via capacitor appears as a localized negative impedance dip at the via's physical location.

## Capacitor Behavior During a Voltage Transition

The parasitic via capacitor reacts exclusively to the transition rate, not the repetition rate.

A 1 MHz clock and a 1 GHz clock driven by the same transistor possess the same physical rise time (perhaps 20 picoseconds). The transistor always transitions from 0V to 1V at the same speed; only the time between transitions differs.

Sitting at a steady 0V or 1V means the voltage is completely static. The rate of change (  $dV/dt$  ) equals zero. The capacitor fully charges, stopping all current flow and acting as a total open circuit. The signal sails past the via undisturbed.

The exact moment the transistor initiates a 20-picosecond transition,  $dV/dt$  spikes to an astronomically high value. The capacitor equation (  $I = C \cdot dV/dt$  ) dictates that this massive rate of change produces a massive instantaneous current spike flowing into the parasitic via capacitor. The via acts as a low-impedance shunt exclusively during those 20 picoseconds. Reaching the flat rail voltage drops  $dV/dt$  back to zero, instantly turning the parasitic via back into a harmless open circuit.

## Edge Cases

---

### Reference Plane Discontinuities

The most damaging parasitic effect in high-speed PCB design occurs when the reference plane itself is interrupted. A slot cut through the ground plane for routing purposes, a split between power domains, or a poorly placed mounting hole all create a physical gap that the high-frequency return current cannot cross.

The return current must detour around the obstacle, drastically expanding the loop area. This loop ballooning simultaneously increases the loop inductance (producing an impedance spike visible on the TDR), creates a radiating magnetic loop antenna (generating electromagnetic interference), and degrades the mutual inductance between the trace and return path. The severity scales with frequency: at DC, the current spreads uniformly and the gap is irrelevant. At multi-gigahertz frequencies, the return current is locked beneath the trace, and even a small gap forces a massive detour.

## **The Separation of Repetition Rate and Transition Rate**

A conceptual trap exists around the relationship between clock frequency and parasitic interaction. A digital clock signal's repetition rate determines how often the transistor toggles (the fundamental clock frequency). The transition rate determines how fast the voltage actually changes during each toggle (the high-frequency Fourier content required to construct the edge).

The parasitic via capacitor reacts exclusively to the transition rate. A 1 MHz clock with a 20-picosecond rise time contains the same high-frequency edge content as a 10 GHz clock with the same rise time. Both create the same instantaneous displacement current burst through the via. The only difference is how frequently the burst occurs, not how severe each individual burst is.

This distinction matters for measurement and analysis. Two designs operating at different data rates but using the same transistor technology produce identical per-transition parasitic effects. The cumulative impact on signal quality depends on both the per-event severity and the event frequency.

## **AC Coupling Capacitor and Baseline Wander**

An intentionally placed series capacitor (such as an AC coupling capacitor in a high-speed serial link) exploits the same physics that makes parasitic capacitance problematic, but in a controlled manner.

The capacitor blocks DC by the same plate-charging mechanism that governs parasitic capacitance. A long run of identical bits (a sustained DC level) causes the plates to charge fully. The back-pressure from the crowded plates eventually equals the driving voltage, and current stops flowing entirely. The measured voltage on the receiver side drains to zero through the termination resistor.

The time required for this collapse depends on the RC time constant: a large capacitor with massive plates takes many identical bits to fill, while a small capacitor fills rapidly. The voltage decay follows a smooth exponential function, not a sudden cutoff. High-speed serial protocols mitigate this baseline wander by enforcing balanced encoding (equal numbers of ones and zeros), ensuring that each polarity reversal resets the capacitor charge before the plates can fill.



PART 2

# Signal Integrity Measurement

## Chapter 6: Time Domain Reflectometry

---

A time domain reflectometer is an electromagnetic wave instrument. It launches a fast voltage step onto a transmission line and records the echoes that return from impedance discontinuities along the path. The measurement is conceptually identical to sonar or radar: a known pulse is transmitted, the system waits for reflections, and the arrival time and amplitude of each returning echo reveal the location and severity of every structural anomaly in the channel.

The physical foundation is wave propagation. The incident step travels through the dielectric material between the trace and the ground plane at a velocity governed by the dielectric constant ( $v_p = c/\sqrt{\epsilon_r}$ ). When this wavefront encounters a change in characteristic impedance, a portion of the wave energy reflects backward while the remainder continues forward. The TDR instrument sitting at the source measures the total voltage at the connector over time, extracting the reflection coefficient from the difference between the measured voltage and the known incident step. The algebraic inversion of the reflection coefficient equation converts each measured voltage perturbation into a local impedance value, producing a continuous impedance profile of the entire channel from a single measurement event.

The result is a spatial map of the transmission line's impedance as a function of physical distance. Every via, connector, trace width change, reference plane gap, and termination discontinuity produces a distinct signature on the TDR display. This guide covers the complete instrument architecture, the physics of reflection and impedance translation, the mathematical derivation of spatial resolution limits, quantitative parasitic extraction from the TDR waveform, and the peeling algorithm that corrects for cumulative wavefront degradation.

### Core Concepts

---

#### The TDR Architecture: Source, Sampler, and Termination

A TDR consists of three functional elements sharing a single physical node at the instrument connector: an internal voltage source (step generator), an internal series resistor, and a high-bandwidth sampling oscilloscope.

The industry standard for the internal source resistance ( $R_s$ ) is exactly 50 ohms. Connecting the TDR to a PCB trace physically links this internal 50-ohm source to the transmission line's characteristic impedance ( $Z_0$ ). For a properly designed 50-ohm trace, the entire system forms a simple voltage divider at the connector.

The internal source generates a clean 1V step. This 1V potential divides across two equal resistances in series: the internal 50-ohm resistor and the 50-ohm active impedance of the trace. Equal resistances split the voltage evenly. Exactly 0.5V drops across the internal TDR resistor as heat, and the remaining 0.5V launches onto the trace as an electromagnetic wave.

The TDR screen displays the voltage measured precisely at the connector. It shows an instantaneous jump to 0.5V. The screen remains perfectly flat at 0.5V while the wavefront is actively traveling down a uniform 50-ohm trace. The required ratio of voltage to current (  $V/I = Z_0$  ) remains constant across the uniform copper geometry, producing zero reflections during propagation.

The internal 50-ohm resistor serves a second critical function beyond forming the voltage divider. When a reflection returns from a downstream discontinuity, the backward-traveling wave encounters this 50-ohm resistor at the source. The arriving wave perceives the resistor as a perfect physical continuation of the 50-ohm trace. Encountering zero change in impedance prevents any new reflections from occurring at the connector boundary. The electrical current of the returning wave flows directly into the physical resistor material, converting the electromagnetic energy into thermal heat. The entire reflection is absorbed, preventing it from bouncing forward again and contaminating the measurement of the rest of the trace.

## Velocity of Propagation and Distance Measurement

The electromagnetic wave propagates through the dielectric at a velocity determined by the dielectric constant:

$$v_p = \frac{c}{\sqrt{\epsilon_r}}$$

Standard FR4 has a dielectric constant of approximately 4.0, reducing the propagation velocity to roughly half the speed of light, or about six inches per nanosecond. Lower-loss laminates such as Megtron 6 have lower dielectric constants (around 3.6), permitting slightly faster propagation.

TDR operates on time of flight. The instrument launches the step, starts a timer, and waits for the echo. The wave must travel to the discontinuity and back, so the distance formula accounts for the round trip:

$$d = \frac{v_p \cdot \Delta t}{2}$$

The instrument does not know the dielectric constant of the board under test. The user interface provides a setting for velocity factor (  $V_f$  ) or dielectric constant (  $D_k$  ) that must be configured manually. Entering the wrong value has a specific consequence: the impedance readings on the vertical axis remain perfectly correct (the reflection coefficient depends only on the voltage ratio, which is measured directly), but the distance scale on the horizontal axis is warped. If the true dielectric constant is unknown, a practical calibration technique is available: measure a trace of physically known length, record the time delay, and calculate the exact  $D_k$  of the fabricated board from the measured flight time.

Varying dielectrics along the signal path introduce cumulative distance error. A wavefront leaving an FR4 trace (  $D_k = 4.0$  ) and passing through a coaxial connector (Teflon,  $D_k = 2.1$  ) accelerates through the connector region. If the TDR assumes a constant velocity for the entire measurement, the spatial X-axis will be distorted after the first material transition. Standard TDR instruments accept this minor error for physically short transitions like vias or connectors, where the time error is measured in single-digit picoseconds. Transitioning to a long cable or substrate with a substantially different dielectric constant, however, will noticeably skew all downstream distance measurements.

## The Reflection Coefficient

When the incident step encounters an impedance mismatch, the fraction of voltage that reflects backward is governed by the reflection coefficient:

$$\Gamma = \frac{Z_L - Z_0}{Z_L + Z_0}$$

$Z_0$  is the transmission line's characteristic impedance (the 50-ohm baseline), and  $Z_L$  is the impedance of the discontinuity.

Three canonical cases define the extremes of TDR behavior:

**Perfectly matched load (  $Z_L = 50\Omega$  ):**  $\Gamma = \frac{50-50}{50+50} = 0$  . No reflection occurs. The energy is entirely absorbed by the load. The TDR screen displays a flat 0.5V line extending indefinitely.

**Open circuit (  $Z_L = \infty$  ):**  $\Gamma = \frac{\infty-50}{\infty+50} = +1$  . 100% of the voltage reflects with a positive sign. The TDR screen, which has been displaying a steady 0.5V, abruptly jumps to 1.0V at the moment the full positive reflection returns. The returning wave encounters the internal 50-ohm source termination and is completely absorbed. The system achieves steady state at 1.0V, holding the full internal source potential across the entire trace length.

**Short circuit (  $Z_L = 0$  ):**  $\Gamma = \frac{0-50}{0+50} = -1$  . 100% of the voltage reflects with an inverted (negative) sign. The TDR screen drops from 0.5V to 0V.

Turning the TDR on without connecting any trace at all causes the connector itself to act as an open circuit. Zero current flows through the internal 50-ohm resistor, so zero voltage drops across it. The full 1V potential from the internal source appears immediately at the connector with no intermediate 0.5V stair, because the round-trip flight time is zero.

## From Measured Voltage to Impedance

The TDR oscilloscope is a voltmeter. It physically measures the total voltage on the line at its port:

$$V_{measured}(t) = V_{incident} + V_{reflected}(t)$$

The reflection coefficient at each point in time is isolated by subtracting the known incident voltage and normalizing:

$$\rho(t) = \frac{V_{measured}(t) - V_{incident}}{V_{incident}}$$

The TDR firmware then applies a standard algebraic inversion of the reflection coefficient formula to convert each  $\rho(t)$  value into the local impedance. Starting from the definition  $\rho = \frac{Z-Z_0}{Z+Z_0}$  , multiplying both sides by  $(Z + Z_0)$  , distributing, grouping terms, and solving for  $Z$  yields:

$$Z(t) = Z_0 \cdot \frac{1 + \rho(t)}{1 - \rho(t)}$$

This equation is the bridge between the raw voltage measurement and the impedance profile displayed on the screen.

**Numerical example:** The TDR launches a 1V step (  $V_{incident} = 1V$  , producing a 0.5V launched wave into a matched line). At  $t = 2 \text{ ns}$  , the scope measures 1.2V total at the connector. The incident wave was 1V (the full source amplitude as seen at the scope node), so  $\rho = \frac{1.2-1.0}{1.0} = 0.2$  . The local impedance is  $Z = 50 \cdot \frac{1.2}{0.8} = 75\Omega$  . The TDR displays a step from the 50-ohm baseline up to 75 ohms at the physical location corresponding to the 2 ns round-trip delay.

## The TDR Staircase: What the Screen Displays

Understanding the TDR display requires internalizing that the screen is not a camera looking sideways at the physical length of the trace. The TDR is a hyper-fast voltmeter permanently fixed at exactly one physical location: the connector. The horizontal X-axis represents time (or, equivalently, round-trip distance), and the vertical Y-axis represents either voltage or impedance.

Consider a 12-inch PCB trace with an open circuit at the far end:

1.  $t < 0$  : The screen displays a flat line at 0V.
2.  $t = 0$  : The internal source fires. The voltage at the connector immediately jumps. The screen draws a vertical line shooting from 0V to 0.5V.
3.  $0 < t < 4 \text{ ns}$  : The wavefront actively travels down the 12-inch trace and reflects back. The electromagnetic wave in FR4 requires approximately 2 ns to travel 12 inches; the round trip is 4 ns. During this entire window, the active transmission line presents a 50-ohm load at the connector, holding the voltage at 0.5V via the voltage divider. The screen draws a flat horizontal line at 0.5V for 4 nanoseconds. This plateau is the literal time-of-flight of the wave through the board.
4.  $t = 4 \text{ ns}$  : The reflection arrives back at the connector. The returning +0.5V wave superimposes on the existing 0.5V. The screen draws a second vertical step from 0.5V to 1.0V.
5.  $t > 4 \text{ ns}$  : The matched internal source absorbs the reflection. The system achieves steady state. The screen extends a flat line at 1.0V indefinitely.

The resulting shape is a two-step staircase. The horizontal width of the 0.5V plateau is the exact round-trip flight time. Engineers use this width to physically calculate the length of the trace.

## Architecture

---

### Capacitive Dips and Inductive Bumps

Every localized impedance discontinuity produces a characteristic signature on the TDR display. The polarity of the reflection directly identifies the physical nature of the parasitic structure.

**Capacitive parasitic (impedance dip):** Excess capacitance arises from physical structures that add metallic surface area or bring conductors closer together. Test point pads, surface-mount device pads, and via barrels passing through internal ground planes act as parallel-plate capacitors connecting the signal path to the ground plane. When the propagating wavefront encounters this excess geometric capacitance, the wide region demands a sudden surge of current to charge the extra metal area. A surge in current for a given voltage means the instantaneous impedance drops (  $Z = V/I$  ). The TDR records this momentary drop as a downward dip on the screen.

The physical mechanism of the negative reflection follows directly from the wavefront dynamics. The parasitic capacitor opens a new parallel path to the ground plane, allowing electrons to flow out of the signal path into ground more readily than the baseline geometry permits. This sudden opening of an additional escape route drops the local electrical pressure below the forward baseline. The localized pressure drop propagates backward as a negative voltage ripple, which the TDR instrument records as a negative dip.

The impedance of a capacitor is:

$$Z_C = \frac{1}{2\pi fC}$$

At DC or very low frequencies, this impedance is effectively infinite. The capacitor acts as an open circuit to ground, and low-frequency signal energy ignores the via entirely. At the high frequencies composing the sharp leading edge of the step function, the impedance drops dramatically. The via looks like a partial short circuit to ground, and the high-frequency wave energy reflects violently.

**Inductive parasitic (impedance bump):** Excess inductance arises from physical structures that restrict the current path or force the return current to take a wider detour. Neck-downs, wirebonds, sharp corners, and gaps in the reference plane force the electromagnetic loop to expand. The elevated local inductance generates a violent opposing back-EMF against the high-frequency wavefront (  $V = L \cdot dI/dt$  ). This opposition establishes a local impedance above the 50-ohm baseline. A positive reflection propagates backward, and the TDR displays an upward bump.

The impedance of an inductor is:

$$Z_L = 2\pi fL$$

At low frequencies, this impedance is negligible, and the restriction is invisible. At high frequencies, the impedance climbs, producing a strong positive reflection.

**Diagnostic power:** Both parasitics degrade high-frequency edge content and produce a rounded transition at the receiver. The distinction matters for diagnosis. On the TDR display, the capacitive parasitic appears as a negative dip (local impedance below  $Z_0$  ), while the inductive parasitic appears as a positive bump (local impedance above  $Z_0$  ). The physical difference lies in the mechanism: the capacitive parasitic shunts energy out of the signal path through a parallel escape route, while the inductive parasitic blocks energy from passing through the signal path with a series barrier.

## Returning to the Baseline After a Defect

A localized parasitic structure (a narrow trace measuring only a few millimeters, a via, a connector pad) has finite physical length. The forward-traveling wavefront hits the start of the anomalous section and immediately begins generating a reflection backward. The wavefront requires only a few picoseconds to traverse those millimeters. Exiting the defect returns the wavefront to the properly designed 50-ohm trace geometry, which generates zero reflections.

The TDR screen plots reflected energy arriving at the instrument against time. Receiving a short burst of reflected energy followed immediately by zero new reflected energy causes the screen to draw a bump (or dip) and then return to the flat 0.5V baseline. The time duration of the feature on the screen directly corresponds to the physical width of the discontinuity:

$$\text{Length} = v_p \cdot \frac{\Delta t}{2}$$

## The Impedance Cancellation Principle

If a discontinuity introduces additional inductance and additional capacitance in the exact proportions required to maintain  $Z_0 = \sqrt{L/C}$ , the impedance remains matched. The TDR trace displays a flat line with no reflection, even though the physical geometry changed. Engineers exploit this phenomenon to optimize signal links. Deliberately routing a narrow, highly inductive trace segment adjacent to a wide, highly capacitive connector pad can maintain a uniform 50-ohm impedance profile through a region where either element alone would produce a severe mismatch.

## Worked Examples

---

### Spatial Resolution: Why Rise Time Limits What the TDR Can See

The voltage step generated by the TDR is not instantaneous. It has a finite rise time ( $t_r$ ), typically defined as the 10% to 90% transition time. This rise time sets a fundamental limit on the instrument's ability to distinguish between two impedance discontinuities that are physically close together.

**The physical length of the rising edge:** While the step function propagates down the trace, the transition from 0V to full voltage is not concentrated at a mathematical point. The rising edge is physically spread across a length of trace:

$$l_{\text{edge}} = v_p \cdot t_r$$

A 30-picosecond rise time at 6 inches per nanosecond produces a rising edge physically spread over 0.18 inches of the PCB.

**Deriving the resolution limit:** Consider two discontinuities, Point A and Point B, separated by physical distance  $d$ .

At  $t = 0$  , the leading tip of the rising edge hits Point A. Point A immediately begins generating a reflection that travels backward toward the TDR. This reflection requires  $t_r$  seconds to finish, because the incident edge itself takes that long to reach full amplitude.

The forward wave continues past Point A and travels distance  $d$  to reach Point B. This takes time  $t_{forward} = d/v_p$  . The edge hits Point B and generates a second reflection, which travels distance  $d$  back to Point A. The total delay between the start of the first reflection and the arrival of the second reflection at Point A is:

$$\Delta t = \frac{2d}{v_p}$$

For the TDR to display two distinctly separate features, the first reflection must completely finish before the second reflection begins to arrive. The first reflection takes exactly  $t_r$  seconds to finish rising. The condition for resolving the two features is:

$$\frac{2d}{v_p} > t_r$$

Solving for the minimum resolvable distance:

$$d > \frac{v_p \cdot t_r}{2} = \frac{l_{edge}}{2}$$

**Two discontinuities are resolvable on a TDR screen only if their physical separation exceeds half the physical length of the rising edge.** If they are closer, the second echo arrives before the first is finished, and they superimpose into a single, unresolved lump.

**Practical numbers:** A Keysight TDR with a 15 ps rise time operating on FR4 (  $v_p \approx 6 \text{ in/ns}$  ) produces  $l_{edge} = 6 \times 0.015 = 0.09 \text{ inches}$  . The minimum resolvable separation is 0.045 inches, roughly 1.1 mm. A 30 ps system resolves only down to 2.3 mm.

## Quantitative Parasitic Extraction: The Area Under the Curve

The TDR waveform contains enough information to extract the actual physical capacitance or inductance of a localized discontinuity. The key is the mathematical integral of the reflection coefficient over time.

**The inductor formula:** The fundamental equation for an inductor is  $V = L \cdot dI/dt$  . Integrating both sides over the entire time duration of the defect converts the instantaneous rate of change into total accumulated values. The integral of voltage over time represents the total magnetic flux change. The integral of  $dI/dt$  simply equals the total current change  $\Delta I$  .

Applying boundary conditions from transmission line theory (the maximum current through the defect equals  $V_{inc}/Z_0$  , and the reflected voltage equals half the total voltage generated across a small series element under the small-reflection approximation) yields, after substituting the reflection coefficient  $\rho(t) = V_{reflected}/V_{incident}$  :

$$L = 2Z_0 \int \rho(t) dt$$

The total inductance of the defect equals twice the baseline impedance multiplied by the area under the positive impedance bump on the  $\rho(t)$  curve.

**The capacitor formula:** The fundamental equation for a capacitor is  $I = C \cdot dV/dt$ . The same integration process, applied to a parallel shunt element with the appropriate sign convention for the negative reflection, yields:

$$C = \frac{2}{Z_0} \int -\rho(t) dt$$

The total capacitance equals twice the reciprocal of the baseline impedance multiplied by the area enclosed within the negative impedance dip.

**Why the area stays constant regardless of rise time:** The capacitance of a via is a fixed physical property of the board. It is determined by the drilled hole geometry and the distance to the ground planes. The integral  $\int -\rho(t) dt$  must therefore be a constant for a given structure. What changes is the shape of the curve. A fast incident edge produces a short interaction time ( $dt$  is small), which forces the reflection amplitude ( $\rho$ ) to spike high, producing a deep, narrow dip. A degraded, slow edge produces a long interaction time, which forces the amplitude down to maintain the same total area, producing a wide, shallow dip. The underlying physical capacitance has not changed, but the instrument's ability to measure the true peak impedance has been destroyed by the degraded rise time.

## Distinguishing $Z_0$ from Reactive Components

The TDR identifies the characteristic impedance  $Z_0$  by reading the longest flat baseline region on the screen. On this flat line, the distributed inductance and capacitance are perfectly balanced in the ratio  $Z_0 = \sqrt{L/C}$ , producing zero net reflection. The TDR cannot separate the baseline inductance and capacitance on that flat line because their effects cancel.

Isolating the reactive components of individual discontinuities (the excess  $L$  or  $C$  above the balanced baseline) requires either the area-integral technique described above (for electrically small defects) or a transition to the frequency domain. A Vector Network Analyzer (VNA) measures S-parameters as a function of frequency, providing both magnitude and phase of the complex reflection coefficient. The Smith Chart plots this complex data: the center horizontal line represents pure resistance, points in the upper half indicate excess inductance, and points in the lower half indicate excess capacitance. Mathematically, a time-domain TDR step response can be converted into a frequency-domain S-parameter response by applying a Fast Fourier Transform to the measured step waveform.

## Edge Cases

---

### Rise Time Degradation and the Need for the Rising Edge

A DC voltage has no concept of time. Applying a flat 1V DC level charges the entire trace to 1V, and only the total resistance can be measured. To measure distance, the instrument needs a time marker. The rising edge is the physical event that starts the stopwatch, acting as a localized wavefront marker traveling down the line.

The sharpness of that edge is equally critical. A fast rising edge contains extremely high-frequency content (Fourier analysis of a 10-picosecond transition shows harmonics extending past 35 GHz). Impedance mismatches caused by tiny parasitic capacitances and inductances only react to high frequencies. If the edge is slow and contains only low-frequency content, it rolls right over these structures without generating a detectable reflection.

The finite bandwidth of the TDR instrument itself determines the steepest possible edge. A high-end TDR with 20 GHz bandwidth can generate harmonics only up to that frequency, limiting the rise time to approximately  $t_r \approx 0.35/BW \approx 17.5 \text{ ps}$ . A perfectly vertical (zero rise time) edge would require infinite bandwidth, which no physical instrument can produce.

### Masking: Upstream Discontinuities Blind the Instrument

Every time the forward wave hits a mismatch, some energy reflects backward. The forward-traveling incident wave shrinks with each reflection. If the wave hits a massive mismatch (a near-open or near-short), almost all the energy reflects, and the forward wave drops to near zero. The TDR becomes blind to everything downstream.

This is the masking effect. The practical rule for signal integrity debugging is absolute: fix the first major mismatch in a trace before trusting what the TDR reports about the rest of the trace.

The masking effect also distorts the quantitative impedance readings. The TDR firmware calculates downstream impedance by assuming the incident wave retains its full original amplitude. If upstream reflections have depleted the wave to 80% of its original value, the firmware applies the full-amplitude calibration to a reduced signal, reporting incorrect impedance values for all downstream structures.

### Spreading: Frequency-Dependent Attenuation Smears Downstream Features

The first parasitic inductor strips high-frequency content from the forward-traveling wave. The surviving wavefront continuing past the inductor possesses a much slower, rounded rising edge. This degraded wavefront becomes the new interrogating signal for the remainder of the trace.

Striking a second parasitic inductor with a slow, rounded edge produces a drastically different response. The inductor generates back-EMF proportional to the rate of current change ( $V = L \cdot dI/dt$ ). A slow, rounded edge has a small rate of change, producing only a weak opposing voltage. The TDR records this weak response as a small bump. Two physically identical inductors will look different on the screen: the first appears as a sharp, deep feature, while the second, hit by the degraded wavefront, appears as a wider, shallower, blurrier feature.

The spatial resolution literally degrades the further down the trace the TDR looks. The channel attenuation mechanisms described in the skin effect and dielectric loss guide (conductor loss scaling as  $\sqrt{f}$ , dielectric absorption scaling linearly with frequency) compound the problem. Even on a perfectly uniform trace with no discrete discontinuities, these distributed losses continuously filter the highest-frequency harmonics from the propagating step. A pristine 20-picosecond edge at the source can degrade to a sluggish 100-picosecond edge by the time it reaches the far end of a long board.

## The Peeling Algorithm

The peeling algorithm (also called impedance profile deconvolution, implemented in Keysight Physical Layer Test System software) is a purely mathematical correction that runs entirely in software. It does not alter the physical voltage on the board.

The algorithm analyzes the first reflection, calculates exactly how much energy was lost from the forward wave, and dynamically recalibrates the mathematics for the second reflection. It then strips the effect of the first discontinuity from the waveform and repeats the process for each subsequent feature, peeling away the cumulative errors layer by layer.

The recalibration accounts for three distinct error sources at each layer: the amplitude depletion of the forward wave (masking), the rise time degradation of the forward edge (spreading), and the secondary reflections generated when energy bouncing between two discontinuities creates phantom features on the screen. The result is a de-embedded impedance profile that represents the true local impedance at each point along the trace, as if each discontinuity had been measured in isolation with a pristine incident edge.

## Chapter 7: Smith Charts

---

The reflection coefficient is the central quantity of transmission line analysis. Time domain reflectometry measures it as a real-valued voltage ratio at each point in time, producing an impedance profile along the physical length of a channel. The same reflection coefficient, however, is inherently a complex number: it carries both a magnitude (how much energy reflects) and a phase angle (the timing relationship between the incident and reflected waves). Capturing both dimensions simultaneously requires a representation in the complex plane, where the horizontal axis carries the real component and the vertical axis carries the imaginary component.

The fundamental challenge is one of scale. The complex impedance that produces a given reflection can range from zero (a short circuit) to infinity (an open circuit). Plotting impedance directly on a Cartesian grid is physically impossible when one axis must extend to infinity. The Smith Chart solves this problem through conformal mapping: a rigorous mathematical transformation that compresses the entire infinite impedance plane into a single finite circle of unit radius. Every passive impedance that could ever exist on a transmission line maps to a unique point inside this circle.

The internal structure of the chart consists of two families of curves overlaid on the reflection coefficient plane. Constant resistance circles and constant reactance arcs create a warped coordinate grid that allows an engineer to read impedance values directly from any plotted point. Vector network analyzers display measured S-parameter data on this chart, and the trajectory of a trace across frequency immediately reveals whether a device under test is resistive, capacitive, inductive, well-matched, or drifting toward a dangerous mismatch.

This guide develops the Smith Chart from first principles: the complex nature of the reflection coefficient, the normalization convention that makes the chart universal, the full algebraic derivation of both circle families via completing the square, the geometric interpretation of every major landmark on the chart, and the extension beyond the unit circle for active devices.

### Core Concepts

---

#### Normalization and the Universal Chart

The reflection coefficient relating load impedance to characteristic impedance is:

$$\Gamma = \frac{Z_L - Z_0}{Z_L + Z_0}$$

Absolute impedance values prevent the chart from being universal. A 100-ohm load produces a different reflection coefficient depending on whether the system impedance is 50 ohms or 75 ohms. Dividing the load impedance by the characteristic impedance produces a dimensionless normalized impedance:

$$z = \frac{Z_L}{Z_0} = r + jx$$

The variable  $r$  represents normalized resistance (always non-negative for passive circuits) and  $x$  represents normalized reactance (positive for inductors, negative for capacitors). Substituting this normalized impedance into the reflection coefficient equation eliminates the system impedance entirely:

$$\Gamma = \frac{z - 1}{z + 1}$$

A normalized impedance of  $z = 1$  (a load matched to the system) always produces  $\Gamma = 0$ , regardless of whether the system is 50 ohms, 75 ohms, or any other value. This normalization is the mechanism that makes a single printed chart applicable to every transmission line system.

### The Reflection Coefficient as a Complex Number

The normalized impedance  $z = r + jx$  is complex. Dividing one complex number by another always produces a third complex number. The reflection coefficient therefore possesses both a real component and an imaginary component, which are assigned the variables  $u$  and  $v$ :

$$\Gamma = u + jv$$

This assignment is not arbitrary. It reflects the physical reality that a reflected wave carries both a magnitude and a phase shift relative to the incident wave, and representing both quantities simultaneously requires two independent coordinates.

The algebraic proof begins with the explicit substitution  $z = r + jx$  into  $\Gamma = (z - 1)/(z + 1)$ :

$$\Gamma = \frac{(r - 1) + jx}{(r + 1) + jx}$$

Removing the imaginary number from the denominator requires multiplying both the numerator and denominator by the complex conjugate of the denominator,  $(r + 1) - jx$ . The denominator becomes purely real:

$$[(r + 1) + jx] \cdot [(r + 1) - jx] = (r + 1)^2 + x^2$$

Expanding the numerator and grouping real and imaginary terms yields:

$$\Gamma = \frac{r^2 - 1 + x^2}{(r + 1)^2 + x^2} + j \frac{2x}{(r + 1)^2 + x^2}$$

The first fraction contains no  $j$  operator and is a purely real number. This is  $u$ . The second fraction multiplies  $j$  and is purely imaginary. This is  $v$ :

$$u = \frac{r^2 - 1 + x^2}{(r + 1)^2 + x^2}$$

$$v = \frac{2x}{(r+1)^2 + x^2}$$

These two quantities serve as the horizontal and vertical plotting coordinates on the Smith Chart. Every complex impedance  $z = r + jx$  maps to exactly one point  $(u, v)$  inside the unit circle.

## The Conformal Mapping: Infinite Plane to Finite Circle

The Smith Chart is the visual result of a conformal mapping. The word "conformal" means angle-preserving: orthogonal intersections in the impedance plane remain orthogonal after transformation. The rectangular grid of constant- $r$  vertical lines and constant- $x$  horizontal lines in the impedance plane warps into families of circles and arcs on the  $\Gamma$  plane, but the perpendicular crossing angles are maintained.

Passive circuits guarantee that the magnitude of the reflection coefficient can never exceed 1. No passive load can reflect more energy than it receives. This physical constraint means all possible positive resistance and reactance values map to points inside a circle of radius 1 centered at the origin of the  $(u, v)$  plane. The conformal mapping compresses an infinite Cartesian coordinate system into this finite circular domain.

## Worked Examples

---

### Deriving the Constant Resistance Circles

The mathematical derivation begins with the fundamental relationship between normalized impedance and the reflection coefficient. Starting from the inverse mapping:

$$z = \frac{1 + \Gamma}{1 - \Gamma}$$

Substituting  $z = r + jx$  and  $\Gamma = u + jv$  :

$$r + jx = \frac{(1 + u) + jv}{(1 - u) - jv}$$

Multiplying numerator and denominator by the complex conjugate of the denominator,  $(1 - u) + jv$ , produces:

$$r + jx = \frac{1 - u^2 - v^2}{(1 - u)^2 + v^2} + j \frac{2v}{(1 - u)^2 + v^2}$$

Equating real parts isolates the resistance equation:

$$r = \frac{1 - u^2 - v^2}{(1 - u)^2 + v^2}$$

Multiplying both sides by the denominator and expanding:

$$r - 2ru + ru^2 + rv^2 = 1 - u^2 - v^2$$

Grouping the  $u$  and  $v$  terms on the left side:

$$u^2(r+1) - 2ru + v^2(r+1) = 1 - r$$

Dividing the entire equation by  $(r+1)$  :

$$u^2 - \frac{2r}{r+1}u + v^2 = \frac{1-r}{r+1}$$

Completing the square for  $u$  requires adding  $\left(\frac{r}{r+1}\right)^2$  to both sides:

$$\left(u - \frac{r}{r+1}\right)^2 + v^2 = \frac{1-r}{r+1} + \frac{r^2}{(r+1)^2}$$

Simplifying the right side by combining fractions over the common denominator  $(r+1)^2$  :

$$\frac{(1-r)(r+1) + r^2}{(r+1)^2} = \frac{1-r^2+r^2}{(r+1)^2} = \frac{1}{(r+1)^2}$$

The final result is the equation of a circle:

$$\left(u - \frac{r}{r+1}\right)^2 + v^2 = \left(\frac{1}{r+1}\right)^2$$

This circle has its center at coordinates  $\left(\frac{r}{r+1}, 0\right)$  on the horizontal axis and radius  $\frac{1}{r+1}$  .

The geometric interpretation is immediate. All constant resistance circles are centered on the horizontal axis ( $v = 0$  ). As  $r$  increases, the center shifts rightward toward  $u = 1$  and the radius shrinks toward zero. At  $r = 0$  , the circle has center  $(0, 0)$  and radius 1, tracing the outer boundary of the entire chart. At  $r = 1$  , the circle has center  $(0.5, 0)$  and radius 0.5, passing exactly through the chart center. As  $r \rightarrow \infty$  , both the center and the circle converge to the single point  $(1, 0)$  on the far right edge.

## Deriving the Constant Reactance Arcs

Equating imaginary parts from the same expanded equation isolates the reactance:

$$x = \frac{2v}{(1-u)^2 + v^2}$$

Multiplying by the denominator and distributing:

$$x - 2xu + xu^2 + xv^2 = 2v$$

Dividing the entire equation by  $x$  and rearranging:

$$u^2 - 2u + v^2 - \frac{2v}{x} = -1$$

Completing the square for  $u$  (adding 1) and for  $v$  (adding  $\frac{1}{x^2}$ ):

$$(u - 1)^2 + \left(v - \frac{1}{x}\right)^2 = \left(\frac{1}{x}\right)^2$$

This circle has its center at coordinates  $\left(1, \frac{1}{x}\right)$  and radius  $\frac{1}{|x|}$ .

All constant reactance circles are centered on the vertical line  $u = 1$  (the right edge of the chart). They appear as arcs rather than full circles because only the portion falling inside the unit circle corresponds to physical passive impedances. As  $|x|$  increases, the center approaches the horizontal axis and the radius shrinks, producing tighter arcs near the right edge. As  $|x| \rightarrow 0$ , the radius approaches infinity, and the arc flattens into the horizontal line  $v = 0$  (the pure resistance axis). Positive reactance arcs ( $x > 0$ ) curve through the upper half of the chart, representing inductive behavior. Negative reactance arcs ( $x < 0$ ) curve through the lower half, representing capacitive behavior.

## Architecture

---

### Reading the Chart: Landmarks and Navigation

The overlaid grid of constant resistance circles and constant reactance arcs defines a complete coordinate system inside the unit circle. Every point on the chart corresponds to a unique normalized impedance, and every normalized impedance maps to a unique point.

**The center point** is the single most important location. It lies at the intersection of the  $r = 1$  resistance circle and the  $x = 0$  reactance line, at coordinates  $(u, v) = (0, 0)$  on the reflection coefficient plane. This point represents  $\Gamma = 0$ : a perfect impedance match with zero reflection. The goal of impedance matching network design is to transform the load impedance toward this center point.

**The far left edge** of the horizontal axis, at  $(u, v) = (-1, 0)$ , represents a short circuit. The resistance is zero and the reactance is zero. The reflection coefficient magnitude is 1, and the phase is 180 degrees. The signal encounters a direct path to ground, reflecting all energy with an inverted voltage.

**The far right edge** of the horizontal axis, at  $(u, v) = (1, 0)$ , represents an open circuit. The resistance is infinite. The reflection coefficient magnitude is 1, and the phase is zero. All constant resistance circles and all constant reactance arcs converge at this single point.

**The horizontal axis** ( $v = 0$ , the line cutting straight through the center) represents zero reactance. Every point on this line is a purely resistive impedance. The printed resistance values along this axis begin at zero on the far left, pass through 1 at the center, and reach infinity at the far right.

**The upper half** of the chart (  $v > 0$  ) contains all impedances with positive reactance (inductive behavior). **The lower half** (  $v < 0$  ) contains all impedances with negative reactance (capacitive behavior). The outer circumference of the chart corresponds to the boundary where  $r = 0$  : purely reactive components with no resistive energy dissipation.

**Reading a specific impedance:** Locating the intersection of the  $r = 0.5$  resistance circle with the  $x = +1.0$  reactance arc defines the normalized impedance  $z = 0.5 + j1.0$  . In a 50-ohm system, this denormalizes to  $Z_L = 25 + j50$  ohms. The physical distance from the chart center to this intersection point is the magnitude of the reflection coefficient. The angle of this point relative to the positive horizontal axis is the phase of the reflection.

## Coordinate Systems and Measurement

The Smith Chart exists simultaneously in two coordinate systems that serve complementary purposes.

The **Cartesian coordinates** ( $u, v$ ) span from  $-1$  to  $+1$  on both axes and define the physical drawing space. The overlaid resistance circles and reactance arcs constitute the impedance grid projected onto this drawing space.

The **polar coordinates** (magnitude and angle) describe the measurement. The radial distance from the center to any plotted point is the magnitude  $|\Gamma|$  of the reflection coefficient. The angle from the positive horizontal axis to that point is the phase of the reflected wave. Measurement equipment captures magnitude and phase natively, and the chart provides the visual translation from these polar quantities back to Cartesian resistance and reactance.

## VNA Trace Interpretation

Plotting S-parameter data from a vector network analyzer onto the Smith Chart provides immediate visual insight into channel behavior across frequency.

A well-matched component appears as a tight cluster of points near the chart center. The closer the trace stays to the center, the smaller the reflection coefficient across the measured frequency range. A trace circling tightly around the center as frequency sweeps indicates a broadband impedance match with minor frequency-dependent variation.

A trace residing entirely in the upper hemisphere reveals a predominantly inductive load. A trace confined to the lower hemisphere indicates a predominantly capacitive load. The trajectory of the trace across frequency shows how the impedance evolves: a spiral moving inward toward the center as frequency increases indicates that the match improves at higher frequencies, while a spiral moving outward indicates deteriorating performance.

The Smith Chart transforms abstract complex algebra into a direct visual map that connects the time-domain reflection coefficient measured by TDR instruments to the frequency-domain S-parameter data captured by network analyzers. Both instruments measure the same physical quantity. TDR presents it as a real-valued impedance profile along the spatial length of the trace. The VNA presents it as a complex-valued trajectory sweeping across frequency on the Smith Chart.

## Edge Cases

---

### The Passive Constraint and the Unit Circle Boundary

The unit circle boundary on the Smith Chart is not an arbitrary choice of scale. It is a direct consequence of energy conservation. A passive network contains no internal energy source. It can absorb energy, store and release energy, or reflect energy, but it cannot create energy. The maximum possible reflection occurs when no energy is absorbed:  $|\Gamma| = 1$ . This physical limit confines all passive impedances to points at or inside the unit circle.

Points exactly on the boundary (  $|\Gamma| = 1$  ) represent purely reactive loads with zero resistance. An ideal lossless inductor or an ideal lossless capacitor reflects all incident energy with zero absorption. The  $r = 0$  circle traces this boundary exactly.

Points inside the boundary (  $|\Gamma| < 1$  ) represent loads with nonzero resistance. Some energy is dissipated as heat in the resistive component, and less than 100% reflects.

### The Extended Smith Chart

Active circuits containing transistor amplifiers or negative-resistance oscillators can generate reflection coefficients with magnitudes exceeding 1. A device that adds energy to the reflected wave produces  $|\Gamma| > 1$ , and its impedance maps to a point outside the unit circle.

The polar coordinate system underlying the Smith Chart is not bounded by the unit circle. Polar coordinates define the entire infinite plane using a radial distance from the origin and an angular rotation from a reference axis. The radial distance can extend from zero to infinity. The extended Smith Chart expands the printed region beyond the standard boundary to accommodate these active-device measurements, preserving the same families of resistance circles and reactance arcs extrapolated into the region where  $|\Gamma| > 1$ .

### Conformal Mapping Preserves Angles but Distorts Distances

The mapping from the impedance plane to the reflection coefficient plane is conformal (angle-preserving) but not isometric (distance-preserving). Equal increments of resistance in the impedance plane do not correspond to equal distances on the chart. The region near the center (around  $r = 1$ ) is expanded relative to the extremes. Impedance values near the matched condition occupy a large visual area, providing high resolution exactly where engineers need it most. Impedance values near short or open circuits are compressed into narrow regions near the chart boundary, reflecting the physical reality that extreme mismatches are easy to detect but difficult to differentiate with precision.

This nonuniform spatial scaling is an inherent feature of the conformal mapping, not a limitation. It concentrates visual resolution on the impedance range where matching network adjustments are most sensitive.

## Chapter 8: S-Parameters and Vector Network Analysis

---

Time domain reflectometry reveals what a transmission line looks like at a single moment: a fast voltage step propagates down the channel, and the instrument records each echo as a spatial impedance profile. This approach excels at localizing physical structures along the trace, but it cannot separate the frequency-dependent behavior of those structures. A via that appears as a simple capacitive dip on a TDR waveform actually presents a complex impedance that varies continuously with frequency. Characterizing this frequency-dependent behavior requires an instrument that operates natively in the frequency domain.

A Vector Network Analyzer (VNA) replaces the broadband voltage step with a precisely controlled, narrowband sine wave. The instrument sweeps this sine wave across a programmable frequency range, and at each frequency step, it measures the magnitude and phase of the waves that reflect back and the waves that transmit through the device under test. These measurements are organized into a matrix of complex ratios called scattering parameters (S-parameters), which completely describe the linear behavior of the network at every measured frequency.

S-parameters are the lingua franca of high-speed channel characterization. Compliance specifications for PCIe, USB, HDMI, and Ethernet all define pass/fail criteria in terms of S-parameter masks. Return loss limits, insertion loss budgets, crosstalk isolation requirements, and impedance matching tolerances are all expressed as S-parameter thresholds plotted against frequency. This guide covers the physical architecture of the VNA, the mathematical definition of scattering parameters, the decibel system used to represent them, multi-port measurement for differential and crosstalk characterization, and the IF bandwidth mechanism that controls dynamic range.

### Core Concepts

---

#### What the VNA Measures: Incident and Scattered Waves

A VNA injects a known sine wave into one port of a device under test and simultaneously measures two quantities: the fraction of energy that reflects back out of the same port, and the fraction that transmits through to each other port. The instrument captures both the magnitude and the phase of each measured wave relative to its own internal reference oscillator.

Scattering parameters organize these measurements into a systematic matrix indexed by port number. The naming convention  $S_{ij}$  means "the ratio of the wave leaving port  $i$  to the wave entering port  $j$ ." In a standard two-port measurement where Port 1 is the transmitter end and Port 2 is the receiver end:

- $S_{11}$  is the signal injected into Port 1 that reflects back to Port 1. This quantity is called return loss. It measures the impedance match at the input.
- $S_{21}$  is the signal injected into Port 1 that successfully transmits through to Port 2. This is insertion loss. It measures how much energy the channel delivers to the receiver.
- $S_{12}$  is the signal injected into Port 2 that transmits backward to Port 1. This is the reverse transmission coefficient.
- $S_{22}$  is the signal injected into Port 2 that reflects back to Port 2. This is the return loss at the output port.

Measuring from both directions is necessary to fully characterize the device. A simple passive copper trace is symmetric:  $S_{21} = S_{12}$  and  $S_{11} = S_{22}$ . This symmetry is a direct consequence of the reciprocity theorem, which guarantees that the electromagnetic coupling between two ports of a passive linear network is identical regardless of direction. Active components like amplifiers or RF isolators break this symmetry, passing energy preferentially in one direction while attenuating the reverse path.

## S-Parameters as Complex Ratios

The VNA measures the incident wave (  $a$  ) entering a port and the scattered wave (  $b$  ) leaving a port. Each S-parameter is the ratio of a scattered wave to an incident wave:

$$S_{21} = \frac{b_2}{a_1}$$

The incident and scattered waves are both measured in the same voltage units, so S-parameters are fundamentally dimensionless. They carry no inherent unit of measurement. Each S-parameter is a complex number containing both magnitude and phase, represented as  $r\angle\theta$ . The magnitude quantifies how much energy reflects or transmits. The phase quantifies the time delay that the signal accumulates in transit, expressed as degrees of rotation of the specific sine wave period at each measurement frequency.

## The Decibel System

S-parameter magnitudes span an enormous dynamic range. Insertion loss through a well-designed trace might be  $-0.5$  dB (89% of the power reaching the receiver), while the crosstalk leaking between two adjacent lanes might be  $-60$  dB (one millionth of the injected power). Comparing these values on a linear scale is impractical. The decibel compresses this range into manageable numbers using a logarithmic transformation.

The decibel is fundamentally a ratio between two quantities, not an absolute unit of physical measurement. It expresses how much larger or smaller a measured value is compared to a reference value. The mathematical definition depends on whether the underlying quantity is power or voltage.

For **power ratios**, the decibel is defined as:

$$\text{dB} = 10 \log_{10} \left( \frac{P_{out}}{P_{in}} \right)$$

For **voltage ratios**, the exponent rule of logarithms converts the squared relationship (  $P \propto V^2$  ) into a factor of 20:

$$\text{dB} = 20 \log_{10} \left( \frac{V_{out}}{V_{in}} \right)$$

S-parameters are voltage ratios, so the 20-multiplier form applies. An  $S_{21}$  of 0 dB indicates perfect unity: 100% of the signal power injected at Port 1 reached Port 2. The ratio is exactly 1, and  $\log_{10}(1) = 0$ . An  $S_{21}$  of  $-3$  dB indicates a 50% reduction in power (the voltage ratio is  $1/\sqrt{2} \approx 0.707$ ). An  $S_{11}$  of  $-20$  dB indicates that only 1% of the injected power reflects back, denoting an excellent impedance match.

Several important reference benchmarks follow directly from the logarithm:

| Magnitude (linear) | dB Value | Physical Meaning                       |
|--------------------|----------|----------------------------------------|
| 1.0                | 0 dB     | Perfect unity (zero loss, zero gain)   |
| 0.707              | $-3$ dB  | Half-power point                       |
| 0.1                | $-20$ dB | 1% power reflection / 99% transmission |
| 0.01               | $-40$ dB | 0.01% power                            |
| 0.001              | $-60$ dB | One-millionth of power                 |

### Absolute Power: dBm

Standard dB expresses a relative ratio between two values that may both be unknown. The unit dBm anchors this ratio to a fixed physical reference: 0 dBm equals exactly 1 milliwatt. The underlying mathematics remain identical. The ratio of 1 milliwatt divided by the 1-milliwatt reference evaluates to 1, and  $10\log_{10}(1) = 0$ .

A critical distinction separates these two units. A relative dB gain or loss can be applied to an absolute dBm power level (a 10 dBm signal passing through a  $-3$  dB attenuator produces a 7 dBm signal at the output), but converting between dB and dBm without knowing at least one absolute reference is mathematically undefined. The two units measure fundamentally different concepts: dB measures a ratio, and dBm measures an absolute power.

## Architecture

### The VNA Swept-Sine Architecture

The physical architecture of a VNA differs fundamentally from a TDR. A TDR launches a broadband voltage step containing energy across the entire spectrum simultaneously, then sorts out the frequency content by mathematics (FFT) after the fact. A VNA operates in the opposite direction. Its internal source generates a single, pure, continuous sine wave at one specific frequency. The instrument injects this sine wave into the device under test, measures the reflected and transmitted responses at that one frequency, then increments the source to the next frequency and repeats. The VNA builds the complete frequency-domain picture one data point at a time across the programmed sweep range.

This swept-sine approach carries a fundamental advantage in measurement precision. At each frequency step, the VNA knows exactly what frequency it injected and can tune its receiver to listen exclusively at that frequency. All energy arriving at any other frequency is rejected as noise. This narrow-band detection is the mechanism that gives VNAs their extraordinary dynamic range, often exceeding 100 dB.

The two instruments measure the same physical quantity. A time-domain TDR step response can be mathematically converted into a frequency-domain S-parameter response by applying a Fast Fourier Transform (FFT) to the measured step waveform. The inverse operation is equally valid: an IFFT applied to S-parameter data reconstructs the equivalent time-domain impulse response. The TDR excels at spatial localization of impedance structures. The VNA excels at isolating frequency-dependent behavior and separating reactive components that the TDR cannot distinguish on a flat baseline.

## Phase Measurement and Display

Each S-parameter carries a phase component alongside its magnitude. The VNA determines the phase by comparing the arrival time of the measured wave against its own internal reference oscillator. The result is expressed as degrees of rotation of that specific sine wave period. At low frequencies, a given physical delay produces a small phase shift. At high frequencies, the same physical delay produces many more degrees of rotation.

Phase is typically plotted on a secondary vertical axis (on the right side of the graph) against the same frequency horizontal axis used for the magnitude trace. Alternatively, phase appears on a completely separate rectangular graph beneath the magnitude plot. Both magnitude and phase can be combined onto a single polar Smith Chart, where the radial distance from the center represents magnitude and the angular position represents phase. The Smith Chart representation is identical to the one described in the companion guide: the impedance corresponding to each measured frequency appears as a point on the chart, and sweeping across frequency traces a trajectory that reveals the evolution of the impedance match.

## Multi-Port and Differential Measurement

A standard two-port VNA characterizes single-ended channels. Differential signaling, which dominates modern high-speed protocols (PCIe, USB, HDMI, Ethernet), requires simultaneous measurement of two coupled signal paths. A 4-port VNA handles this naturally. A single differential pair requires four ports: two at the transmitter end (corresponding to the D+ and D – signals) and two at the receiver end. This configuration produces a 4x4 S-parameter matrix with 16 entries.

The raw single-ended S-parameters from the 4-port measurement are mathematically transformed into mixed-mode S-parameters that decompose the channel behavior into differential mode (the desired signal) and common mode (the noise mode). Differential insertion loss (  $S_{dd21}$  ) measures how efficiently the differential signal propagates through the channel. Mode conversion parameters (  $S_{cd21}$  ,  $S_{dc21}$  ) measure how much energy converts between differential and common modes during transit, a phenomenon caused by asymmetries in the physical trace routing (length mismatch, unequal coupling to ground planes, or asymmetric via structures).

High port-count VNAs (8, 16, or 32 ports) extend this framework to characterize complex multi-lane buses. Crosstalk between adjacent lanes is measured by injecting a signal into one lane and recording the coupled energy appearing at the ports of neighboring lanes. Near-End Crosstalk (NEXT) measures the energy coupling backward to the aggressor's source end. Far-End Crosstalk (FEXT) measures the energy coupling forward to the victim's receiver end. Compliance specifications for multi-lane protocols define NEXT and FEXT isolation thresholds that must be met across the entire operating frequency range.

## IF Bandwidth and Dynamic Range

The precision of each frequency-domain measurement depends on how long the VNA dwells at each frequency step. The instrument controls this through its Intermediate Frequency (IF) bandwidth setting. The IF bandwidth defines the width of the detection filter centered on the measurement frequency.

A wider IF bandwidth allows the instrument to sweep quickly but admits more thermal noise into the measurement. Thermal noise is random, Gaussian, and zero-mean. The VNA test signal is a steady, deterministic sine wave. If the instrument measures for a very short duration, a random noise spike might superimpose on the signal, producing a false amplitude reading. Narrowing the IF bandwidth forces the analyzer to dwell longer at each frequency step, collecting more samples. The positive and negative random noise spikes mathematically cancel toward zero over longer averaging windows, while the deterministic sine wave remains constant. This noise cancellation lowers the measurement noise floor, allowing the analyzer to detect extremely faint signals buried deep in the noise.

The practical consequence is a direct tradeoff between sweep speed and dynamic range. A wide IF bandwidth produces a fast sweep with a shallow noise floor, adequate for measuring insertion loss on low-loss channels. A narrow IF bandwidth produces a slow sweep with a deep noise floor, necessary for measuring weak crosstalk signals or verifying the isolation of high-quality shielding.

## Worked Examples

---

### Reading an S-Parameter Plot

Consider a 12-inch PCB differential pair measured with a 4-port VNA from 10 MHz to 20 GHz. The  $S_{dd21}$  (differential insertion loss) trace begins near 0 dB at low frequencies, indicating that nearly all the signal power reaches the receiver. As frequency increases, dielectric absorption and skin effect losses progressively attenuate the signal. The trace slopes downward, reaching  $-3$  dB at approximately 8 GHz (the half-power frequency for this trace length) and  $-15$  dB at 20 GHz. The compliance specification for the protocol requires that  $S_{dd21}$  remain above the specification mask at every frequency point within the operating band.

The  $S_{dd11}$  (differential return loss) trace begins at a large negative value (for example,  $-30$  dB at low frequencies, indicating excellent matching). As frequency increases, parasitic resonances from vias and connectors create peaks where the return loss degrades. A spike reaching  $-10$  dB at 12 GHz indicates that 10% of the injected power reflects back at that frequency, signaling a localized impedance mismatch that worsens at that specific resonant frequency.

### Converting Between Domains: FFT and IFFT

The mathematical connection between TDR and VNA measurements is the Fourier transform. A Fast Fourier Transform (FFT) decomposes a time-domain signal into its constituent frequency components. A sharp time-domain step function consists of an infinite sum of continuous sine waves, and the FFT calculates the exact magnitude and phase of each sine wave required to construct that step.

Applied to TDR data, the FFT converts the measured step response into an equivalent  $S_{11}$  frequency-domain trace. The Inverse Fast Fourier Transform (IFFT) performs the reverse operation: it takes VNA S-parameter data and reconstructs the equivalent time-domain waveform. This bidirectional conversion allows engineers to view the same physical channel through either the spatial lens of TDR or the spectral lens of the VNA, choosing whichever representation best illuminates the specific defect under investigation.

## Edge Cases

---

### Passive Reciprocity and Its Violations

The reciprocity property (  $S_{21} = S_{12}$  ) holds strictly for any passive, linear, time-invariant network. A copper trace, a connector, a via, a passive filter, or any combination of these elements will always exhibit identical forward and reverse transmission. Violating reciprocity requires an active element (a transistor amplifier providing gain in one direction), a nonlinear element (a diode whose impedance changes with signal amplitude), or a non-reciprocal component (a ferrite circulator or isolator that exploits magnetic bias to break symmetry).

Verifying reciprocity in measured data serves as a built-in sanity check. If a VNA measurement of a passive PCB trace shows  $S_{21} \neq S_{12}$  beyond the instrument's measurement uncertainty, the discrepancy indicates a calibration error, a connector problem, or a systematic measurement artifact rather than a genuine physical asymmetry.

### The Tradeoff: Sweep Speed vs. Noise Floor

Narrowing the IF bandwidth improves dynamic range but proportionally increases the total sweep time. A 10 Hz IF bandwidth provides approximately 40 dB more dynamic range than a 10 kHz IF bandwidth, but the sweep takes 1000 times longer. Production environments performing high-volume compliance testing optimize this tradeoff by using a wide IF bandwidth for quick pass/fail screening and switching to a narrow IF bandwidth only when investigating marginal failures or measuring weak crosstalk signals that approach the noise floor.

### What S-Parameters Cannot Capture

S-parameters describe only the linear, time-invariant behavior of a network. They cannot represent phenomena that depend on signal amplitude (compression in active devices, dielectric nonlinearity at extreme voltages) or phenomena that change over time (thermal drift during a long measurement, aging of connector contacts). Systems that exhibit any of these behaviors require either large-signal measurement techniques or time-stamped repeated sweeps to track the variation.

## Chapter 9: Fourier Analysis

---

Every digital signal on a high-speed transmission line is simultaneously a time-domain waveform and a frequency-domain spectrum. Fourier analysis provides the mathematical framework that connects these two representations, proving that any voltage waveform can be decomposed into a unique set of sine waves with specific frequencies, amplitudes, and phases. The Fast Fourier Transform (FFT) and its inverse (IFFT) are the computational engines that perform this decomposition and reconstruction.

This duality is not merely academic. The entire discipline of signal integrity measurement rests on the ability to move between the time domain and the frequency domain. A TDR step response reveals impedance structures localized in physical space. An S-parameter sweep from a VNA reveals how those same structures behave as a function of frequency. The FFT is the mathematical bridge that converts between these two views of the same physical channel. Every loss mechanism, every parasitic interaction, and every equalization strategy becomes clearer when viewed through the appropriate domain.

This guide traces the Fourier framework from its mathematical foundations through its physical implications for signal propagation. The treatment begins with how the FFT algorithm operates, moves through the critical distinction between discrete and continuous spectra, examines why frequency content is determined by transition speed rather than data rate, and concludes with the practical consequences of bandwidth limitation, including the Gibbs phenomenon and spatial resolution degradation in time-domain measurements.

### Core Concepts

---

#### The FFT Algorithm: Correlation as Detection

The FFT operates through a process of mathematical correlation. The algorithm multiplies the incoming time-domain signal by a mathematically perfect test sine wave of a specific frequency, then integrates the product over time. If the target signal contains energy at that frequency, the integral evaluates to a large value. If it does not, the positive and negative lobes of the product cancel to zero. The algorithm sweeps across a full spectrum of test frequencies, computing the correlation at each one, to build the complete frequency-domain representation.

The result at each test frequency is a complex number encoding two pieces of information: the magnitude (how much energy the signal contains at that frequency) and the phase (the timing offset of that frequency component relative to a reference). Together, these form the complete spectral portrait of the signal.

The Inverse Fast Fourier Transform (IFFT) performs the exact reverse operation. It takes a frequency spectrum and reconstructs the original time-domain waveform by mathematically summing all the indicated sine waves at their respective amplitudes and phase shifts. This process is called Fourier synthesis. The FFT and IFFT are lossless complements: applying the FFT followed by the IFFT recovers the original waveform exactly.

## Discrete Spectra vs. Continuous Spectra

A repeating signal and a non-repeating signal produce fundamentally different frequency-domain representations.

A repeating clock signal or square wave contains only discrete, widely spaced frequency components. A 1 GHz square wave contains energy at 1 GHz, 3 GHz, 5 GHz, 7 GHz, and so on, with nothing in between. These individual spectral lines appear as distinct vertical spikes on an FFT plot, separated by empty gaps.

A single, non-repeating voltage step possesses a completely continuous frequency spectrum. The signal contains energy at 1 GHz, at 1.000001 GHz, at 1.000002 GHz, and at every other conceivable fractional frequency. This continuous spectrum begins at exactly 0 Hz (the DC offset that establishes the final steady-state voltage) and extends without interruption toward infinity.

The distinction matters for measurement. A spectrum analyzer sweeping across a repeating clock signal detects energy only at the harmonic frequencies and measures silence between them. The same analyzer sweeping across a step function detects energy at every frequency in an unbroken continuum.

## Why Square Waves Contain Only Odd Harmonics

A perfect square wave possesses a geometric property called half-wave symmetry: the positive half of each cycle is the exact inverted mirror image of the negative half. The Fourier mathematics of this symmetry dictate which harmonics survive and which cancel.

A fundamental sine wave shares this mirror symmetry. Adding a sine wave at exactly three times the fundamental frequency (the 3rd harmonic) with one-third the amplitude pushes the rounded peaks of the fundamental downward, flattening the top while steepening the vertical transitions. Adding the 5th harmonic at one-fifth the amplitude flattens the waveform further. Each successive odd harmonic brings the composite shape closer to a perfect square wave.

Even harmonics (the 2nd, 4th, 6th, and so on) complete an integer number of full cycles within one fundamental period in a way that destroys the mirror symmetry. Adding a 2nd harmonic pushes the positive half of the fundamental wave upward, but it also pushes the negative half upward by the same amount. The resulting waveform is skewed and asymmetric. The Fourier integral, applied to a shape with perfect half-wave symmetry, evaluates to exactly zero for every even harmonic coefficient. The mathematics enforce the symmetry of the physics.

## Frequency Content Is Determined by Transition Speed

A fundamental result of Fourier mathematics connects the speed of a voltage transition to its frequency content. A signal that takes 1 nanosecond to rise from 0V to 1V contains frequency components extending to approximately 350 MHz. A signal that snaps through the same transition in 10 picoseconds contains components extending past 35 GHz.

The practical formula relating rise time to maximum usable frequency is the knee frequency:

$$F_{\text{knee}} = \frac{0.35}{T_{\text{rise}}}$$

A physical TDR with a 20-picosecond rise time generates a step function containing a continuous, uninterrupted spectrum from 0 Hz up to approximately 17.5 GHz. Past this knee frequency, the amplitudes of the higher-frequency components decay precipitously until they effectively vanish. The absence of these extreme high-frequency components prevents the physical signal from achieving a perfectly sharp 90-degree corner at the transition, forcing a finite slope instead.

Two distinct frequency scales coexist on every high-speed transmission line:

- **The data rate frequency** describes how often the driver toggles between states. An alternating 10101010 pattern at 10 Gbps has a fundamental frequency of 5 GHz. A long run of identical bits drops to 0 Hz for that duration.
- **The edge bandwidth** describes the high-frequency harmonics required to construct the shape of each rising or falling transition. This bandwidth is set entirely by the transistor's switching speed, independent of the bit pattern.

Both scales propagate simultaneously on the same physical trace. Any mechanism that treats high frequencies differently from low frequencies will cause the composite waveform to distort and smear as it propagates.

## All Frequencies Travel Simultaneously

The Fourier decomposition of a digital signal reveals a counter-intuitive physical truth: all frequency components occupy the same physical space at the same time. A digital bitstream is not a signal that changes its frequency moment by moment as the bit pattern shifts. It is the superposition of hundreds of continuous, unbroken sine waves running simultaneously through the transmission line.

A long run of identical bits (such as 1111100000 ) is dominated by low-frequency components that hold the flat voltage levels. A rapid alternating pattern ( 10101010 ) is dominated by high-frequency components. The sharp vertical edge of every transition, regardless of the surrounding bit pattern, requires extreme high-frequency components to construct its steep slope.

The trace physically carries low-frequency waves (sustaining the flat voltage plateaus) and high-frequency waves (building the steep transitions) at every point along its length. Frequency-dependent physical mechanisms such as the skin effect and dielectric absorption do not selectively target one "part" of the signal. They attenuate the high-frequency components everywhere, simultaneously, across the entire trace. The composite waveform distorts because its constituent sine waves are no longer arriving at the receiver with their original amplitude relationships intact.

## Worked Examples

---

### Fourier Decomposition of Common Waveforms

The FFT and IFFT produce characteristic signatures for several fundamental waveforms that appear repeatedly in signal integrity work:

**Pure sine wave:** A single-frequency sine wave in the time domain produces a single, sharp vertical spike at that exact frequency on the FFT. No other frequency components exist.

**Square wave:** A time-domain square wave produces a spike at the fundamental frequency, plus a series of progressively smaller spikes at every odd harmonic (3f, 5f, 7f, ...), with each harmonic's amplitude decaying as  $1/n$  where  $n$  is the harmonic number. The even harmonic positions are empty.

**Dirac delta (ideal impulse):** An infinitely narrow, infinitely tall time-domain impulse produces a perfectly flat, horizontal line on the FFT. Equal energy exists at every frequency from zero to infinity. This transform pair is the mathematical reason that impulse response testing can characterize a channel across its entire frequency range.

**Brick-wall bandpass filter:** A frequency spectrum with a perfectly flat passband and a sudden, vertical cutoff translates into a sinc function (  $\sin(x)/x$  ) in the time domain. The sinc shape features a main pulse flanked by symmetrically rippling precursor and post-cursor lobes that decay gradually with distance from the peak.

**Gaussian pulse:** A Gaussian amplitude envelope in the frequency domain transforms into a Gaussian pulse in the time domain. The Gaussian is the only mathematical shape that is identical in both domains, making it uniquely well-behaved for signal processing.

## Step Function Synthesis: From Sine Waves to a Voltage Wall

Building a step function from its Fourier components illustrates how constructive and destructive interference produce a sharp transition followed by a flat plateau. The process proceeds as follows.

The first component is the 0 Hz DC offset. On a graph with time on the horizontal axis and voltage on the vertical axis, this appears as a perfectly flat horizontal line at exactly half the step amplitude. For a 0-to-1V step, the DC baseline sits at 0.5V. This component establishes the average value that all subsequent sine waves ride upon.

The second component is the lowest-frequency sine wave in the spectrum. This wave possesses the largest amplitude of any component. It is centered symmetrically on the 0.5V baseline and aligned in phase so that it crosses the baseline with a positive slope at exactly  $t = 0$ . Adding only this single wave to the DC offset produces a smooth, rolling oscillation between 0V and 1V.

Each subsequent component is a sine wave at a progressively higher frequency with a progressively smaller amplitude. Every wave must cross the 0.5V baseline with a positive slope at exactly  $t = 0$ . This phase alignment is not arbitrary; it is mathematically required by the shape of the target step function.

Adding these components one by one transforms the waveform geometrically. At  $t = 0$ , every sine wave is crossing zero with a positive slope simultaneously. This extreme constructive interference steepens the transition, forcing the composite slope to approach a vertical wall. At  $t > 0$ , the sine waves are oscillating at their vastly different frequencies and rapidly falling out of synchronization. The destructive interference between them causes the ripples on the plateau to become progressively more rapid and smaller in amplitude. As more components are added, the ripples shrink toward invisibility, and the plateau converges toward a flat line at 1V. At  $t < 0$ , the same destructive interference flattens the waveform toward 0V.

The flat line at 1V after the transition does not result from the sine waves ceasing to exist. They continue oscillating indefinitely. The specific recipe of decaying amplitudes and synchronized phases causes every positive peak of one component to be precisely canceled by the negative troughs of carefully calculated combinations of other components. The alternating high-frequency ripples destroy each other, but they do not destroy the underlying DC offset. The destructive interference erases only the transient oscillations, leaving the 0 Hz component exposed as the steady plateau.

## Superposition and Phase: Why the Result Is Not Zero

Summing an infinite number of sine waves does not automatically produce zero. The outcome depends entirely on the phase relationships among the components.

If the phases are completely random and uncorrelated, the superposition generates white noise with an average value of zero. Random phases cause random constructive and destructive interference at every point, producing a signal with no coherent structure.

A step function requires the opposite condition. The electronics that generate the step inject every frequency component with a precisely determined phase. At  $t = 0$ , every sine wave in the spectrum crosses the zero axis with a positive slope simultaneously. This extreme synchronization prevents the components from canceling to zero at the transition point. They constructively interfere to build the vertical wall. After the transition, the synchronized phases cause the specific pattern of destructive interference that produces the flat plateau rather than continued oscillation.

The critical distinction is between random superposition (which produces noise) and deterministic superposition (which produces coherent waveforms). Fourier analysis works precisely because the decomposition preserves the exact phase of every component, allowing the IFFT to reconstruct the original waveform without loss.

## The TDR-to-S-Parameter Conversion

The FFT provides the mathematical bridge between time-domain reflectometry and frequency-domain network analysis. A TDR measures the reflected voltage as a function of time, recording how the impedance profile of a channel varies along its physical length. Applying an FFT to this time-domain step response decomposes the waveform into its frequency components, producing the equivalent  $S_{11}$  reflection coefficient as a function of frequency.

The IFFT performs the reverse conversion. Starting from VNA S-parameter data measured across a swept frequency range, the IFFT reconstructs the equivalent time-domain impulse response. This bidirectional conversion allows engineers to view the same physical channel through the spatial lens of the TDR or the spectral lens of the VNA, choosing whichever representation best illuminates the defect under investigation.

The practical value of this bridge extends beyond convenience. Certain physical phenomena are opaque in one domain but transparent in the other. A small parasitic capacitance from a via barrel may produce only a subtle dip on a TDR trace, difficult to distinguish from nearby reflections. The same capacitance, viewed as a frequency-dependent  $S_{11}$  response on the VNA, reveals its characteristic  $1/(2\pi fC)$  impedance roll-off clearly separated from resistive or inductive effects.

## Architecture

---

### Physical Reality vs. Mathematical Decomposition

A conceptual trap surrounds the Fourier description of signal generation. Describing a step function as "composed of" an infinite spectrum of sine waves can create the impression that the generating instrument (a TDR, an oscilloscope calibrator, or a SerDes transmitter) physically synthesizes and sums thousands of individual sine wave oscillators.

The physical mechanism is far simpler. A DC power supply holds a steady voltage. A high-speed transistor acts as a gate. At  $t = 0$ , the transistor closes, connecting the voltage supply directly to the trace. The flat plateau at  $t > 0$  exists because the transistor remains closed, permanently connecting the trace to the supply. No sine wave generators are involved.

Fourier analysis acts as a mathematical lens applied after the fact. The physical voltage step exists first. The Fourier Transform equation proves that the shape of that step intrinsically contains a continuous spectrum of frequency components. The analogy to optical spectroscopy is precise: a beam of white light exists as a single physical entity. Passing it through a glass prism reveals its constituent wavelengths, but no one constructed the white light by aligning millions of individual colored lasers. The decomposition reveals structure that was always present in the original signal.

This distinction is essential when reasoning about loss mechanisms. The skin effect does not wait for a Fourier analyzer to decompose the signal before selectively attenuating high frequencies. The physical copper and dielectric interact with the electromagnetic wavefront as it passes. The high-frequency content is present in the steep slope of the transition itself. The skin effect attenuates the rapid field variations that constitute that steep slope, softening the edge. Fourier analysis provides the quantitative framework to predict exactly how much softening will occur at each frequency, but the physical interaction operates on the waveform directly, not on its mathematical decomposition.

### How Bandwidth Limitation Degrades Spatial Resolution

The connection between frequency content and spatial resolution in time-domain measurements follows directly from Fourier principles. A TDR step response can resolve impedance features only as small as the spatial extent of its rising edge. The rising edge, in turn, is determined by the highest-frequency content present in the step.

A TDR launches a step function with a finite rise time set by the instrument's analog bandwidth. As this step propagates down a lossy transmission line, the skin effect and dielectric absorption progressively attenuate the highest-frequency components. The physical consequence is that the rising edge spreads out and slows down. A 20-picosecond edge at the launcher might degrade to a 100-picosecond edge after traversing several inches of FR4.

This degraded edge interacts differently with impedance discontinuities than the original sharp edge did. A via barrel that acts as a parasitic capacitance has an impedance of  $Z = 1/(2\pi fC)$ . At extremely high frequencies, this impedance drops dramatically, and the via appears as a strong partial short circuit to ground, generating a sharp, deep reflection. At lower frequencies, the impedance rises toward infinity, and the via becomes electrically invisible. The sharp original edge, rich in high-frequency content, produces a crisp, narrow dip on the TDR trace that precisely marks the via location. The degraded edge, depleted of its high-frequency content, produces only a shallow, smeared response.

The practical consequence is that a TDR loses spatial resolution progressively along the length of the trace. Identical vias at different distances from the launcher appear differently on the TDR screen. The first via, struck by the sharpest edge, produces a well-defined impedance dip. A via further down the trace, struck by an edge already degraded by intervening losses, produces a broader and shallower response. The second via has not changed physically. The measurement resolution has degraded because the probe signal has lost the high-frequency content needed to interact with small structures.

## Edge Cases

---

### The Gibbs Phenomenon

Performing Fourier synthesis with a finite number of frequency components reveals a mathematical artifact at every sharp discontinuity. Stopping the summation at a finite bandwidth limit (for example, at 20 GHz to simulate the bandwidth of a physical instrument) prevents the vertical transition from becoming perfectly sharp. More significantly, the truncated synthesis produces a distinct overshoot and a localized oscillatory ringing immediately adjacent to the discontinuity before the waveform settles to its steady-state value.

This ringing is the Gibbs phenomenon. It is the visible confirmation that the summation was terminated before reaching mathematical infinity. The overshoot amplitude converges to approximately 9% of the step height regardless of how many terms are included; adding more frequency components narrows the ringing but does not eliminate the overshoot. Only the inclusion of frequency components extending to true mathematical infinity would produce a perfectly clean transition.

The Gibbs phenomenon has a direct physical counterpart in measurement instrumentation. The overshoot and ringing visible on an oscilloscope when a fast step function encounters the instrument's analog bandwidth limit are physically equivalent to the truncated Fourier synthesis. The instrument's bandwidth acts as the finite summation cutoff. Signal integrity engineers must distinguish between Gibbs ringing (an artifact of measurement bandwidth) and physical ringing caused by actual impedance mismatches on the transmission line. The two produce similar waveform signatures but arise from fundamentally different mechanisms.

### The DC Component and Amplitude Independence

The Fourier decomposition of a step function requires a 0 Hz DC component to establish the steady-state voltage level. Summing an infinite number of pure, un-offset sine waves centered around 0V can never produce a flat line at 1V. The high-frequency components construct only the shape of the transition; the DC component sets the amplitude of the plateau.

This separation means that the Fourier mathematics governing the transition shape are independent of the physical voltage level. The same recipe of harmonics, with the same relative amplitudes and phases, constructs a step from 0V to 0.5V, from 0V to 1V, or from 0V to 3.3V. Only the DC offset and the overall scaling factor change. The voltage divider that determines the launched amplitude (the TDR source impedance and the line impedance forming a resistive divider) operates independently of the Fourier spectral content that determines the edge shape.

## **When Fourier Analysis Breaks Down**

The Fourier framework assumes linearity and time invariance. A linear system produces an output proportional to its input, and superposition holds: the response to the sum of two signals equals the sum of the individual responses. A time-invariant system responds identically regardless of when the input is applied.

Physical transmission lines at moderate signal levels satisfy both conditions well. The skin effect, dielectric loss, and impedance mismatches are all linear phenomena whose effects can be analyzed independently for each frequency component and then superposed. Nonlinear elements such as semiconductor junctions, saturating ferrite cores, or ESD protection diodes violate the superposition principle. The Fourier decomposition of the input signal remains valid, but predicting the output by independently processing each frequency component and summing the results does not produce the correct answer. Nonlinear analysis requires time-domain simulation methods that process the composite waveform directly.



PART 3

# Jitter Analysis

## Chapter 10: Jitter Decomposition and Measurement

---

Jitter is the deviation of a signal edge from its ideal position in time. At modern serial data rates, no single number for "total jitter" provides enough information to diagnose a failing link. A signal that misses its timing budget could be failing because of thermal noise in the silicon, because of impedance discontinuities on the PCB, because of switching power supply coupling, or because of asymmetric transistor drive strength in the transmitter. Each of these root causes produces a physically distinct timing perturbation, and each requires a different corrective action.

The jitter decomposition framework separates total jitter into categories based on the statistical behavior and physical origin of each component. Random Jitter arises from thermodynamic processes that are fundamentally unbounded and Gaussian. Deterministic Jitter arises from structural and environmental sources that are bounded and repeatable. Within Deterministic Jitter, further sub-classification isolates pattern-dependent effects (Intersymbol Interference, Duty Cycle Distortion) from pattern-independent periodic disturbances (power supply coupling, EMI, clock crosstalk). This taxonomy transforms an opaque timing failure into a set of actionable hardware diagnostics.

This guide covers the physical origins of each jitter component, the measurement infrastructure that captures them, and the statistical properties that make decomposition possible.

### Core Concepts

---

#### The Decomposition Taxonomy

All jitter falls into two primary categories, distinguished by their probability distributions:

**Random Jitter (RJ)** is caused by fundamental thermodynamic processes: thermal noise (Johnson-Nyquist noise from resistive elements), shot noise (discrete charge carriers crossing semiconductor junctions), and flicker noise (1/f noise in active devices). These microscopic processes involve the aggregate behavior of billions of independent electrons. The Central Limit Theorem guarantees that the sum of billions of independent random events produces a macroscopic distribution that follows a Gaussian probability curve with mathematical precision.

The critical property of RJ is that it is unbounded. The Gaussian distribution has tails that extend to infinity in both directions. Measuring for a longer duration does not change the underlying distribution; it simply reveals increasingly rare events further out in the tails. A peak-to-peak measurement of RJ is therefore meaningless because the observed peak-to-peak value grows without limit as the sample size increases. RJ is quantified exclusively by its standard deviation,  $RJ_{rms}$ , which is a fixed physical constant determined by the temperature and material properties of the silicon.

**Deterministic Jitter (DJ)** is caused by predictable, systemic characteristics of the hardware, channel, and electromagnetic environment. DJ is bounded: it has strict physical minimum and maximum limits that do not grow with observation time. DJ is quantified as a peak-to-peak value,  $DJ_{p-p}$ .

DJ decomposes further into sub-categories that isolate specific hardware mechanisms:

- **Data Dependent Jitter (DDJ)** is correlated to the transmitted bit sequence. It encompasses Intersymbol Interference (ISI) and Duty Cycle Distortion (DCD).
- **Periodic Jitter (PJ)** is uncorrelated to the data but repeats at a fixed frequency driven by an external oscillating source.

## Random Jitter: Physics and Statistical Invariance

Thermal noise originates from the random thermal agitation of electrons colliding with atoms in the silicon lattice and copper interconnect. This is a voltage (amplitude) perturbation, not a timing perturbation. The physical mechanism that converts amplitude noise into timing jitter is the finite slew rate of the signal edge.

A high-speed signal edge is not a perfect vertical transition; it traverses from the low voltage rail to the high voltage rail over a finite time interval, producing a slope ( $dV/dt$ ). When random voltage noise superimposes on the data signal during this transition, the combined voltage (  $V_{\text{received}} = V_{\text{data}} + V_{\text{noise}}$  ) shifts the point at which the edge crosses the receiver's decision threshold. A negative noise spike during a rising edge pulls the combined voltage downward, forcing the signal to travel further along its slope before reaching the threshold. The crossing point shifts to the right in time. A positive noise spike accelerates the crossing to the left. This geometric relationship between voltage perturbation and timing shift is called amplitude-to-phase conversion.

The Gaussian distribution of the underlying voltage noise maps directly through this conversion into a Gaussian distribution of timing jitter. The shape of the distribution is preserved because the conversion is linear over the small perturbation range near the threshold crossing.

RJ is fundamentally uncorrelated and memoryless. The thermal agitation that shifts the current edge has no physical mechanism to influence the position of the next edge. Each edge crossing is an independent statistical trial drawn from the same Gaussian distribution. This is the defining contrast with ISI, where the timing of the current edge is strictly dictated by the history of preceding bits.

## $RJ_{\text{rms}}$ as a Physical Time Coordinate

$RJ_{\text{rms}}$  is not a count of edge crossings. It is a physical time distance on the horizontal axis of the TIE histogram, measured in picoseconds.

$RJ_{\text{rms}}$  equals exactly one standard deviation (  $1\sigma$  ) of the Gaussian distribution. Geometrically, it defines the distance from the center peak of the Gaussian to the inflection point where the curve transitions from concave-down to concave-up. The mathematics of the Gaussian function dictate that 34.1% of all edge crossings will fall within one  $\sigma$  on either side of the mean, totaling 68.2% within the  $\pm 1\sigma$  window.

The stability of  $RJ_{\text{rms}}$  as a measurement derives from a fundamental property of the Gaussian distribution: its variance converges. Variance is the expected value of the squared distance from the mean, weighted by the probability of landing at each distance. While the squared distance grows quadratically (  $x^2$  ) as you move further from the center, the probability of landing at that distance decays exponentially (  $e^{-x^2}$  ). Exponential decay dominates quadratic growth at every scale. By the time an outlier event occurs far out in the tail, its probability is so vanishingly small that multiplying it by the large squared distance contributes effectively zero to the total variance sum. The mathematical series converges, and the standard deviation reaches a fixed value determined by the physics of the noise source, not by how long the oscilloscope runs.

This convergence is what allows  $RJ_{\text{rms}}$  to serve as the input to BER extrapolation. The static  $RJ_{\text{rms}}$  constant is multiplied by a Q-factor (derived from the complementary error function for a target BER) to predict the peak-to-peak Random Jitter spread at any desired confidence level. For a BER of  $10^{-12}$ , the Q-factor is 7.03, and the total multiplier is  $2Q \approx 14$ , accounting for both the early and late tails of the distribution.

## Intersymbol Interference (ISI)

ISI is the dominant source of jitter in modern high-speed serial links. It arises because the physical channel (PCB traces, vias, cables, connectors) acts as a low-pass filter that attenuates high-frequency signal components more aggressively than low-frequency components.

Three physical mechanisms produce this frequency-dependent attenuation:

1. **Skin effect:** At high frequencies, eddy currents force the signal current to the extreme outer perimeter of the copper conductor. The usable cross-sectional area shrinks, and the effective resistance increases proportionally. The skin depth is inversely proportional to the square root of frequency, so a 10 GHz signal component encounters dramatically higher resistance than a 1 GHz component traveling the same trace.
2. **Dielectric loss:** The FR4 or advanced laminate material absorbs electromagnetic energy and dissipates it as heat. The absorption increases with frequency, further attenuating high-frequency components.
3. **Parasitic capacitance:** The copper trace above the ground plane forms a distributed capacitor that resists instantaneous voltage changes. Charging a capacitor instantaneously requires infinite current; the finite drive strength of the transmitter limits the achievable slew rate, governed by the local RC time constant.

The interaction between these mechanisms and the data pattern is what creates ISI:

A long run of identical bits (such as **00000**) acts as a near-DC signal. The channel's distributed capacitance has ample time to fully charge or discharge to the absolute voltage rail. The signal reaches maximum amplitude.

An alternating **10101** pattern represents the highest-frequency content the link carries (the Nyquist rate). The transmitter pulls the voltage upward, but before the parasitic capacitance can charge to full amplitude, the transmitter reverses direction and pulls the voltage back down. The peak-to-peak voltage swing of the alternating pattern is substantially attenuated compared to the low-frequency run.

When the data transitions from a long run of zeros to a one (**000001**), the voltage must climb from a deeply discharged state. It has a longer voltage distance to travel to reach the receiver's decision threshold compared to a transition following an alternating pattern, which started from a shallower voltage. The transmitter's slew rate is fixed by the silicon, so traversing a longer voltage distance requires more time. The edge crosses the threshold later than it would in the alternating-pattern case. This pattern-dependent timing shift is ISI jitter.

## Duty Cycle Distortion (DCD)

DCD occurs when the duration of a logic "1" does not equal the duration of a logic "0." Two distinct physical mechanisms produce DCD:

**Transmitter slew rate asymmetry:** If the P-channel pull-up transistor in the CMOS output driver has different drive strength than the N-channel pull-down transistor, rising edges will slew at a different rate than falling edges. The faster edge crosses the threshold earlier; the slower edge crosses later. The resulting jitter histogram shows a systematic offset between rising-edge and falling-edge crossing times.

**Receiver threshold offset:** Even if the transmitter generates a perfectly symmetrical waveform, a DC offset in the receiver's decision threshold voltage will create apparent DCD. If the threshold is shifted upward from the true center of the waveform, the receiver intersects the rising edge later (higher on its slope) and the falling edge earlier (higher on its slope). The receiver interprets the logic "1" as shorter in duration and the logic "0" as longer. The hardware is flawless; the measurement reference point is wrong.

## Periodic Jitter (PJ)

Periodic Jitter is uncorrelated to the data pattern but repeats at a fixed, identifiable frequency. It is caused by external oscillating sources coupling into the signal path:

- Switching power supply noise at the regulator's switching frequency
- Electromagnetic interference from nearby circuits or external sources
- Crosstalk from an adjacent clock signal routed on the PCB
- A PLL that has not achieved stable phase lock, causing the recovered clock to oscillate around its target frequency

PJ appears in the frequency domain as sharp spectral spikes at the interfering source's frequency, rising distinctly above the broadband noise floor of RJ.

## RJ Symmetry vs. DJ Asymmetry

The decomposition taxonomy rests on a fundamental physical asymmetry between the two primary categories.

Random Jitter is guaranteed to follow a perfectly symmetric Gaussian curve. The microscopic thermodynamic processes that generate it possess no directional bias. An electron is exactly as likely to experience a positive thermal perturbation as a negative one. The resulting timing distribution is centered on the ideal crossing time with identical tails on both sides.

Deterministic Jitter can be profoundly asymmetric. DCD from a stronger pull-up transistor will concentrate rising-edge crossings early and falling-edge crossings late, producing a lopsided histogram. ISI from a specific via reflection may interfere strongly with one bit transition pattern ( 001 ) while leaving the complementary pattern ( 110 ) unaffected, creating a heavy concentration of delayed edges on one side of the distribution.

This physical asymmetry is precisely why measurement algorithms cannot assume a simple symmetric distribution for total jitter. The Dual-Dirac model, covered in detail in [Dual-Dirac Model and BER Extrapolation](#), addresses this by placing two independent impulses that need not be equidistant from the ideal crossing, nor carry equal populations of hits.

## Worked Examples

---

### Time Interval Error (TIE) Measurement

The measurement infrastructure for jitter decomposition begins with Time Interval Error.

The oscilloscope digitizes the incoming voltage waveform at high sample rates and mathematically interpolates the precise moment each edge crosses the differential zero-volt threshold (or the single-ended decision threshold). Simultaneously, the instrument runs a software Phase-Locked Loop on the captured data to generate a perfect, ideal recovered clock.

For each edge, the instrument subtracts the ideal clock time from the actual edge crossing time. The resulting difference is the TIE for that specific transition. A positive TIE means the edge arrived late (lagging the ideal clock). A negative TIE means the edge arrived early (leading the ideal clock). The zero point on the TIE axis is not absolute zero time; it is the relative reference defined by the ideal recovered clock.

The oscilloscope accumulates millions of TIE values and bins them into a histogram. The horizontal axis is time (divided into femtosecond or picosecond bins), and the vertical axis is the count of edge crossings that landed in each bin.

### Normalizing to a Probability Density Function

To perform mathematical operations on the histogram, the oscilloscope normalizes the raw hit counts into a Probability Density Function (PDF).

Each bin's hit count is divided by the total number of edges captured in the entire measurement run. The vertical axis transforms from integer counts to probability density. The total area under the normalized curve equals exactly 1.0, representing the certainty that every edge must land somewhere on the time axis.

The resulting PDF is the convolution of the DJ distribution with the RJ Gaussian distribution. The full measured histogram reflects the combined effect of all jitter sources acting simultaneously on every edge crossing.

### Extracting $RJ_{\text{rms}}$ from the Histogram Tails

The extraction of Random Jitter exploits the bounded nature of Deterministic Jitter.

DJ has absolute physical limits. No matter how many edges are captured, no DJ mechanism can push an edge beyond its maximum displacement. Therefore, the extreme left and right tails of the TIE histogram are entirely free of DJ influence. Every edge crossing that lands far enough from center arrived there solely because of unbounded Random Jitter.

The oscilloscope algorithm isolates these far tails and fits a standard Gaussian function to each one using least-squares regression. The fit solves for the standard deviation  $\sigma$ , which is  $RJ_{\text{rms}}$ , and the mean of each fitted Gaussian, whose separation defines  $DJ_{\delta\delta}$  in the Dual-Dirac framework. The algorithm deliberately ignores the complex, multi-modal center of the histogram where DJ and RJ contributions are entangled.

## TJ Extrapolation

With  $RJ_{\text{rms}}$  and  $DJ_{\delta\delta}$  extracted, the oscilloscope calculates Total Jitter at any target BER without requiring hours of measurement to observe trillions of bits directly:

$$TJ(BER) = DJ_{\delta\delta} + 2 \cdot Q \cdot RJ_{\text{rms}}$$

For a BER of  $10^{-12}$ ,  $Q \approx 7.03$ , giving  $2Q \approx 14$ . The total jitter at  $10^{-12}$  BER is  $DJ_{\delta\delta} + 14 \cdot RJ_{\text{rms}}$ .

This formula reflects the physical reality that a bit error occurs when an edge shifts past the receiver's sampling threshold. The probability of that event is the area under the Gaussian tail beyond the threshold coordinate. The Q-factor, derived from the complementary error function (erfc), is the number of standard deviations required to leave exactly the target BER's worth of area in the tail. The detailed derivation of the Q-factor and erfc is covered in [Dual-Dirac Model and BER Extrapolation](#).

## Architecture

### The Jitter Tree

The full decomposition hierarchy forms a tree that maps every timing perturbation to a measurable root cause:

```
Total Jitter (TJ)
├─ Random Jitter (RJ)
│   └─ Gaussian, unbounded, quantified by RJ_rms
│       Sources: thermal noise, shot noise, flicker noise
├─ Deterministic Jitter (DJ)
│   └─ Data Dependent Jitter (DDJ)
│       └─ Intersymbol Interference (ISI)
│           └─ Channel low-pass filtering, reflections
│       └─ Duty Cycle Distortion (DCD)
│           └─ Slew rate asymmetry, threshold offset
└─ Periodic Jitter (PJ)
    └─ Power supply, EMI, clock crosstalk, PLL instability
```

Each leaf of the tree corresponds to a distinct physical mechanism and a distinct corrective action. ISI is addressed through channel design improvements or equalization. DCD is addressed through transmitter calibration or receiver threshold adjustment. PJ is addressed by identifying and decoupling the interfering oscillating source. RJ is addressed only by reducing the operating temperature or changing the semiconductor process, neither of which is typically practical, so RJ is treated as a fixed noise floor in the jitter budget.

## Measurement Flow: From TIE to Root Cause

The oscilloscope decomposes total jitter through a layered measurement process:

1. **Capture TIE values** for every edge crossing in the measurement window.
2. **Construct the TIE histogram** and normalize to a probability density function.
3. **Fit Gaussians to the extreme tails** to extract  $RJ_{\text{rms}}$  and  $DJ_{\delta\delta}$ .
4. **Record the TIE Track** (chronological TIE values vs. time) to preserve temporal ordering.
5. **Apply synchronous pattern averaging** on the TIE Track over many repetitions of a known test pattern (such as PRBS-7) to isolate DDJ. Uncorrelated components (RJ and PJ) average to zero; only the pattern-correlated ISI and DCD survive.
6. **Subtract DDJ from the TIE Track** to produce the uncorrelated jitter residual.
7. **Apply FFT to the residual** to separate PJ (sharp spectral spikes at specific frequencies) from RJ (flat broadband noise floor).

The full detail of steps 4 through 7, including test pattern selection (PRBS-7/15/31), pattern averaging mathematics, and FFT spectral analysis, is covered in [Test Patterns and Jitter Isolation](#).

## Differential Noise Cancellation

Differential signaling provides a natural defense against common-mode interference. If electromagnetic noise couples equally onto both the D+ and D- traces of a differential pair, it appears as common-mode noise. The differential receiver measures only the difference between the two traces ( $V_{D+} - V_{D-}$ ). Equal noise on both lines cancels exactly, leaving the data signal undisturbed.

An analogous mechanism operates in source-synchronous architectures such as DDR memory. If low-frequency jitter affects both the data line and the forwarded clock identically, the receiver's sampling clock shifts in perfect lockstep with the data drift. The phase relationship between data and clock remains constant, and the jitter is effectively neutralized. This cancellation fails when the jitter source affects data and clock differently, as occurs when high-frequency jitter exceeds the bandwidth of the forwarded clock path.

## Edge Cases

---

### ISI and the Closed Eye

In a functioning link, ISI shifts the threshold crossing by a fraction of the Unit Interval. The signal still reaches the decision threshold within the allotted bit period, but it arrives late (or early), contributing deterministic jitter to the timing budget.

If channel attenuation is severe enough, the signal may fail to reach the threshold entirely. The voltage begins climbing from a deeply discharged state after a long run of opposite bits, the bit period expires, and the transmitter reverses direction before the voltage crosses the threshold. The receiver samples a value on the wrong side of the decision boundary and records a bit error. On an oscilloscope, the eye diagram is physically closed: the voltage trajectories overlap and no open region remains for the receiver to sample cleanly.

This failure mode is the direct motivation for transmitter pre-emphasis (Feed-Forward Equalization) and receiver equalization (CTLE and DFE). Equalizers artificially boost high-frequency transition energy to guarantee threshold crossing regardless of the preceding bit history. Equalization is covered in [Transmitter FFE](#) and [CTLE and DFE](#).

## Peak-to-Peak RJ: A Defined Quantity, Not a Measured One

A common source of confusion is the notion that Random Jitter should have a measurable peak-to-peak value. The Gaussian distribution extends to infinity; there is no maximum displacement. Any peak-to-peak RJ number reported by an instrument is not a direct measurement but a statistical calculation:

$RJ_{p-p} = 2Q \cdot RJ_{rms}$ , where  $Q$  is chosen for a specific target BER. Changing the target BER changes the reported peak-to-peak value, even though the underlying physical noise source has not changed. The  $RJ_{rms}$  constant is the only measurement that reflects the true physics; the peak-to-peak value is a derived statistical projection.

## The Dual-Dirac Approximation and Its Limits

The Dual-Dirac model simplifies the entire complex DJ distribution into two Dirac delta impulses separated by  $DJ_{\delta\delta}$ . This is a worst-case simplification: it concentrates all DJ energy at two extreme points, producing the widest possible convolution with the RJ Gaussian. The model is mathematically convenient and adopted by industry standards (PCIe, Fibre Channel), but it can overestimate total jitter when the actual DJ distribution is spread across many smaller values rather than concentrated at two extremes. The full treatment of Dual-Dirac model mechanics, including asymmetric delta placement, constrained fitting, and Q-factor normalization, is in [Dual-Dirac Model and BER Extrapolation](#).

## Chapter 11: Dual-Dirac Model and BER Extrapolation

---

High-speed serial standards specify compliance at Bit Error Rates that are impossible to observe directly. PCIe Gen 4 requires  $10^{-12}$ , meaning only one bit in a trillion may arrive outside the timing window. At 16 GT/s, accumulating a trillion bits takes over 60 seconds of continuous transmission, and observing even a single error at that rate requires statistical patience that exceeds practical measurement timescales by orders of magnitude. The fundamental problem is not measurement speed; it is that the BER target describes an event so rare that the Gaussian tail governing its probability has never been directly populated in the captured histogram.

The Dual-Dirac model resolves this by exploiting a structural property of the Gaussian distribution: the shape of the unmeasured tail is mathematically determined by the measurable center. The model decomposes the total jitter histogram into two Gaussian populations anchored at the Deterministic Jitter boundaries, fits the visible tail slopes to extract a single Random Jitter standard deviation, and then uses the complementary error function to project that standard deviation outward to the target BER. The result is a Total Jitter number at  $10^{-12}$  (or any other target) derived from a measurement of a few hundred thousand edges, not a trillion.

This guide covers the physical and mathematical machinery of the Dual-Dirac model: why the Gaussian tail shape is locked by the center mass, how the complementary error function maps between tail area and boundary position, how the Q-factor normalizes this mapping into a hardware-independent constant, how the constrained fitting algorithm enforces physical consistency, and how the complete framework produces a PCIe Gen 4 jitter budget from oscilloscope measurements.

### Core Concepts

---

#### Why the Gaussian Tail Is Predictable from the Center

Inside the silicon and copper interconnect, thermal energy causes electrons to vibrate chaotically. Each vibration generates a microscopic voltage pulse. At any given picosecond during transmission, the total noise voltage on the wire is the sum of billions of these individual pulses.

Most of the time, roughly half the electrons experience a positive thermal perturbation and half experience a negative perturbation. Their voltage contributions cancel. This high physical likelihood of cancellation creates the massive peak at the center of the bell curve, representing near-zero net noise voltage. For a large voltage spike to occur, a massive majority of those billions of electrons must coincidentally vibrate in the same direction at the same instant. The laws of entropy dictate that this level of spontaneous coordination is extraordinarily improbable.

This physical reality produces the Gaussian shape. The center peak exists because cancellation is highly probable. The tails decay exponentially because large-scale coordination of billions of independent electrons is highly improbable. The critical insight is that the physics generating the dense center and the physics generating the sparse tails are identical. There is no change in mechanism at any point along the curve. A single continuous mathematical equation (  $e^{-x^2/2\sigma^2}$  ) describes the entire spectrum from center to infinity.

Once the oscilloscope measures the visible central mass of the histogram and calculates the standard deviation (  $\sigma$  ), all variables in that equation are locked. The shape of the unmeasured tail at  $7\sigma$  or  $10\sigma$  is not an extrapolation in the sense of projecting a trend; it is a consequence of the same thermodynamic process that produced the measured data. The tail's trajectory is mathematically bound to the center mass because both regions arise from the same physical source.

## Standard Deviation as a Geometric Coordinate

The standard deviation (  $\sigma$  ) is not an abstract statistical summary. It is a physical distance on the horizontal axis of the TIE histogram, measured in picoseconds.

The Gaussian bell curve starts flat at its peak, accelerates steeply downward, and then begins to level out as it approaches the tail. The exact coordinate where the curve stops accelerating downward and begins to level out is the inflection point, the transition from concave-down to concave-up curvature. The horizontal distance from the center peak to this inflection point is exactly one standard deviation.

The mathematics of the Gaussian function lock the area between  $-1\sigma$  and  $+1\sigma$  at precisely 68.2689% of the total distribution. This is not an engineering approximation. It is a fundamental mathematical constant derived by integrating  $e^{-x^2/2}$  between those inflection boundaries, in the same way that  $\pi$  is derived from the geometry of a circle. The 68.2% proportion holds regardless of whether  $\sigma$  is 0.5 picoseconds or 50 picoseconds. It is a property of the Gaussian function itself, not of any particular hardware.

This geometric rigidity is what makes the Q-factor framework possible. Measuring  $\sigma$  at the inflection point fixes the entire curve. The area under the tail at  $7\sigma$  , at  $10\sigma$  , or at any other multiple is a calculable consequence of that single measured parameter.

## Bit Error Rate as Tail Area

A bit error does not occur at a single specific jitter coordinate. A bit error occurs when an edge shifts anywhere past the receiver's sampling threshold.

If the receiver samples data exactly 31.25 picoseconds after the ideal edge crossing (half the Unit Interval for PCIe Gen 4), then an edge delayed by 32 picoseconds causes a bit error. An edge delayed by 50 picoseconds also causes a bit error. The probability of a bit error is therefore the total probability that the edge arrives anywhere from 31.25 ps outward to infinity.

Reading the probability density at the 31.25 ps coordinate alone would report only the likelihood of landing in that specific infinitesimal time bin. The actual BER requires summing the probability across all bins from 31.25 ps to infinity. In calculus, summing a continuous curve over a range is an integral. The Bit Error Rate is mathematically identical to the area under the Gaussian tail beyond the sampling threshold.

## The Complementary Error Function (erfc) and Its Inverse

The Gaussian probability density is defined by the exponential function  $e^{-x^2}$  . Integrating this function from a boundary point  $x$  out to infinity cannot be solved using elementary algebra. Mathematicians created the complementary error function (erfc) to represent this specific integral.

The erfc computes the area under the tail of a Gaussian distribution beyond a specified boundary. It answers the question: if a random variable is drawn from a Gaussian distribution, what is the probability that it falls beyond  $x$  standard deviations from the mean?

In jitter analysis, the problem runs in the opposite direction. The standard (PCIe, IEEE, Fibre Channel) dictates the required BER. The area is known. What is unknown is how far the physical timing jitter must spread before the probability drops to that BER level. The problem is solving for the x-axis boundary, not the y-axis probability.

This requires the inverse erfc. An inverse function is not a reciprocal (  $1/\text{function}$  ); it is a mathematical reversal of inputs and outputs. The forward erfc takes a boundary coordinate as input and outputs the tail area. The inverse erfc takes a tail area as input and outputs the boundary coordinate. The relationship mirrors basic geometry: the forward function  $A = \pi r^2$  computes area from radius; the inverse function  $r = \sqrt{A/\pi}$  computes radius from area.

The inverse erfc receives the target BER (such as  $10^{-12}$  ) and returns the exact number of standard deviations from the mean where the tail area equals that probability. This number of standard deviations is the Q-factor.

## The Q-Factor

Every physical system has a different amount of thermal noise. One oscillator's Gaussian distribution might be very narrow (low noise), while another's is very wide (high noise). The erfc cannot output a result in picoseconds because it operates in dimensionless mathematical space, normalized to standard deviations.

The Q-factor is the output of the inverse erfc: the boundary coordinate measured in units of  $\sigma$  . For a target BER of  $10^{-12}$  , the inverse erfc returns  $Q = 7.03$  . This means that for any Gaussian distribution in the universe, regardless of its physical scale, the tail area beyond  $7.03\sigma$  from the mean contains exactly  $10^{-12}$  of the total probability.

Common Q-factor values across standard BER targets:

| Target BER | Q-Factor | Peak-to-Peak Multiplier ( $2Q$ ) |
|------------|----------|----------------------------------|
| $10^{-9}$  | 6.00     | 12.00                            |
| $10^{-12}$ | 7.03     | 14.06                            |
| $10^{-15}$ | 7.94     | 15.88                            |

To convert the universal Q-factor into a physical time boundary for a specific piece of hardware, the Q-factor is multiplied by the measured  $RJ_{\text{rms}}$  of that hardware:

$$\text{Physical Boundary (ps)} = Q \cdot RJ_{\text{rms}}$$

This multiplication scales the dimensionless mathematical model to the physical reality of the silicon. The result is the distance from the center of the RJ distribution to the point where the tail probability equals the target BER.

The factor of 2 in the Total Jitter formula arises because an edge can shift either too early (the left tail) or too late (the right tail).  $Q \cdot RJ_{\text{rms}}$  measures the distance from center to one boundary. The total span of possible arrival times destroyed by random jitter is twice that distance:

$$RJ_{\text{p-p}} = 2 \cdot Q \cdot RJ_{\text{rms}}$$

## The Dual-Dirac Model Structure

The measured TIE histogram is the convolution of the Deterministic Jitter distribution with the Random Jitter Gaussian. The actual DJ distribution is complex: it contains discrete spectral lines for PJ and DCD, and bounded multi-modal shapes for ISI. Deconvolving the pure RJ Gaussian from this entangled composite is computationally intensive and highly susceptible to noise.

The Dual-Dirac model replaces the complex DJ distribution with two Dirac delta impulses separated by a distance  $DJ_{\delta\delta}$ . The total jitter histogram is then modeled as the convolution of these two deltas with a single Gaussian, producing a bimodal distribution: two overlapping bell curves, one anchored at the left DJ boundary ( $\mu_L$ ) and one anchored at the right DJ boundary ( $\mu_R$ ).

This is a worst-case simplifying assumption adopted by industry standards. Concentrating all DJ energy at two extreme points produces the widest possible convolution with the RJ Gaussian, yielding a conservative (overestimating) total jitter number. The mathematical convenience lies in reducing the DJ description from an arbitrary complex shape to two scalar parameters: the separation distance ( $DJ_{\delta\delta}$ ) and the population split between the two impulses.

## Asymmetric Delta Placement and Population Weighting

The two Dirac deltas do not need to be equidistant from the ideal zero-crossing, and they do not need to carry equal populations of hits.

Random Jitter is guaranteed to follow a perfectly symmetric Gaussian. The thermodynamic processes that generate it possess no directional bias. An electron is exactly as likely to experience a positive perturbation as a negative one.

Deterministic Jitter, by contrast, can be profoundly asymmetric. DCD from a stronger pull-up transistor concentrates rising-edge crossings early and falling-edge crossings late. ISI from a specific via reflection may interfere aggressively with one bit pattern while leaving the complementary pattern unaffected. These hardware-specific directional biases create lopsided timing distributions.

The Dual-Dirac model accommodates this asymmetry by allowing the left and right impulses to sit at unequal distances from center and to carry unequal fractions of the total population. If DCD causes 80% of edges to cross early and 20% to cross late, the left Gaussian population contains 800,000 hits out of a million, and the right Gaussian contains 200,000. The left curve appears four times taller, but both curves share the same horizontal width ( $\sigma$ ). The height difference reflects the population weighting multiplier (0.8 vs. 0.2), not a difference in noise magnitude.

The 68.2% rule applies locally to each sub-population. Of the 800,000 edges in the left group, exactly 68.2% fall within  $\pm 1\sigma$  of the left mean. Of the 200,000 edges in the right group, exactly 68.2% fall within  $\pm 1\sigma$  of the right mean.

## The Single-Sigma Constraint

Random Jitter originates from thermal noise in the silicon. The silicon possesses a single temperature and a single fundamental noise floor. The thermodynamic variance of the noise is identical for all edges, regardless of whether DCD or ISI has shifted the mean crossing time of a particular sub-population.

The Dual-Dirac model enforces this physical reality by requiring both Gaussian curves to share the same standard deviation. The model does not use two different Gaussians with independent widths. It uses the same Gaussian, duplicated and anchored at two different mean positions. The vertical amplitudes (population weights) may differ substantially, but the horizontal shapes (  $\sigma$  ) are mathematically identical.

In practice, independently fitting the left tail and the right tail will yield slightly different  $\sigma$  values due to finite sample sizes, instrument noise, or differing slew rates between rising and falling edges (which slightly alter the amplitude-to-phase conversion ratio). Oscilloscope algorithms such as Keysight EZJIT resolve this discrepancy through a constrained least-squares fit. The software applies a mathematical constraint that  $\sigma_{\text{left}}$  must equal  $\sigma_{\text{right}}$  and processes both tails simultaneously to find the single unified  $RJ_{\text{rms}}$  that minimizes the total fitting error across the entire histogram.

This single unified  $RJ_{\text{rms}}$  represents the true thermal baseline of the hardware. It is this number that the Q-factor scales to produce the Total Jitter budget.

## Worked Examples

---

### PCIe Gen 4 Jitter Budget

Consider a PCIe Gen 4 link operating at 16 GT/s, measured with a Keysight Infiniium oscilloscope and EZJIT+ analysis software.

#### Step 1: Establish the time budget.

At 16 GT/s, the Unit Interval (UI) is the physical time duration of a single bit:

$$UI = \frac{1}{16 \times 10^9} = 62.5 \text{ ps}$$

#### Step 2: Capture and extract the Dual-Dirac parameters.

The oscilloscope captures several hundred thousand edge crossings and generates a TIE histogram. The EZJIT+ algorithm isolates the extreme tails, fits Gaussian curves using the single-sigma constraint, and extracts:

- $DJ_{\delta\delta} = 12.0$  ps (the separation between the two Gaussian means)
- $RJ_{\text{rms}} = 1.5$  ps (the unified standard deviation from the constrained fit)

#### Step 3: Define the compliance target.

The PCI Express specification mandates compliance at  $BER = 10^{-12}$ . The corresponding Q-factor is 7.03.

#### Step 4: Calculate peak-to-peak Random Jitter.

$$RJ_{p-p} = 2 \cdot Q \cdot RJ_{rms} = 2 \times 7.03 \times 1.5 \text{ ps} = 21.09 \text{ ps}$$

**Step 5: Calculate Total Jitter at  $10^{-12}$  BER.**

$$TJ = DJ_{\delta\delta} + RJ_{p-p} = 12.0 \text{ ps} + 21.09 \text{ ps} = 33.09 \text{ ps}$$

**Step 6: Evaluate the timing margin.**

The total jitter consumes 33.09 ps out of the 62.5 ps Unit Interval, approximately 53% of the available bit period. The remaining 29.41 ps (47% of the UI) is the horizontal eye opening, the time window within which the receiver must successfully sample the data.

## Gaussian Tail Fitting: The $3\sigma$ -to- $5\sigma$ Region

The Dual-Dirac algorithm does not fit a Gaussian to the entire histogram. It fits exclusively to the far tails, typically in the region between  $3\sigma$  and  $5\sigma$  from each Gaussian mean.

The rationale is physical. DJ is bounded, so it has absolute maximum limits. Every edge that lands beyond the DJ boundary arrived there solely because unbounded Random Jitter pushed it there. The tails beyond  $3\sigma$  are well outside the DJ contamination zone and reflect pure RJ behavior.

The practical challenge is that events in this region are rare. A  $4\sigma$  event occurs roughly 6 times per 100,000 bits transmitted. To populate the fitting region with enough data points for a stable regression, the oscilloscope must capture a large total sample. Modern Dual-Dirac implementations typically require 100,000 to 1,000,000 edge crossings. The algorithm sacrifices the overwhelming majority of data points (those in the center, where DJ and RJ are entangled) and builds the entire extrapolation from the sparse but physically clean tail region.

Once the curve is fit to the visible tail edges between  $3\sigma$  and  $5\sigma$ , the Gaussian equation extends the curve to  $7\sigma$ ,  $10\sigma$ , or any other multiple with mathematical certainty, because the curve shape beyond  $5\sigma$  is fully determined by the  $\sigma$  measured within the visible portion.

## Architecture

---

### The TJ Extrapolation Formula

The complete Total Jitter extrapolation formula encapsulates the entire Dual-Dirac framework:

$$TJ(BER) = DJ_{\delta\delta} + 2 \cdot Q(BER) \cdot RJ_{rms}$$

Each term has a direct physical interpretation:

- $DJ_{\delta\delta}$  is the deterministic floor: the irreducible timing spread caused by bounded hardware mechanisms (ISI, DCD, PJ), modeled as the separation between two Dirac delta impulses.
- $RJ_{rms}$  is the thermal baseline: the standard deviation of the Gaussian noise contribution, measured as a single unified value across both tails via constrained least-squares fitting.
- $Q(BER)$  is the statistical amplifier: the number of standard deviations required to reach the target error probability, computed as the inverse of the complementary error function.
- The factor of 2 accounts for both the early and late tails of the distribution.

The formula is additive because it combines a bounded deterministic component with the peak-to-peak extent of the unbounded Gaussian at the target confidence level. It produces a single number in picoseconds that can be compared directly against the Unit Interval to determine pass/fail compliance.

## From Histogram to Compliance: The Algorithmic Pipeline

The EZJIT+ software executes the following sequence to produce a TJ number at the target BER:

1. **Acquire TIE values** for several hundred thousand edges (input from the measurement infrastructure described in [Jitter Decomposition and Measurement](#)).
2. **Construct the histogram** and normalize to a probability density function.
3. **Identify the tail regions** beyond the DJ bounds on both the left and right extremes.
4. **Fit Gaussian curves** to both tails simultaneously under the single-sigma constraint ( $\sigma_L = \sigma_R$ ), solving for the unified  $RJ_{rms}$  and the two mean positions ( $\mu_L, \mu_R$ ).
5. **Compute  $DJ_{\delta\delta}$**  as  $\mu_R - \mu_L$ .
6. **Look up or compute the Q-factor** for the target BER using the inverse erfc.
7. **Apply the TJ formula** to produce the final compliance number.

The entire pipeline transforms a physical histogram into a single picosecond value that determines whether the link meets its specification.

## Edge Cases

---

### The Overestimation Tendency of Dual-Dirac

The Dual-Dirac model concentrates all DJ energy at exactly two points, creating the widest possible convolution footprint with the RJ Gaussian. In physical hardware, the DJ distribution is typically spread across many intermediate values (ISI creates a range of delays for different bit patterns, not just two extremes). The actual DJ distribution, when convolved with RJ, produces a total jitter that is narrower than the Dual-Dirac prediction.

This makes the Dual-Dirac model inherently conservative. It will report a TJ number that equals or exceeds the true total jitter, providing a safety margin. Standards bodies adopted this model precisely because its overestimation tendency means that hardware passing the Dual-Dirac compliance test is guaranteed to meet the BER target under real operating conditions.

The conservatism becomes problematic only when the overestimation is so large that functional hardware fails the specification. In practice, this is rare: the separation between the Dual-Dirac TJ and the true TJ is typically a few picoseconds, well within the design margin of links that are close to the performance boundary.

## The Single-Temperature Assumption

The single-sigma constraint assumes that the silicon operates at a uniform temperature, producing a single noise floor for all edges. In practice, power dissipation creates thermal gradients across the die. Regions near high-activity logic may run hotter than quiet regions, producing locally elevated noise. If the transmitter's output stage experiences significant temperature variation during a measurement window (as might occur during thermal transients at power-up), the left and right tails of the histogram could genuinely have slightly different intrinsic  $\sigma$  values.

The constrained fit averages this variation into a single number, which remains a good approximation as long as the thermal gradient is small relative to the overall noise floor. For links where thermal transient effects are suspected, capturing data after thermal equilibrium is established eliminates this concern.

## Sample Size and Tail Population

The accuracy of the Dual-Dirac extraction depends entirely on populating the  $3\sigma$  -to-  $5\sigma$  tail region with enough data points to constrain the Gaussian fit. With too few points, the regression is underdetermined and the extracted  $\sigma$  is noisy.

In a system with only Random Jitter, a few thousand edges would suffice to determine  $\sigma$  accurately, because the central mass directly reveals the standard deviation. In a real system contaminated with DJ, the algorithm cannot use the central mass (it is corrupted by the DJ distribution) and must rely exclusively on the sparse tail population. This is why the sample size requirement jumps from thousands to hundreds of thousands: the vast majority of captured edges are discarded by the algorithm, and only the rare tail events contribute to the fit.

Insufficient sample size manifests as measurement-to-measurement variation in the reported  $RJ_{\text{rms}}$  and TJ values. If repeated acquisitions produce inconsistent results, the first diagnostic step is to increase the capture depth.

## Chapter 12: Test Patterns and Jitter Isolation

---

The jitter decomposition framework described in [Jitter Decomposition and Measurement](#) establishes that total jitter separates into Random Jitter and Deterministic Jitter, and that Deterministic Jitter subdivides further into Data Dependent Jitter and Periodic Jitter. The [Dual-Dirac Model and BER Extrapolation](#) provides the statistical machinery to project these components to any target BER. Both of those guides operate on the measured histogram and the extracted numerical parameters (  $RJ_{rms}$  ,  $DJ_{\delta\delta}$  ). Neither addresses the physical measurement techniques that produce the separation in the first place.

This guide covers the practical machinery of jitter isolation: how test patterns are selected to make specific jitter components visible and measurable, how synchronous averaging extracts Data Dependent Jitter from the raw timing data, how FFT spectral decomposition identifies Periodic Jitter sources, how DCD and ISI are separated within the DDJ measurement, and how specialized patterns target specific hardware failure modes. The techniques described here correspond directly to steps 4 through 7 of the seven-step measurement flow introduced in the Architecture section of [Jitter Decomposition and Measurement](#).

### Core Concepts

---

#### From Histogram to Timeline: The TIE Track

The TIE histogram used by the Dual-Dirac model collapses all edge crossings into a probability distribution. This representation is powerful for extracting  $RJ_{rms}$  and  $DJ_{\delta\delta}$  through tail fitting, but it completely destroys the chronological order of the edge crossings. Decomposing Deterministic Jitter into its sub-components requires a fundamentally different data representation.

The TIE Track is a chronological sequence of Time Interval Error values plotted continuously against time. Each data point records the timing error of a specific edge crossing in the order it was captured. The horizontal axis is time (or equivalently, bit position within the capture window), and the vertical axis is the TIE value in picoseconds. The TIE Track preserves the temporal structure that the histogram discards: which edges were early, which were late, and how the timing errors evolved over millions of consecutive bit transitions.

The TIE Track is the raw material from which synchronous averaging and FFT spectral analysis extract the individual DJ components.

#### Why Test Patterns Must Be Pseudo-Random

A live communication link transmits scrambled data that possesses a mathematically random distribution. The transmitted bit sequences contain a complex mix of high-frequency alternating patterns and low-frequency long bit runs that trigger different physical responses in the channel. To accurately characterize a channel's jitter behavior, the test stimulus must replicate this broadband spectral content and excite all possible Intersymbol Interference conditions.

The test pattern cannot be truly random. Synchronous averaging, the core technique for extracting Data Dependent Jitter, requires the oscilloscope to capture thousands of identical repetitions of the same pattern and average the timing error at each bit position across all repetitions. A truly random sequence never repeats, making alignment and averaging impossible.

A Pseudo-Random Bit Sequence resolves this conflict. PRBS patterns provide the complex frequency spectrum of live data while maintaining the exact, repeating mathematical structure required by the averaging algorithm.

## PRBS Pattern Generation and Structure

PRBS patterns are generated by Linear Feedback Shift Registers (LFSRs). A PRBS-N pattern has a length of  $2^N - 1$  bits and contains every possible N-bit sequence except the all-zeros state. The pattern repeats identically after  $2^N - 1$  bits, providing the deterministic structure needed for synchronous averaging.

The maximum run length of consecutive identical bits within a PRBS-N pattern is exactly N bits. A PRBS-7 pattern contains 127 bits, and its longest run of identical bits is seven (such as 1111111). A PRBS-15 pattern contains 32,767 bits with a maximum run of fifteen. A PRBS-31 pattern contains approximately 2.1 billion bits with a maximum run of thirty-one.

The maximum run length is the critical parameter for ISI testing. Long runs of identical bits allow the channel's distributed capacitance to charge or discharge fully, establishing the maximum voltage swing. When the pattern transitions out of a long run, the signal must traverse the full voltage distance to the opposite rail, exposing the worst-case ISI condition. Longer maximum runs create more severe ISI stress.

## PRBS-7: The Workhorse Pattern

PRBS-7 is the most widely used pattern in jitter characterization for two reasons.

It is short enough to allow an oscilloscope to capture tens of thousands of complete repetitions within a standard memory buffer. A 127-bit pattern repeating 10,000 times requires only 1.27 million bits of capture depth. This accumulation provides the statistical depth needed for synchronous averaging to suppress RJ and PJ by a factor of  $\sqrt{N}$ , where N is the number of averaged repetitions.

Its maximum run length of seven directly stresses the encoding boundaries of foundational high-speed standards. PCIe Gen 1, SATA, and Gigabit Ethernet use 8b/10b encoding, which mathematically prevents any data run longer than five identical consecutive bits. A PRBS-7 pattern with its run length of seven exceeds this limit, providing a necessary physical guardband that proves the system survives worst-case ISI conditions across all temperature variations, voltage droops, and manufacturing tolerances.

Modern standards using 128b/130b encoding (PCIe Gen 3 and beyond) permit much longer identical bit sequences. Validating these links requires PRBS-15 or PRBS-31 patterns to simulate worst-case ISI channel charging at low frequencies. The tradeoff is measurement time: a PRBS-31 pattern requires billions of bits per repetition, and accumulating thousands of repetitions for synchronous averaging demands deep oscilloscope memory and long acquisition windows.

## Why PRBS-5 Is Avoided

PRBS-5 contains only 31 bits before repeating. While it does include a run length of five, a 31-bit sequence lacks the spectral density required to simulate live data. Real communication traffic contains a massive variety of transition histories and frequency combinations that excite different physical resonances in the printed circuit board traces. A repeating 31-bit loop is too simple to trigger those complex channel interactions.

The practical barrier is equally significant. The LFSR polynomial required to generate PRBS-7 is physically hardcoded into the silicon of nearly every commercial Bit Error Ratio Tester and internal SerDes test block. PRBS-5 is rarely implemented in hardware generation logic, making it impractical for standardized laboratory testing.

## Worked Examples

---

### Synchronous Averaging for DDJ Extraction

Synchronous averaging is the technique that separates Data Dependent Jitter from all uncorrelated jitter sources. The procedure exploits the defining property of DDJ: it is correlated to the bit pattern, meaning the same bit sequence produces the same timing error on every repetition.

**Step 1: Acquire multiple pattern repetitions.** The hardware under test transmits a repeating PRBS-7 pattern (127 bits per repetition). The oscilloscope captures thousands of repetitions of the complete pattern and measures the TIE at every edge transition.

**Step 2: Align and average.** The software identifies the pattern boundary and mathematically aligns every captured repetition. For each bit position within the pattern (bit 1, bit 2, through bit 127), the algorithm averages the TIE values recorded across all repetitions.

**Step 3: Observe the cancellation.** Random Jitter and Periodic Jitter are entirely uncorrelated to the bit sequence. At any given bit position, the RJ contribution is equally likely to be positive or negative on successive repetitions. Averaging thousands of these positive and negative swings drives their contribution toward zero. After  $N$  averages, the uncorrelated jitter amplitude reduces by a factor of  $\sqrt{N}$ . With 10,000 repetitions, uncorrelated jitter is suppressed by a factor of 100.

**Step 4: Isolate the DDJ.** The only timing error that survives the averaging process is the jitter physically caused by the data pattern itself. If bit position 47 consistently arrives 2.3 ps late (because it follows a long run of zeros and the channel's parasitic capacitance requires extra charging time), that 2.3 ps offset appears in the averaged waveform regardless of how many repetitions are captured. This isolated remainder is pure Data Dependent Jitter, encompassing both Intersymbol Interference and Duty Cycle Distortion.

### DCD and ISI Separation

Once the software isolates the total DDJ track, separating DCD from ISI is a sorting operation based on edge polarity.

Duty Cycle Distortion is caused by a physical asymmetry between rising edges and falling edges. The oscilloscope software calculates the average timing error of every rising edge in the pattern and compares it to the average timing error of every falling edge. If rising edges cross the threshold consistently earlier than falling edges (as occurs when the pull-up transistor has stronger drive strength than the pull-down transistor), the software extracts that static time delta and classifies it as DCD.

After subtracting the DCD baseline, the remaining timing variations fluctuate based strictly on the historical sequence of bits leading up to each transition. This residual variance is classified entirely as Intersymbol Interference.

## FFT Spectral Decomposition for Periodic Jitter

The oscilloscope takes the isolated DDJ values and subtracts them from the original, raw TIE Track. The remaining data track now contains strictly uncorrelated jitter: a mixture of Random Jitter and Periodic Jitter.

To separate them, the oscilloscope executes a Fast Fourier Transform on the remaining track, converting the timing data from the time domain into the frequency domain.

Random Jitter is broadband white noise. Its energy spreads evenly across all frequencies, creating a flat, low-level noise floor at the bottom of the FFT graph. This flat floor is the spectral signature of the Gaussian process described in [Jitter Decomposition and Measurement](#): billions of independent microscopic thermal events producing a featureless distribution of timing perturbations.

Periodic Jitter is driven by specific oscillating sources. A 33 MHz switching power supply, a 100 MHz adjacent clock trace, or a ringing PLL each produces a sinusoidal modulation of the edge timing at its characteristic frequency. In the frequency domain, these periodic interferences appear as sharp, distinct spectral spikes rising well above the RJ noise floor.

The algorithm identifies each spike that exceeds the noise floor by a specified margin (typically 6 to 10 dB), records its frequency and amplitude, and quantifies the Periodic Jitter contribution. The frequency of each spike directly identifies the interference source, transforming an opaque timing problem into an actionable hardware diagnostic.

## Architecture

---

### Specialized Jitter Test Patterns

Standard PRBS patterns exercise the full spectral range of a channel, but specific hardware failure modes are more efficiently exposed by purpose-built patterns.

**Clock Pattern (101010...):** This pattern operates exclusively at the highest fundamental frequency of the link (the Nyquist rate). It contains no long consecutive bit runs, so the channel's distributed capacitance never fully charges or discharges. ISI is effectively eliminated. Engineers transmit a clock pattern to isolate and measure pure Duty Cycle Distortion or to evaluate the Random Jitter floor without channel-dependent interference. The measurement reveals only the transmitter's intrinsic timing asymmetry and the system's thermal noise baseline.

**CJTPAT (Continuous Jitter Test Pattern):** This pattern is designed to stress the receiver's Clock and Data Recovery circuit. It forces abrupt transitions between extremely high-frequency alternating bits and maximum-length low-frequency bit holds. The sudden shifts in frequency content challenge the CDR's Phase-Locked Loop to maintain phase lock. CJTPAT is used in receiver compliance testing to verify that the CDR can track the data under maximum stress without drifting or losing lock.

**Low-Frequency Square Wave (e.g., 64 ones followed by 64 zeros):** This pattern forces the transmitter to hold a static voltage state for an extended duration. It exposes power distribution network flaws by revealing vertical voltage droop that occurs when the transmitter slowly drains the local decoupling capacitors during a long sequence of identical bits. The pattern is diagnostic for supply integrity, not for timing jitter.

## The Complete Measurement Flow: Steps 4 Through 7

The seven-step jitter measurement flow begins with signal acquisition, clock recovery, and TIE computation (steps 1 through 3, covered in [Jitter Decomposition and Measurement](#)). Steps 4 through 7 apply the techniques described in this guide to decompose the measured jitter into its root causes.

**Step 4: Pattern Synchronization.** The oscilloscope locks to the repeating test pattern to enable synchronous averaging. This requires knowledge of the pattern type and length. The software identifies pattern boundaries by correlating the captured data with the known PRBS polynomial.

**Step 5: Deterministic Jitter Extraction.** Synchronous averaging separates DDJ from all uncorrelated jitter. The averaged TIE Track isolates the pattern-correlated timing errors. DCD is separated from ISI by comparing rising-edge and falling-edge averages. The remaining uncorrelated jitter residual (original TIE Track minus DDJ) contains only RJ and PJ.

**Step 6: Spectral Decomposition.** FFT analysis of the uncorrelated residual separates Periodic Jitter (sharp spectral spikes at identifiable frequencies) from Random Jitter (flat broadband noise floor). Each PJ spike is cataloged with its frequency and amplitude.

**Step 7: Statistical Validation.** The decomposition is validated by confirming internal consistency. The total jitter reconstructed from decomposed components (  $DJ_{p-p} + 2Q \cdot RJ_{rms}$  ) should match the directly measured total jitter within measurement uncertainty. The RJ histogram should be approximately Gaussian (verified by normality tests). DJ should be bounded and repeatable across successive measurements. Where the observation window is long enough to produce a measurable BER, the extrapolated BER from the Dual-Dirac model should agree with the directly counted error rate.

## Edge Cases

---

### Pattern Length vs. Averaging Depth Tradeoff

Longer PRBS patterns provide richer spectral content and exercise more ISI conditions, but they require proportionally more measurement time to achieve the same averaging depth. A PRBS-7 pattern at 10 Gbps repeats approximately 79 million times per second. Accumulating 10,000 averages takes a fraction of a millisecond. A PRBS-31 pattern at the same data rate repeats approximately 4.7 times per second. Accumulating 10,000 averages would require over 35 minutes of continuous acquisition.

The practical consequence is that PRBS-31 measurements typically use fewer averaging repetitions, which reduces the suppression of uncorrelated jitter in the DDJ extraction. The residual RJ contamination in the DDJ measurement is higher, reducing the accuracy of the decomposition. Engineers must balance the need for realistic spectral content (longer patterns) against the need for clean DDJ isolation (more averages of shorter patterns).

## Measurement Floor and Instrument Noise

The oscilloscope itself contributes jitter to every measurement. Trigger jitter, sampling clock phase noise, and analog front-end noise all add to the measured TIE values. These instrument contributions set a measurement floor below which the device under test's jitter cannot be resolved.

For Random Jitter measurements, the instrument's own  $RJ_{\text{rms}}$  adds in quadrature with the DUT's  $RJ_{\text{rms}}$  :

$$RJ_{\text{measured}} = \sqrt{RJ_{\text{DUT}}^2 + RJ_{\text{instrument}}^2}$$

If the instrument floor is 200 femtoseconds and the DUT's RJ is 1.5 picoseconds, the instrument contribution is negligible (less than 1% error). If the DUT's RJ is 300 femtoseconds, the instrument floor dominates and the measurement becomes unreliable.

## CDR Bandwidth and Its Effect on Jitter Measurement

The Clock and Data Recovery bandwidth setting critically affects which jitter components are captured. The CDR acts as a high-pass filter on jitter: timing variations slower than the CDR bandwidth are tracked by the recovered clock (and therefore invisible to the TIE measurement), while variations faster than the CDR bandwidth appear as measured jitter.

A narrow CDR bandwidth tracks only very slow drift, exposing virtually all jitter components to measurement. A wide CDR bandwidth tracks moderate-frequency variations, effectively filtering out low-frequency PJ components. The CDR bandwidth must match the application's specification to produce compliance-relevant results. A measurement taken with a different CDR bandwidth than the standard requires is not comparable to the specification limit.

PART 4

# Equalization and RX Architecture

## Chapter 13: Transmitter Feed-Forward Equalization

---

Copper traces behave as physical low-pass filters. The skin effect and dielectric absorption progressively strip high-frequency energy from a propagating wavefront, rounding the sharp edges of digital transitions and smearing the voltage of one bit into the time slots of its neighbors. This inter-symbol interference (ISI) degrades the voltage margin and timing margin at the receiver, closing the eye diagram and driving the bit error rate toward failure.

Feed-forward equalization (FFE) attacks this problem at the source. Rather than waiting for the receiver to reconstruct a damaged waveform, the transmitter intentionally pre-distorts the signal before launching it onto the channel. The circuit boosts the voltage of high-frequency transitions and suppresses the amplitude of low-frequency static sequences, producing a waveform that looks severely distorted at the transmitter pad but arrives at the receiver with a restored frequency balance. The channel loss and the transmitter pre-distortion are designed to cancel each other, yielding a clean, open eye at the far end.

The mathematical engine behind this pre-distortion is a finite impulse response (FIR) filter implemented as a digital shift register operating on adjacent bits in the time domain. The filter does not explicitly separate the signal into frequency bands the way an analog equalizer does. It achieves the equivalent frequency shaping indirectly, by adjusting the output voltage of each bit based on the values of its immediate neighbors. The static tap weights that control these adjustments must be calibrated to the specific physical channel, and modern high-speed protocols use an active negotiation process called link training to optimize them in real time.

### Core Concepts

---

#### The FIR Shift Register

The transmitter implements feed-forward equalization using a digital shift register that holds a sequential queue of bits sliding through time. At every clock cycle, the filter simultaneously reads three specific positions in this queue: the current bit, the immediately preceding bit, and the immediately following bit. Each position is assigned a static mathematical multiplier called a tap weight. The filter multiplies each bit value by its assigned weight, sums the three products, and drives the resulting voltage onto the transmission line.

The three taps correspond to three distinct physical roles:

**Main cursor.** The main cursor tap provides the primary baseline voltage drive for the current bit. Its weight is the largest of the three, typically normalized to 1.0, and it sets the overall signal amplitude.

**Post-cursor.** The post-cursor tap reads the value of the bit that just finished transmitting. A physical transmission line stores electromagnetic energy from the previous bit, and that energy bleeds forward into the current time slot as trailing ISI. The post-cursor tap subtracts a calculated fraction of the previous bit's contribution to cancel this lingering energy.

**Pre-cursor.** The pre-cursor tap reads the value of the bit that is about to transmit next. Energy can also bleed backward in time, creating leading ISI that arrives slightly before the main wavefront of the upcoming bit. The pre-cursor tap adjusts the current bit's voltage to cancel this early interference from the approaching transition.

## How Time-Domain Tap Math Creates Frequency Shaping

Adjusting voltages based on neighboring bit values produces the same mathematical effect as applying a high-pass frequency filter, but the FIR circuit operates entirely in the time domain.

Long sequences of identical bits (such as 11111 or 00000) represent the lowest-frequency content on the link. The signal is static, and no transitions occur. The post-cursor and pre-cursor taps both detect matching neighbors and subtract energy from the main cursor output, reducing the total amplitude. This amplitude reduction during static sequences is physically equivalent to attenuating low-frequency content.

Rapid alternating patterns (such as 101010) represent the highest-frequency content. At every bit transition, the pre-cursor and post-cursor taps detect opposite-polarity neighbors. The negative tap weights, multiplied by opposite-sign bit values, produce positive contributions that add to the main cursor voltage rather than subtracting from it. The output amplitude at transitions is preserved at full strength while static sequences are suppressed. The resulting frequency response preferentially preserves high-frequency edge content and attenuates low-frequency runs, which is exactly the inverse of the channel's low-pass loss characteristic.

## The Tap Weight Formula

The transmitter assigns mathematical values to binary states: a logic 1 equals  $+1$ , and a logic 0 equals  $-1$ . The FIR output voltage at any given clock cycle is the weighted sum:

$$V_{out} = (c_{pre} \times b_{next}) + (c_{main} \times b_{current}) + (c_{post} \times b_{previous})$$

where  $c_{pre}$ ,  $c_{main}$ , and  $c_{post}$  are the static tap weights, and  $b_{next}$ ,  $b_{current}$ , and  $b_{previous}$  are the signed bit values ( $+1$  or  $-1$ ). The tap weights are fixed numbers programmed into the silicon, but the output voltage changes dynamically with every clock cycle as the signed bit values slide through the shift register.

## Worked Examples

---

### The "111" Sequence: De-emphasis

Consider a standard 3-tap FIR with weights:  $c_{pre} = -0.25$ ,  $c_{main} = 1.0$ ,  $c_{post} = -0.25$ .

Evaluating the middle bit of a 111 sequence: the previous bit is  $+1$ , the current bit is  $+1$ , and the next bit is  $+1$ .

$$V_{out} = (-0.25 \times 1) + (1.0 \times 1) + (-0.25 \times 1) = -0.25 + 1.0 - 0.25 = 0.5$$

The filter reduced the output voltage from a full-strength 1.0 to 0.5. A long run of identical bits risks overcharging the transmission line's distributed capacitance, saturating the dielectric and making the eventual transition sluggish. Reducing the drive voltage during static sequences actively prevents this overcharging.

### The "110" Sequence: Pre-emphasis at a Transition

Evaluating the final 1 in a 110 sequence: the previous bit is  $+1$ , the current bit is  $+1$ , and the next bit is  $-1$ .

$$V_{out} = (-0.25 \times -1) + (1.0 \times 1) + (-0.25 \times 1) = 0.25 + 1.0 - 0.25 = 1.0$$

The pre-cursor term detected an incoming opposite-polarity bit. Multiplying the negative weight by the negative bit value produced a positive contribution, boosting the voltage from 0.5 back to 1.0 in the clock cycle immediately before the transition.

Evaluating the 0 (the next clock cycle): the previous bit is  $+1$ , the current bit is  $-1$ , and the next bit is  $-1$ .

$$V_{out} = (-0.25 \times -1) + (1.0 \times -1) + (-0.25 \times 1) = 0.25 - 1.0 - 0.25 = -1.0$$

The voltage plunged from  $+1.0$  to  $-1.0$  in a single clock cycle. The completely static mathematical weights automatically generated a maximized voltage swing at the exact bit boundary where the signal needs the most energy to punch through the parasitic capacitance and establish a clean transition. The pre-emphasis boosted the bit before the transition, and the de-emphasis on the other side pulled the voltage to its extreme, producing the steepest possible edge.

### Interpreting the Pre-Distorted Waveform

The signal departing the transmitter pad looks aggressively distorted. Static bit runs are suppressed to half amplitude. Transitions are amplified to full swing. An oscilloscope probe placed at the transmitter output would display what appears to be a badly corrupted waveform with wildly varying bit amplitudes.

This appearance is intentional. The physical channel will attenuate the high-frequency transitions far more heavily than the static runs. The pre-distorted signal and the channel loss are designed to cancel. The waveform arriving at the receiver, after propagating through the full length of copper, should exhibit roughly equal amplitude for both static and transitioning bit patterns, producing a vertically symmetric and horizontally centered eye.

## Architecture

---

### Channel-Specific Tap Calibration

Every physical transmission line presents a unique parasitic profile. A two-inch top-layer trace on a high-quality laminate introduces minimal capacitance and negligible high-frequency loss. A ten-inch backplane channel traversing multiple vias and connectors accumulates massive distributed capacitance and severe dielectric absorption.

Applying the same aggressive tap weights to both channels produces opposite failure modes. The short trace suffers from over-equalization: the maximized voltage swings at transitions blow past the receiver thresholds, creating artificial noise, unnecessary reflections, and degraded timing margin. The long trace suffers from under-equalization: the voltage swing lacks sufficient energy to overcome the heavy loss, and the edges arrive too slow to trigger the receiver cleanly.

The tap weights must be tuned to match the specific physical severity of the channel. A universal default set of weights cannot account for the diversity of board layouts, trace lengths, via counts, and laminate materials encountered across different system designs.

## Tap Asymmetry and the Dielectric Wake

The pre-cursor and post-cursor weights rarely carry the same magnitude. The physical impulse response of a copper channel is temporally asymmetric: the leading edge of a pulse suffers only moderate rounding, while the trailing edge is smeared into a long, decaying tail.

This asymmetry originates from the dielectric absorption mechanism in the PCB laminate. The rising edge of the signal brings an intense, highly organized electric field that physically forces the polar molecules inside the fiberglass to snap into strict alignment with the field. This forced polarization is fast because the full power of the driving voltage actively accelerates the molecular rotation.

The falling edge of the signal removes that external force. The transmitter drops the voltage, but the molecules do not snap back to their resting state instantly. They must rely entirely on ambient thermal energy to slowly jostle out of their rigid alignment and drift back toward random orientation. Random thermal agitation is an extremely weak and slow physical process compared to a directed electromagnetic force. The molecules release their stored energy gradually back into the copper trace over a period that extends well beyond the original pulse duration, creating a trailing energy wake analogous to the wake behind a moving boat.

The receiver at the far end measures the total combined voltage arriving at any given instant. A new pulse rides directly on top of the lingering, decaying wake of the previous pulse. The old trailing energy adds to the new voltage, corrupting the receiver threshold and creating severe trailing ISI.

The pre-cursor tap corrects a short-duration effect: the slight rounding of the incoming leading edge. A relatively modest negative weight is sufficient. The post-cursor tap must combat a long, stubborn trailing wake that persists across multiple bit periods. A much stronger negative weight is required. A realistic production configuration might use a pre-cursor weight of  $-0.1$  paired with a post-cursor weight of  $-0.3$ , reflecting the physical reality that the trailing edge demands significantly more aggressive correction than the leading edge.

## The Equalization Chain

The transmitter FIR filter is the first stage in a multi-stage equalization system. It pre-compensates for channel loss before the signal enters the copper. At the receiver end, a continuous time linear equalizer (CTLE) applies analog high-frequency gain to further restore edge amplitude, and a decision feedback equalizer (DFE) uses past bit decisions to mathematically subtract the residual ISI echo from the incoming waveform. Forward error correction (FEC) provides a final layer of protection at data rates where analog equalization alone cannot achieve an acceptable raw bit error rate.

The distribution of equalization work between the transmitter FFE, the receiver CTLE, and the receiver DFE is not fixed. The optimal partitioning depends on the channel characteristics: channels dominated by smooth insertion loss respond well to CTLE, while channels with complex reflections and resonances benefit from stronger DFE correction. Link training negotiates this partitioning dynamically.

## Link Training and Production Adaptation

Engineers use S-parameter measurements during the design phase to calculate theoretical FIR tap weights. Laboratory software converts measured insertion loss data into a mathematical impulse response and derives the tap values required to flatten the channel. These calculated weights serve as starting points, but they cannot be hardcoded into production hardware.

Physical mass production introduces variation at every level. The fiberglass weave shifts slightly between board lots. The copper etching process leaves traces wider or narrower than the design target. Ambient temperature fluctuates across operating conditions. The silicon transmitter itself suffers from manufacturing variation: a weight of **1.0** programmed into one chip may produce a slightly different physical voltage than the identical weight on a chip cut from a different wafer. A rigid set of FIR values derived from one laboratory measurement on one board at one temperature will fail on a statistically significant fraction of production units.

The optimal transmitter taps also depend on the receiver's equalization state. The receiver CTLE and DFE are simultaneously adapting to the same channel, and the best FFE configuration changes depending on how aggressively the receiver is compensating on its end.

Modern high-speed protocols (PCIe Gen 4 and above, 100G+ Ethernet) solve this through link training. The transmitter and receiver boot up and exchange standardized test patterns over a low-speed backchannel. The receiver continuously measures the incoming signal quality after its own equalization. If the eye opening is insufficient, the receiver sends commands back to the transmitter requesting specific adjustments: increase the pre-cursor weight, decrease the post-cursor weight, or shift the balance between FFE and CTLE. The two devices negotiate iteratively until they converge on the optimal real-time combination for that specific physical connection at that specific operating temperature.

The training sequence itself must be a pseudo-random bit pattern (PRBS). The static tap weights are designed to work for any possible data sequence, and validating their performance requires exercising worst-case patterns. A PRBS contains every possible bit transition combination up to its pattern length, simultaneously stressing maximum-frequency alternating patterns and maximum-length static runs. This ensures the negotiated taps successfully combat both capacitive overcharging during static runs and trailing dielectric relaxation during rapid transitions.

## Edge Cases

---

### Over-Equalization and Noise Amplification

Excessive tap weights on a low-loss channel amplify the voltage swings beyond what the receiver expects. The transitions overshoot the target threshold, and the settling oscillations create artificial deterministic jitter. In extreme cases, the over-equalized edges ring past the opposite decision threshold, producing phantom crossings that the clock recovery circuit may interpret as false bit transitions.

The risk is particularly acute with strong pre-cursor taps. Pre-cursor correction acts on bits that have not yet arrived at the receiver, so errors in pre-cursor magnitude directly distort the timing of the upcoming transition rather than merely affecting the amplitude of a past echo.

### Link Training Convergence and Local Minima

Link training algorithms search a multi-dimensional parameter space (pre-cursor weight, post-cursor weight, CTLE peaking gain, DFE tap magnitudes) for the combination that maximizes the eye opening. The search landscape contains local minima: parameter combinations that appear optimal within a narrow neighborhood but are significantly worse than the true global optimum.

Training state machines can become trapped in these false optima, particularly on channels with complex reflection profiles where the relationship between tap weights and eye opening is non-monotonic. Engineers must define appropriate initial seed values and algorithmic step sizes during the silicon design phase to give the training algorithm the best chance of converging on the global optimum. If the total equalization capability designed into the silicon is insufficient for the channel loss (for example, a CTLE capable of 15 dB of restoration on a channel with 35 dB of loss), link training will fail regardless of the search algorithm, because the hardware lacks the physical capacity to close the link.

### FIR Limitations vs. DFE Strengths

The FIR filter operates blindly. It pre-distorts the signal based on static tap weights without any knowledge of the actual waveform arriving at the receiver. The FIR mixes analog voltage levels together and launches the result into the channel, hoping the pre-distortion and channel loss cancel appropriately.

The DFE at the receiver operates with a fundamentally different advantage: digital certainty. The DFE slicer makes a firm binary decision on each received bit, then uses that rigid digital value (not the messy analog voltage) to calculate the exact echo that bit will produce on subsequent time slots. The DFE replaces the analog uncertainty of the received waveform with a clean mathematical integer, enabling precise ISI cancellation that the blind FIR approach cannot match.

This architectural difference makes the FIR most effective at compensating for the smooth, predictable components of channel loss (the general downward slope of insertion loss with frequency), while the DFE excels at canceling the irregular, bit-pattern-dependent echoes that arise from reflections and resonances. The two work together as complementary halves of the equalization system.

## Chapter 14: Continuous Time Linear Equalization and Decision Feedback Equalization

---

The transmitter feed-forward equalizer pre-distorts the signal before it enters the channel, compensating for the predictable downward slope of insertion loss. The signal still arrives at the receiver degraded. The sharp voltage transitions have been rounded by skin effect resistance, the high-frequency energy has been absorbed by the dielectric, and the trailing remnants of previous bits have smeared into the current time slot. Receiver-side equalization must reconstruct the original digital information from this physically damaged waveform.

Two complementary circuits divide this reconstruction work. The continuous time linear equalizer (CTLE) is an analog filter that amplifies high-frequency signal components relative to low-frequency components, restoring the general spectral balance that the channel destroyed. The decision feedback equalizer (DFE) takes a fundamentally different approach: it makes a firm binary decision on each received bit, then uses that rigid digital certainty to calculate and subtract the exact inter-symbol interference that bit will inject into subsequent time slots.

The CTLE operates on the signal before any digital decision occurs. It treats the entire incoming waveform as an analog entity and applies a broadband frequency correction. The DFE operates after the initial digital decision, using the discrete mathematical value of past bits to perform precise echo cancellation. The two stages work in series, with the CTLE cleaning the waveform enough for the DFE slicer to make reliable initial decisions, and the DFE removing the residual bit-pattern-dependent interference that the analog filter cannot resolve.

Neither circuit determines its own operating parameters. A separate adaptation engine monitors the signal quality during a training phase and iteratively calculates the optimal settings for both the CTLE peaking gain and the DFE tap weights using the Least Mean Squares (LMS) algorithm. The physical channel, the manufacturing variation in the silicon, and the thermal environment all change from unit to unit and moment to moment. The adaptation engine tracks these changes continuously, adjusting the equalization in real time to maintain a functional link.

### Core Concepts

---

#### CTLE: Pole-Zero Frequency Shaping

A copper transmission line attenuates high-frequency components far more severely than low-frequency components, producing a frequency response that slopes downward with increasing frequency. The CTLE counteracts this slope by applying selective gain to the high-frequency portion of the spectrum. Two mathematical parameters from filter theory govern the shape of this correction.

A *zero* is a frequency at which the circuit's transfer function begins to increase. The CTLE places a zero at a frequency chosen to match the onset of significant channel loss. Below this frequency, the signal requires little correction. Above it, the CTLE gain rises progressively, amplifying the weakened high-frequency transitions to restore their amplitude relative to the low-frequency static states.

A *pole* is a frequency at which the gain curve levels off and begins to decrease. The CTLE places a pole above the signal's Nyquist frequency to cap the high-frequency boost. Without this ceiling, the amplification would continue rising into frequency ranges dominated by thermal noise and crosstalk. The CTLE would amplify noise as aggressively as signal content, degrading the signal-to-noise ratio rather than improving it.

The gain applied between the zero and the pole is called the peaking gain. A short, low-loss channel requires minimal peaking because the high-frequency content survives the transit with adequate amplitude. A long backplane channel with severe insertion loss requires aggressive peaking to compensate for the heavy attenuation of edge energy. The peaking gain is a configurable parameter negotiated during link training.

The fundamental limitation of the CTLE is that it operates as a linear analog filter. It generates a smooth, sweeping gain curve. A physical channel's insertion loss profile, however, is irregular and jagged: full of resonances from via stubs, reflections from impedance discontinuities, and the complex interactions of multiple connectors. The smooth analog curve counteracts the general downward slope of dielectric loss effectively, but it completely lacks the resolution to map and neutralize the channel's microscopic structural anomalies. Removing bit-pattern-dependent interference requires the DFE.

## DFE: Echo Cancellation Through Digital Certainty

The voltage tail that a transmitted bit leaves behind in subsequent time slots is the ISI echo. A logic 1 transmitted as a positive voltage pulse arrives at the receiver as a rounded lump with a long trailing decay. That decay does not vanish when the clock advances to the next bit period. If the receiver attempts to read a subsequent bit while the tail of the previous bit still contributes residual voltage to the wire, the receiver's measured voltage is a composite of the intended signal and the historical echo. The echo shifts the sampled voltage toward the previous bit's polarity, pushing it dangerously close to the decision threshold.

The DFE eliminates this echo by computing its exact magnitude and subtracting it from the incoming waveform before the receiver evaluates the next bit.

**The slicer.** The receiver contains a high-speed voltage comparator triggered by the clock recovery circuit at the precise center of each bit period. The slicer compares the analog voltage against a reference threshold (zero volts for a differential NRZ link) and outputs a rigid digital decision:  $+1$  for a logic 1 and  $-1$  for a logic 0. This decision is absolute and instantaneous, regardless of how noisy or distorted the analog waveform appears.

**The echo calculation.** The DFE maintains a set of calibrated tap weights, each representing the exact voltage that a single transmitted pulse contributes to a specific subsequent time slot on the particular physical channel. Tap 1 quantifies the echo bleeding from the immediately preceding bit. Tap 2 quantifies the echo from two bits ago. A modern PCIe Gen 5 receiver may carry 15 or more taps, capturing the full temporal extent of the channel's impulse response decay.

After the slicer locks in a digital decision on bit  $N$ , the DFE multiplies that decision value ( $+1$  or  $-1$ ) by the tap 1 weight to calculate the exact echo that bit  $N$  will inject into the time slot of bit  $N + 1$ . The DFE performs the same multiplication for every active tap, computing the individual echo contributions from bits  $N - 1$  through  $N - 14$  (for a 15-tap system) simultaneously.

**The summing node.** An analog subtraction circuit sits directly in front of the slicer. The DFE routes the total calculated echo into this summing node with inverted polarity. The circuit physically subtracts the predicted echo from the raw incoming analog waveform. The slicer evaluates the cleaned result rather than the corrupted original.

**Signed arithmetic and logic zeros.** High-speed differential links represent a logic 0 as a negative voltage, not as zero volts. A logic 0 produces a negative trailing echo that drags the subsequent waveform downward. The slicer assigns a mathematical value of  $-1$  to a logic 0 decision. Multiplying  $-1$  by the tap weight produces a negative echo prediction. The summing node subtracts this negative value, which is mathematically equivalent to adding voltage. The circuit pulls the incoming waveform upward to compensate for the downward drag of the previous logic 0. The signed arithmetic handles both polarities through a single unified multiplication, requiring no separate logic path for ones and zeros.

## The Power of Discarding Analog Reality

The DFE's superiority over a purely analog equalizer rests on a single architectural decision: it discards the received analog voltage and replaces it with a perfect mathematical integer.

The analog voltage of a received logic 1 varies wildly depending on the surrounding bit pattern. A logic 1 following five consecutive zeros may cross the threshold at barely  $+0.2$  V. A logic 1 following five consecutive ones may ride the accumulated charge to  $+1.2$  V. An analog equalizer attempting to calculate the echo from such variable peaks faces an impossible measurement problem.

The DFE avoids this problem entirely. The slicer examines the messy  $+0.2$  V waveform and declares: "This is a logic 1." At that instant, the DFE stops examining the physical wire. It knows the transmitter launched a consistent, pristine pulse for every logic 1 (the transmitter is a rigid machine). The channel is a linear time-invariant (LTI) system, so the echo from that standardized launched pulse is always the same, regardless of what other voltages are simultaneously present on the wire. The DFE outputs a pure  $+1$ , multiplies it by the calibrated tap weight, and routes the result to the summing node.

The entire system replaces the uncertain, noisy analog reality with a clean digital integer and a pre-calibrated physical constant. This is why the DFE achieves precise ISI cancellation that no linear analog filter can match.

## AC Coupling, DC Blocking, and Baseline Wander

High-speed serial links connect transmitters and receivers manufactured by different vendors, often operating at different common-mode DC bias voltages. An Intel CPU might bias its transmitter at 1.2 V while a Broadcom network controller biases its receiver at 0.8 V. A direct copper connection would force massive DC current from the higher-voltage chip into the lower-voltage chip, destroying the receiver's input transistors or saturating their bias points beyond functionality.

A series capacitor placed on the trace between the transmitter and receiver solves this problem. The capacitor blocks the static DC voltage difference between the two chips while allowing the high-frequency AC data transitions to pass through. This is AC coupling, and it is mandatory in PCIe, Ethernet, and USB.

**Capacitor physics.** The capacitor consists of two metal plates separated by an insulating dielectric. Electrons never physically cross the gap. When the transmitter applies a positive voltage, it strips electrons from the transmitter-side plate, creating a strong electric field across the dielectric. This field pulls electrons onto the receiver-side plate from the receiver's wiring. The resulting electron movement in the wires on both sides constitutes displacement current, and to the rest of the circuit, it appears identical to current flowing through a continuous conductor. As the receiver-side plate accumulates electrons, the electrostatic repulsion between them creates increasing back-pressure against the driving field. When the back-pressure exactly equals the driving force, all electron movement stops. Current drops to zero. The capacitor has blocked the DC component.

AC data transitions reverse the driving field before the plates reach equilibrium. The electrons slosh back and forth in the connecting wires without ever crossing the dielectric gap, maintaining continuous effective current flow.

**Baseline wander.** A long unbroken string of identical bits presents a constant voltage to the capacitor. The capacitor interprets this constant voltage as DC, charges up exponentially, and the effective current through the link decays toward zero. As the capacitor charges, the voltage on the receiver side drifts toward zero volts. The receiver's sampling threshold is calibrated against a stable baseline; the drifted baseline displaces the entire waveform, shrinking the eye opening. When the opposite bit value finally arrives, it starts from this collapsed baseline rather than the proper reference, and the receiver reads an incorrect value.

The severity of baseline wander depends on the RC time constant: the product of the capacitor value and the termination resistance. A large coupling capacitor has massive plates that require many electrons to reach equilibrium, tolerating long runs of identical bits before the voltage drifts significantly. A small capacitor fills rapidly and may cause visible wander after only a handful of identical bits. The rate at which the capacitor charges is controlled by the termination resistance through which the current flows.

High-speed protocols mitigate baseline wander through line coding. Schemes like 8b/10b encoding and PRBS scrambling guarantee an approximately equal density of ones and zeros over time, preventing the capacitor from accumulating a sustained charge in either direction. Line coding is covered in [Line Coding, FEC, and Protocol Framing](#).

## Worked Examples

---

### One-Tap DFE Echo Cancellation

Consider a minimal one-tap DFE processing two consecutive bits on a channel where a single logic 1 leaves a trailing echo of  $+0.3$  V in the next time slot.

Bit 1 arrives. The analog voltage is above zero. The slicer locks in the decision: bit 1 is a logic 1 (  $+1$  ).

Bit 2 arrives. The transmitter sent a logic 0 (physically  $-1.0$  V differential). The true signal from the transmitter physically collides with the  $+0.3$  V echo still lingering from bit 1. The raw analog voltage at the receiver is  $-1.0 + 0.3 = -0.7$  V. This voltage sits dangerously close to the zero-volt decision threshold, where a modest noise spike could cause a bit error.

The DFE intervenes before the slicer evaluates bit 2. It recalls that bit 1 was a  $+1$  and that the calibrated tap 1 weight is  $0.3$  V. The predicted echo is  $+1 \times 0.3 = +0.3$  V. The summing node subtracts this echo from the raw waveform:  $-0.7 - 0.3 = -1.0$  V. The slicer evaluates a clean  $-1.0$  V and confidently declares bit 2 is a logic 0.

If bit 1 had been a logic 0 ( $-1$ ), its negative pulse would have left a  $-0.3$  V echo dragging the subsequent waveform downward. The DFE calculates:  $-1 \times 0.3 = -0.3$  V. The summing node subtracts  $-0.3$  V, which adds  $+0.3$  V to the incoming signal, pulling the waveform upward to compensate for the downward drag. The signed arithmetic handles both polarities transparently.

## 15-Tap DFE Processing a Single Bit

Freeze time at the picosecond that bit  $N$  (the 16th bit in a sequence) arrives at a receiver with a 15-tap DFE whose adaptation engine has already converged on the final tap weights.

**Step 1: Shift register state.** The receiver maintains a memory bank containing the past 15 slicer decisions (bits  $N - 1$  through  $N - 15$ ). The stored values form a signed integer array, such as  $[+1, -1, -1, +1, +1, -1, \dots]$ .

**Step 2: Parallel echo calculation.** Fifteen digital-to-analog converters simultaneously multiply each stored decision by its assigned tap weight:

- Tap 1 reads bit  $N - 1$  ( $+1$ ), weight =  $0.2$  V. Echo contribution:  $+1 \times 0.2 = +0.2$  V.
- Tap 2 reads bit  $N - 2$  ( $-1$ ), weight =  $0.1$  V. Echo contribution:  $-1 \times 0.1 = -0.1$  V.
- Tap 15 reads bit  $N - 15$  ( $+1$ ), weight =  $0.01$  V. Echo contribution:  $+1 \times 0.01 = +0.01$  V.

**Step 3: Aggregation.** All 15 individual echo contributions feed into a single analog summing node. The total predicted echo is the algebraic sum:  $+0.2 - 0.1 + \dots + 0.01 = +0.15$  V.

**Step 4: Subtraction.** The raw analog voltage for bit  $N$  arrives at  $+0.05$  V (dangerously close to the zero-volt threshold due to the accumulated echoes from 15 prior bits). The summing node subtracts the total predicted echo:  $+0.05 - 0.15 = -0.10$  V.

**Step 5: Decision.** The slicer evaluates  $-0.10$  V. The voltage is below zero. Bit  $N$  is declared a logic 0 ( $-1$ ).

**Step 6: Shift.** The new  $-1$  decision is pushed into the first slot of the memory bank. Every other entry shifts one position. The oldest decision (bit  $N - 15$ ) is discarded. The clock advances, bit  $N + 1$  arrives, and the six-step loop repeats.

## Receiver Target Voltage and DC Insertion Loss

The DFE error slicer does not compare the received voltage against the transmitter's launch amplitude. A copper trace is a physical resistor that dissipates energy as heat even at DC. If the transmitter launches  $1.0$  V, the trace's DC resistance may absorb 40% of the voltage, leaving an absolute maximum of  $0.6$  V at the receiver. No amount of equalization can recover energy that has been permanently converted to heat.

The adaptation engine determines this baseline DC loss by monitoring the average peak amplitude of the incoming waveform over thousands of bits. If the measured average peak is **0.6 V**, the engine sets the internal reference target of the error slicer to  $\pm 0.6$  V rather than  $\pm 1.0$  V. The error signal then reflects only the damage caused by frequency-dependent loss and ISI, not the unrecoverable DC attenuation. This adaptive target voltage is sometimes referred to as the main cursor amplitude or  $h_0$  in DSP literature.

## Architecture

---

### The Data Path and the Control Path

The DFE circuit itself is a high-speed execution engine with no intelligence. It consists of delay registers, DAC multipliers, and a summing node, all operating at the full data rate (32 Gbps or higher). It blindly multiplies slicer decisions by whatever tap weights are stored in its configuration registers and subtracts the result from the incoming waveform, billions of times per second.

The intelligence resides in a completely separate, slower digital logic block: the adaptation engine. This separation defines two distinct signal paths within the receiver.

**The data path** carries the live, high-speed analog waveform through the CTLE, into the summing node, through the slicer, and out as recovered digital data. Every component on this path operates at the full line rate.

**The control path** operates at a much lower clock rate. The adaptation engine monitors the signal quality by comparing the slicer output against a known reference, computes updated tap weights using the LMS algorithm, and writes the converged values into the configuration registers that the high-speed data path reads.

### How the Adaptation Engine Measures the Echo

The DFE tap weights cannot be determined while transmitting random user data, because the receiver cannot distinguish an echo from a noise spike or from the intended signal of the current bit. The measurement requires a controlled environment.

During the link training phase, the transmitter sends a known pseudo-random binary sequence (PRBS). The receiver knows exactly what bit pattern to expect. Alongside the main slicer (which outputs binary decisions for the data path), the receiver employs a parallel circuit called the error slicer. The error slicer measures the actual raw analog voltage of the incoming waveform and compares it against the adaptive target voltage.

If the main slicer declares the current bit is a logic 1 (target **+0.6 V**) but the error slicer measures **+0.4 V**, the adaptation engine records an error signal of  $+0.4 - 0.6 = -0.2$  V. This error represents the total residual distortion at the current sampling instant, comprising echoes from all preceding bits, thermal noise, and crosstalk.

## DAC Implementation and Current-Steering Execution

The adaptation engine stores each converged tap weight as a digital binary number in a static configuration register. The incoming waveform is physical analog electricity. A digital binary value cannot directly subtract from an analog voltage. The digital-to-analog converter bridges this domain boundary.

The DAC reads the binary number from the register and generates a proportional steady-state electrical current. The DFE does not perform high-speed binary multiplication on the fly. Instead, the physical execution relies on current steering through transistor switches operating at the full line rate.

For each tap, a shift register holds the slicer's past binary decision for the corresponding bit period. That decision controls a pair of differential transistor switches connected to the DAC's steady current source. If the previous bit was a logic 1, one switch closes and steers the tap current (  $+I_{\text{tap}}$  ) into the correction wire. If the previous bit was a logic 0, the complementary switch closes and steers the opposite polarity current (  $-I_{\text{tap}}$  ) into the wire. The transistor pair acts purely as a polarity selector; the magnitude of the current is fixed by the DAC.

All tap output wires from all active taps physically connect to a single summing node (a resistive load or transimpedance amplifier). By Kirchhoff's Current Law, the positive and negative currents from all taps algebraically combine at that node, producing a net correction voltage. This summing happens passively in the analog domain, with no digital logic in the critical timing path. The net correction voltage is injected into the signal path just picoseconds before the slicer evaluates the next incoming bit, subtracting the predicted echo from the raw waveform.

The distinction between analog and mixed-signal implementations lies in how the tap weight itself is stored. In a purely analog system, the tap weight can be maintained as a charge on a capacitor, producing a DC voltage that controls the current source directly. In modern mixed-signal implementations (the dominant architecture), a slow digital state machine computes the LMS updates, stores the tap weight in a digital register, and uses the static DAC to output a constant current. The digital register provides indefinite storage without the charge leakage problems of a capacitor, and the adaptation engine can update the weight at its own slow clock rate without disturbing the high-speed data path.

## The LMS Algorithm

The Least Mean Squares algorithm iteratively drives each tap weight toward the value that minimizes the mean squared error between the received signal and the ideal signal. The update equation for each tap is:

$$w_k^{\text{new}} = w_k^{\text{old}} + \mu \cdot e[n] \cdot d[n - k]$$

where  $w_k$  is the weight of tap  $k$ ,  $\mu$  is the learning rate (a static constant programmed into the hardware),  $e[n]$  is the error signal measured on the current bit  $n$ , and  $d[n - k]$  is the slicer decision for the bit  $k$  time slots in the past.

**Initialization.** The adaptation engine sets all tap weights to zero.

**Error measurement.** The DFE applies the current (initially zero) tap weights to the incoming bit. The main slicer makes its decision. The error slicer compares the raw analog voltage against the adaptive target and records the error. For the first iteration with zero taps, the error captures the full, uncorrected distortion of the channel.

**Parallel correlation.** The engine simultaneously multiplies the error signal by the slicer decisions for each of the past 15 bits. Tap 1 correlates the error with bit  $N - 1$ . Tap 2 correlates the same error with bit  $N - 2$ . All 15 correlations execute in parallel.

**The learning rate.** If the full error were added directly to a single tap weight, the system would assume that one past bit caused the entire error. The next clock cycle would massively overcorrect, producing a large error of opposite sign, and the system would oscillate and diverge. The learning rate  $\mu$  is a small constant (typically on the order of 0.001 to 0.01) that scales each update to a microscopic fraction of the error. The algorithm reacts only to consistent, long-term trends and ignores volatile single-cycle anomalies.

The effective step size is self-regulating. The update magnitude is the product of  $\mu$ , the error, and the past bit decision. When the link first boots and the error is large, each update is proportionally large, and the taps converge quickly. As the taps approach their optimal values and the error shrinks, the updates shrink proportionally. The system decelerates smoothly and settles into equilibrium without requiring an explicit stopping criterion.

**Tap separation through statistical correlation.** If all 15 taps process the same error signal with the same learning rate during the same clock cycle, the mechanism that causes each tap to converge on a different value is the statistical independence of the PRBS training sequence.

Tap 1 correlates the error with bit  $N - 1$ . The physical echo from bit  $N - 1$  is large (it is the most recent predecessor). Every time bit  $N - 1$  is a  $+1$ , the echo pulls the error in a consistent, predictable direction. Every time bit  $N - 1$  is a  $-1$ , the echo pulls it in the opposite direction. Over thousands of iterations, tap 1 accumulates consistent, reinforcing updates that drive it toward a large converged value.

Tap 15 correlates the error with bit  $N - 15$ . On most physical channels, the impulse response has decayed to negligible amplitude by 15 bit periods. Bit  $N - 15$  has no systematic physical influence on the current error. The PRBS ensures that bit  $N - 15$  is equally likely to be  $+1$  or  $-1$ , so the correlation products are equally likely to be positive or negative. Over thousands of iterations, these random updates cancel each other, and tap 15 hovers near zero.

Each tap organically converges to the exact magnitude of the channel's impulse response at its corresponding time delay. The full set of converged tap weights reconstructs the physical decay profile of the channel through purely statistical accumulation.

**Continuous tracking.** In production hardware, the adaptation engine rarely freezes the tap weights. The error residual drops to near zero as the taps converge, causing the updates to become vanishingly small. The engine continues running in the background at all times, allowing the taps to track slow environmental changes such as silicon heating, power supply drift, and mechanical stress on connectors.

## Statistical Correlation and Noise Rejection

The LMS algorithm's ability to extract the true echo from a composite error signal that contains thermal noise, crosstalk, and overlapping echoes from dozens of prior bits relies on the pseudo-random nature of the training data.

Thermal noise is completely random. A noise spike is equally likely to add or subtract voltage. Over millions of iterations, the random noise contributions to the tap updates cancel to zero. Echoes from distant bits (those beyond the channel's impulse response duration) are also uncorrelated with the current error, and their contributions cancel similarly.

The only physical phenomenon that consistently, predictably pushes the error signal in the same direction every time a specific past bit has a specific polarity is the actual physical echo of that bit at that specific time delay. The algorithm isolates this echo organically by multiplying the error by the past bit decision and averaging over millions of iterations. The true impulse response rises out of the noise floor as a statistical certainty.

## Link Training Negotiation

The CTLE peaking gain, the DFE tap weights, and the transmitter FFE tap weights are all interdependent. The optimal CTLE setting depends on how much pre-compensation the FFE provides. The optimal DFE taps depend on the residual ISI left after CTLE correction. The optimal FFE taps depend on how aggressively the receiver is equalizing on its end.

Modern high-speed protocols resolve this interdependence through a backchannel negotiation process. The receiver measures the incoming signal quality after its own CTLE and DFE have settled. If the eye opening remains insufficient, the receiver sends digital commands backward to the transmitter through a dedicated low-speed sideband channel, requesting specific adjustments: increase the FFE pre-cursor weight, decrease the post-cursor, or shift the balance between transmitter pre-compensation and receiver restoration. The transmitter and receiver iterate through this negotiation until the combined system converges on the best achievable eye opening for that specific physical link.

Link training occurs at every system boot, every wake-from-sleep event, and every hot-plug cable insertion. The physical characteristics of the channel change with temperature, and the silicon characteristics change between manufactured lots. The negotiation must adapt to the immediate physical reality, not rely on static design-phase calculations.

Training algorithms search a multi-dimensional parameter space and can become trapped in local minima: parameter combinations that appear optimal within a narrow neighborhood but are significantly worse than the global optimum. Engineers must define appropriate initial seed values and algorithmic step sizes during silicon design to give the training state machine the best probability of convergence. If the total equalization capability designed into the silicon is physically insufficient for the channel loss, link training will fail regardless of the search algorithm.

## Receiver Topologies: Analog Front-End vs. DSP-Based

Traditional SerDes receivers (and BERT error detectors) follow the analog front-end architecture described throughout this guide. The degraded waveform passes through the CTLE, enters the analog summing node where the DFE subtracts its echo estimate, and the slicer makes a binary decision directly on the analog voltage. No analog-to-digital converter exists in this path. The slicer is a pure comparator that reports only whether the voltage crossed the threshold, not what the voltage was.

Ultra-high-speed receivers (112G+ Ethernet, 224G development) employ a fundamentally different DSP-based architecture. The incoming waveform feeds directly into a high-speed ADC that digitizes the raw analog voltage into a stream of multi-bit samples. A digital signal processor performs all equalization mathematically on the digitized data: CTLE filtering, DFE echo cancellation, and the slicer decision all execute as numerical operations in the digital domain. No analog slicer exists. This architecture trades massive power consumption and silicon area for the flexibility and precision of digital processing, and it enables equalization algorithms that would be physically impossible to implement as analog circuits.

The DFE implementation illustrates the architectural contrast most clearly. In the analog front-end, each tap weight is converted to a steady current by a DAC, and physical transistor switches steer that current into a summing node based on the slicer's past decisions. The correction happens in the analog domain, subject to thermal noise from the transistors, mismatch in the current sources, and parasitic capacitance that limits the switching speed. Scaling to a high tap count (15+ taps) demands proportionally more analog circuitry, and each additional current source contributes its own noise to the summing node.

In the DSP-based receiver, the entire DFE computation is numerical. The tap weights are stored as binary numbers in static RAM registers. The ADC has already converted the incoming waveform into a binary sample (for example, the integer ~~+45~~ ). A shift register holds the previous slicer decisions as digital ones and zeros. A bank of digital multiplier circuits (built from thousands of logic gates) multiplies each stored tap weight by the corresponding past bit decision. A digital adder computes the total correction value. The DSP subtracts this correction from the ADC's incoming sample, and a numerical threshold comparison replaces the analog slicer: if the corrected value is above zero, the DSP registers a logic 1; below zero, a logic 0. No analog summing node exists. No physical current is steered. The entire process is arithmetic on integers, and once the signal is digitized, no further analog noise enters the computation. A DSP receiver can run a 40-tap DFE without the compounding thermal noise, physical layout constraints, or current-source mismatch of an analog summing node.

## Edge Cases

---

### Error Propagation in the DFE

The accuracy of DFE echo cancellation depends entirely on the slicer making correct decisions. If the slicer misidentifies a logic 1 as a logic 0, the DFE computes the wrong echo polarity for that bit and subtracts an incorrect voltage from the subsequent waveform. This incorrect subtraction corrupts the next bit's voltage, increasing the probability that the slicer makes another wrong decision. The errors can cascade forward through multiple bit periods.

A single slicer error in a multi-tap DFE corrupts the echo estimate for every tap that references the incorrect decision, potentially damaging up to  $k$  subsequent bits for a  $k$ -tap system. Forward error correction (FEC) exists partly to catch these burst errors. The DFE's self-correcting property limits the damage: once the erroneous decision shifts out of the tap memory, the echo estimates from subsequent correct decisions steer the system back toward accurate cancellation. Channels with marginal eye openings are most susceptible to error propagation, because the probability of the initial slicer error is highest when the voltage margin is smallest.

## CTLE Noise Amplification

The CTLE amplifies all content in its boost frequency range indiscriminately, including thermal noise, crosstalk from adjacent traces, and power supply ripple. Excessive peaking gain on a low-loss channel amplifies noise far beyond the signal content in those frequency bands, degrading the signal-to-noise ratio. The optimal CTLE setting balances signal restoration against noise injection, and the link training algorithm searches for this balance as part of the multi-dimensional parameter optimization.

## Impulse Response Measurement Without a Single Isolated Pulse

The mathematically ideal method for characterizing a channel's echo profile is to transmit a single isolated pulse followed by a long string of zeros and directly measure the trailing decay. Three physical constraints prevent this approach in production hardware.

The clock recovery circuit at the receiver requires continuous voltage transitions to maintain synchronization. A single pulse followed by silence causes the CDR to lose timing lock, blinding the receiver. The AC coupling capacitors on the link interpret a long string of identical bits as DC and begin charging, causing the baseline voltage to wander and collapse. A single analog measurement of the tail is also overwhelmed by random thermal noise, producing unreliable tap weight estimates.

The PRBS training sequence solves all three problems simultaneously. Its dense transitions keep the CDR locked. Its statistically balanced density of ones and zeros keeps the coupling capacitors neutral. Its millions of repeated measurements, processed through the LMS correlation, average out all random noise components and extract the true impulse response as a statistical certainty.

## Chapter 15: PAM4 Signaling and Gray Coding

---

Every high-speed serial link faces a physical ceiling: the analog bandwidth of the copper channel. Skin effect resistance and dielectric absorption attenuate high-frequency signal components progressively, and no amount of equalization can restore energy that has been permanently converted to heat. Raising the symbol rate to push more bits through the channel requires wider analog bandwidth, which the physical copper cannot provide at modern data rates.

PAM4 (Pulse Amplitude Modulation, 4-level) breaks through this ceiling by encoding two digital bits into a single voltage symbol. The transmitter subdivides the available voltage swing into four evenly spaced levels, each representing a unique two-bit combination. The symbol rate stays constant while the data rate doubles. A 64 Gb/s PAM4 link operates at the same 32 GBaud symbol rate, and requires the same 16 GHz analog bandwidth, as a 32 Gb/s NRZ link. The bandwidth problem is solved by packing more information into each voltage measurement rather than switching faster.

The cost is severe. Compressing four voltage levels into the same total swing that previously held two reduces each individual eye opening to one-third of the NRZ eye height, inflicting a 9.5 dB signal-to-noise ratio penalty. The receiver must evaluate the incoming waveform against three separate decision thresholds instead of one, and the voltage margin at each threshold is so small that even minor noise spikes can push the sampled voltage into the wrong level. Aggressive equalization and forward error correction become mandatory.

The mapping between voltage levels and digital bit patterns is itself a critical engineering decision. A naive sequential binary assignment places a two-bit boundary at the most noise-vulnerable threshold in the center of the signal, causing a single analog error to flip both digital bits simultaneously. Gray coding eliminates this vulnerability by constraining adjacent voltage levels to differ by exactly one bit, ensuring that the most common type of analog error produces only a single digital bit flip.

### Core Concepts

---

#### NRZ Signaling and the Return-to-Zero Distinction

The term "non-return-to-zero" (NRZ) describes the temporal behavior of the waveform, not its physical voltage topology. Older signaling protocols used return-to-zero (RZ) encoding, which forced the voltage back to a neutral zero-volt baseline in the middle of every bit period to guarantee a clock transition. A sequence of three logic ones in an RZ scheme pulses high, drops to zero, pulses high, drops to zero, pulses high, and drops to zero.

An NRZ transmitter abandons this forced neutral state entirely. The transmitter holds the voltage at the designated logical level for the full duration of the bit period. A sequence of three logic ones produces a continuous high voltage for three complete bit times. The signal never returns to a neutral baseline between identical bits.

This design eliminates the transition overhead of RZ encoding but creates a specific timing challenge: long runs of identical bits contain no voltage transitions for the receiver to track. The clock and data recovery (CDR) circuit must infer the timing from the statistical properties of the data stream. Engineers commonly use the term "NRZ" interchangeably with "PAM2" to distinguish this two-level modulation from four-level PAM4 systems.

High-speed serial links use differential signaling. The physical link consists of two wires carrying equal and opposite voltages around a constant common-mode DC average. The common-mode voltage never swings. The receiver subtracts the negative trace from the positive trace to evaluate the differential signal, rejecting any noise that appears equally on both wires.

## Data Rate, Symbol Rate, and the Nyquist Frequency

Three distinct metrics describe the performance of a serial link, and confusing any two of them leads to incorrect bandwidth calculations.

**Data rate** measures the total number of digital bits transmitted per second. This is the throughput metric that protocols and applications consume.

**Symbol rate** (baud rate) measures the number of physical voltage changes on the wire per second. Each voltage transition constitutes one symbol. For NRZ (PAM2), one symbol carries one bit, so the data rate and symbol rate are numerically identical. For PAM4, one symbol carries two bits, so the data rate is exactly twice the symbol rate.

**Nyquist frequency** is the fundamental analog frequency of the fastest possible physical transition pattern. A continuous alternating sequence (101010 in NRZ, or cycling through all four levels in PAM4) completes one full analog cycle every two symbol periods. The Nyquist frequency is therefore exactly half the symbol rate. The physical channel must possess adequate analog bandwidth at this frequency to deliver an open eye diagram to the receiver.

A 32 Gb/s NRZ data stream switches at 32 GBaud. The Nyquist frequency is 16 GHz. Moving to 64 Gb/s PAM4 keeps the symbol rate at 32 GBaud and the Nyquist frequency at 16 GHz. The system doubles the digital throughput without demanding any additional analog bandwidth from the copper channel.

The Shannon-Hartley theorem formalizes the absolute boundary governing this tradeoff. The maximum theoretical data rate scales directly with the available analog bandwidth and logarithmically with the signal-to-noise ratio. PAM4 trades vertical voltage margin for horizontal bandwidth efficiency: the data rate doubles, but the noise tolerance drops by 9.5 dB. Modern standards such as PCIe Gen 6 and IEEE 802.3ck operate on channels that often lack sufficient analog bandwidth to perfectly preserve the Nyquist frequency even at the PAM4 symbol rate. The silicon must reconstruct the original data using complex equalization algorithms alongside forward error correction to compensate for the missing bandwidth.

## PAM4 Voltage Levels and the 9.5 dB Penalty

A perfectly linear PAM4 signal divides the maximum differential voltage swing into three equal intervals, producing four evenly spaced levels. Using the normalized representation, these levels sit at  $+1$ ,  $+1/3$ ,  $-1/3$ , and  $-1$ . Engineers typically refer to the four states using the integer labels  $+3$ ,  $+1$ ,  $-1$ , and  $-3$  to denote their uniform spacing.

The physical voltages on the wire are much smaller than these normalized values. A typical high-speed transmitter outputs a total peak-to-peak differential swing of 800 mV. The four levels distribute as:

- Peak positive: **+400** mV
- Inner positive: **+133** mV
- Inner negative: **−133** mV
- Peak negative: **−400** mV

This four-level architecture creates three distinct vertical eye openings stacked between adjacent voltage levels, replacing the single large eye of an NRZ signal. Each individual eye spans only one-third of the total voltage swing. The reduction in vertical eye height from the full NRZ swing to one-third of that swing imposes a signal-to-noise ratio penalty of  $20 \cdot \log_{10}(1/3) \approx 9.5$  dB.

The practical consequence is that the receiver must evaluate incoming voltages with three times the precision it needed for NRZ. A noise spike that would have been harmless in an NRZ system can easily push a PAM4 voltage across an adjacent threshold. The equalization chain (FFE, CTLE, DFE) must not only restore the frequency balance of the incoming waveform but must do so with enough residual voltage margin that the three compressed eyes remain open and distinguishable.

### The Triple Slicer and Thermometer Code

The receiver deploys three separate voltage slicers to evaluate the three PAM4 eyes. Each slicer is a high-speed digital comparator that outputs a pure logic 1 when the incoming analog voltage exceeds its threshold and a pure logic 0 when the voltage falls below it. No slicer produces an analog value or a fractional voltage.

The three thresholds sit at the midpoints between adjacent voltage levels:

- Top threshold: **+266** mV (between **+400** and **+133** mV)
- Middle threshold: **0** V (between **+133** and **−133** mV)
- Bottom threshold: **−266** mV (between **−133** and **−400** mV)

The receiver reads the three binary outputs concurrently as a logic array. The digital ones fill from the bottom upward, producing a pattern known as thermometer code (analogous to mercury rising in a glass thermometer):

| Incoming Voltage                  | Thermometer Code (Top, Mid, Bot) | Interpretation                     |
|-----------------------------------|----------------------------------|------------------------------------|
| <b>+400</b> mV (level <b>+3</b> ) | 1-1-1                            | Above all three thresholds         |
| <b>+133</b> mV (level <b>+1</b> ) | 0-1-1                            | Below top, above middle and bottom |
| <b>−133</b> mV (level <b>−1</b> ) | 0-0-1                            | Below top and middle, above bottom |
| <b>−400</b> mV (level <b>−3</b> ) | 0-0-0                            | Below all three thresholds         |

A hardcoded digital logic table maps these four unique thermometer arrays directly to the final two-bit data payload. The hardware extracts the voltage state and the corresponding data bits simultaneously without performing any mathematical addition or analog measurement.

## Gray Coding and the Hamming Distance Constraint

The mapping between the four thermometer codes and the two-bit data payload determines how much damage a single analog noise event inflicts on the digital data stream.

**The sequential binary vulnerability.** Standard binary counts sequentially: **00** , **01** , **10** , **11** . Mapping these values directly to the four voltage levels from bottom to top (or top to bottom) places the **01** → **10** boundary at the zero-volt middle threshold. This is the most noise-vulnerable location in the entire signal, where the two inner voltage levels sit closest together and thermal noise has the highest probability of pushing a sample across the threshold.

A noise spike at the zero-volt boundary causes the receiver to misidentify the inner positive level as the inner negative level (or the reverse). Under sequential binary mapping, the inner positive level decodes as **01** and the inner negative level decodes as **10** . Both bits flip simultaneously from a single analog error. The bit error rate artificially doubles for the most common type of receiver mistake, and the resulting two-bit burst error is more difficult for FEC to correct than a single-bit error.

The direction of the sequential counting (bottom-to-top or top-to-bottom) does not eliminate this flaw. The transition between the decimal values 1 and 2 always forces the binary shift between **01** and **10** , and aligning this two-bit boundary with the middle threshold guarantees double-bit errors at the most vulnerable location regardless of which end of the voltage scale starts at **00** .

**The Gray code solution.** Gray code is a binary sequence constructed so that every adjacent pair of values differs in exactly one bit position. This property is described formally as a Hamming distance of 1 between all neighboring entries. The standard PAM4 Gray code mapping assigns:

| Voltage Level  | Thermometer Code | Gray Code Payload |
|----------------|------------------|-------------------|
| +400 mV ( +3 ) | 1-1-1            | 1-0               |
| +133 mV ( +1 ) | 0-1-1            | 1-1               |
| -133 mV ( -1 ) | 0-0-1            | 0-1               |
| -400 mV ( -3 ) | 0-0-0            | 0-0               |

A noise spike at the zero-volt threshold now causes the receiver to misread **11** as **01** (or the reverse). Only a single digital bit changes. The FEC algorithm handles single-bit errors routinely, while multi-bit cascade errors often overwhelm the mathematical parity checks and cause dropped packets. Every adjacent-level error at every threshold in the PAM4 signal produces at most one bit flip under Gray coding.

## Worked Examples

---

### Noise Spike at the Zero-Volt Threshold

An incoming signal at the inner positive level should arrive at **+133** mV. A negative thermal noise spike of 150 mV hits at the exact sampling instant, dragging the sampled voltage to **−17** mV. The receiver sees a voltage below the zero-volt middle threshold and categorizes it as the inner negative level.

**Under sequential binary mapping:** The inner positive level decodes as **01** . The inner negative level decodes as **10** . Both the left bit and the right bit flipped. One analog mistake produced two digital errors.

**Under Gray code mapping:** The inner positive level decodes as **11** . The inner negative level decodes as **01** . The left bit flipped from 1 to 0. The right bit remained 1. One analog mistake produced exactly one digital error.

The physical noise event was identical in both scenarios. The encoding scheme determined whether that single analog mistake propagated into one digital error or two.

### NRZ vs. PAM4 Bandwidth Calculation

A system architect needs to transport 64 Gb/s of data through a copper channel whose insertion loss profile supports adequate signal quality up to approximately 16 GHz.

**NRZ approach:** At one bit per symbol, 64 Gb/s requires 64 GBaud. The Nyquist frequency is 32 GHz. The channel cannot support this bandwidth. The link fails.

**PAM4 approach:** At two bits per symbol, 64 Gb/s requires only 32 GBaud. The Nyquist frequency is 16 GHz. The channel supports this bandwidth. The link is physically viable, but the 9.5 dB SNR penalty requires aggressive equalization and FEC to achieve acceptable error rates.

This calculation is the fundamental engineering motivation behind the industry transition from NRZ to PAM4 at data rates above approximately 56 Gb/s per lane.

## Architecture

---

### Gray Code Hardware Implementation

The conversion from a three-bit thermometer array to a two-bit Gray code payload is implemented in physical silicon using combinatorial logic gates. The IEEE standards for high-speed Ethernet and PCIe dictate the specific truth table to guarantee uniform Gray coding across the industry.

The mapping yields a remarkably efficient hardware realization. Examining the truth table reveals that the first data bit equals 1 for the top two voltage levels and 0 for the bottom two levels. The middle slicer output exhibits exactly this behavior: it outputs 1 whenever the voltage exceeds zero and 0 whenever the voltage falls below zero. The hardware simply wires the middle slicer output directly to the first data bit. No logic gates are required.

The second data bit equals 0 for the peak positive level, 1 for both inner levels, and 0 for the peak negative level. This pattern matches the output of an exclusive-OR (XOR) gate fed by the top and bottom slicers. The XOR gate outputs 1 when its two inputs differ and 0 when they match:

- Peak positive ( **+3** ): Top slicer = 1, Bottom slicer = 1. Inputs match. XOR output = 0.
- Inner positive ( **+1** ): Top slicer = 0, Bottom slicer = 1. Inputs differ. XOR output = 1.
- Inner negative ( **-1** ): Top slicer = 0, Bottom slicer = 1. Inputs differ. XOR output = 1.
- Peak negative ( **-3** ): Top slicer = 0, Bottom slicer = 0. Inputs match. XOR output = 0.

The entire Gray code decoder consists of a single direct wire and a single XOR gate. The protocol committee chose this specific truth table assignment, and the silicon engineer discovered that the assignment translates into the simplest possible physical circuit. The conversion executes instantaneously with negligible propagation delay and zero power overhead beyond the gate itself.

## The Equalization Mandate

The 9.5 dB SNR penalty transforms the equalization chain from a performance optimization into an absolute necessity. An NRZ receiver with a full-height eye opening can tolerate moderate residual ISI and still make correct decisions. A PAM4 receiver with one-third the vertical margin cannot.

The transmitter FFE must pre-compensate for channel loss with greater precision, because any residual loss that closes the already-compressed eyes risks merging adjacent voltage levels into an indistinguishable blur. The receiver CTLE must amplify the high-frequency edge content without amplifying noise into the narrow voltage margins. The DFE must cancel ISI echoes that would be tolerable in NRZ but are catastrophic when the decision boundaries are spaced only 266 mV apart.

Forward error correction provides the final safety net. At PAM4 data rates, the raw bit error rate after equalization is typically orders of magnitude worse than what the application layer requires. The FEC encoder adds mathematical redundancy (typically Reed-Solomon codes) that allows the receiver to detect and correct a bounded number of bit errors per block. The combination of equalization and FEC operates as a system: equalization reduces the raw BER to a level the FEC can handle, and FEC closes the remaining gap to the application's target error rate.

## Edge Cases

---

### The Zero-Volt Threshold as the Dominant Error Source

The three PAM4 decision thresholds do not contribute equally to the total bit error rate. The middle threshold at zero volts separates the two inner voltage levels ( **+133** mV and **-133** mV), which are spaced only 266 mV apart. The upper and lower thresholds each separate an inner level from a peak level ( **+133** to **+400** mV, or **-133** to **-400** mV), with 534 mV of spacing.

The inner levels are twice as close to each other as they are to the outer levels. Thermal noise and residual ISI are far more likely to push a sample across the tightly spaced middle boundary than across either of the wider outer boundaries. The zero-volt threshold dominates the error statistics, which is precisely why Gray coding's protection at this boundary is so critical.

## CDR Challenges with PAM4

The clock and data recovery circuit in a PAM4 receiver must extract timing information from a signal with three stacked eyes rather than one. The transition density visible to the CDR depends on which voltage levels are transitioning. A transition between the two inner levels (  $+1$  to  $-1$  ) crosses the zero-volt threshold cleanly. A transition between an inner level and the adjacent peak level (  $+1$  to  $+3$  ) never crosses zero at all.

The CDR must track transitions across all three eyes simultaneously, and the timing extraction algorithms are more complex than their NRZ counterparts. The reduced voltage margin at each eye also means that the timing uncertainty (jitter) of the recovered clock is larger for PAM4 than for NRZ at the same symbol rate, further tightening the system's jitter budget.



PART 5

# SerDes Architecture and Clocking

## Chapter 16: SerDes Architecture and Eye Diagrams

---

A serializer/deserializer (SerDes) link is the fundamental architecture underlying every modern high-speed interface: PCIe, Ethernet, USB, HDMI, and DisplayPort all transmit data as a continuous stream of serial bits across a single differential pair, then reconstruct the parallel data words inside the receiving chip. The entire system exists because of a speed mismatch between the physical channel and the digital logic. A copper trace or optical fiber can carry analog transitions at 112 Gbps, but no digital logic gate can switch at 112 GHz. The SerDes architecture resolves this contradiction by confining extreme-speed analog processing to a tiny, specialized front-end circuit and then gearing the data down to a clock rate that standard digital silicon can handle.

The eye diagram is the primary visual tool for evaluating whether the analog portion of that system is working. It represents the accumulated voltage and timing margin available to the receiver at the moment it commits to a bit decision. Three fundamentally different instruments produce eye diagrams through three different physical mechanisms: real-time oscilloscopes capture continuous waveforms and fold them mathematically, sampling oscilloscopes build dot-by-dot reconstructions using a hardware trigger, and bit error ratio testers (BERTs) sweep a decision threshold across a time-voltage grid and map the statistical boundary where errors begin. Understanding how each instrument constructs its eye, and what that eye physically represents, is essential for interpreting signal integrity measurements at any data rate.

### Core Concepts

---

#### The SerDes Signal Path

The transmitter and receiver halves of a SerDes link divide naturally into analog and digital domains, with the boundary between them defining the critical interface where signal integrity measurements take place.

**Transmitter (TX).** The digital core generates parallel data words at a relatively low clock rate. A multiplexer (MUX) serializes those parallel words into a single high-speed bitstream. For NRZ signaling, the output stage is a current-mode differential driver that steers a constant current between two complementary outputs. To send a logic 1, the driver routes current so that D+ swings positive and D- swings negative, producing a positive differential voltage (for example, +300 mV). To send a logic 0, the driver reverses the current steering, producing a negative differential voltage (-300 mV). The current flows continuously; information is encoded in the direction of steering, not in the presence or absence of a signal. The differential receiver threshold sits at 0 V, the natural midpoint of this symmetric swing. Modern PAM4 transmitters replace the simple current-steering driver with a digital-to-analog converter (DAC) that produces four distinct voltage levels from a 2-bit input symbol, as described in [PAM4 Signaling and Gray Coding](#).

**Channel.** The differential pair (copper traces, cables, connectors, vias) carries the analog signal from transmitter to receiver. Channel loss from skin effect and dielectric absorption progressively degrades the signal, closing the eye diagram.

**Receiver (RX).** The receiver faces the inverse problem. A continuous stream of high-speed serial bits arrives with no accompanying clock signal. The receiver must recover the embedded clock, make binary decisions on each bit, and convert the serial stream back into parallel words that the digital core can process. This chain of operations proceeds through four functional stages:

1. **Clock and Data Recovery (CDR).** The first block the incoming signal encounters is the CDR, a phase-locked loop (PLL) circuit that locks onto the data transitions and generates a local clock aligned to the center of each bit period. The mechanics of the PLL, including loop bandwidth, damping, and jitter transfer, are covered in [CDR and PLL Loop Dynamics](#).
2. **Slicer.** At each tick of the recovered clock, an analog voltage comparator (the slicer) examines the instantaneous differential voltage. If the voltage exceeds the threshold (0 V for NRZ), the slicer outputs a digital 1. If the voltage falls below the threshold, it outputs a digital 0. The slicer has no knowledge of the actual voltage magnitude. It reports only whether the signal crossed the boundary. This single-bit binary decision is the moment the analog world ends and the digital world begins.
3. **Deserializer (DEMUX).** The recovered bits arrive at the full serial line rate, far too fast for digital logic to process individually. The DEMUX acts as a gearbox: it collects bits from the slicer and distributes them across multiple parallel output lanes, stepping the clock rate down proportionally. A 112 Gbps serial stream split across 128 parallel lanes requires only an 875 MHz clock on each lane, well within the capability of standard CMOS digital logic.
4. **Word alignment.** The DEMUX produces a continuous flow of parallel bits with no inherent framing. The digital logic scans for predetermined alignment markers embedded in the data stream (comma symbols in 8b/10b-encoded protocols like early PCIe generations, or sync headers in 64b/66b-encoded protocols like 100G+ Ethernet). Once the alignment marker is found, the logic locks the byte boundaries and hands structured data to the upper protocol layers. The encoding schemes and protocol framing details are covered in [Line Coding, FEC, and Protocol Framing](#).

## The PMA/PCS/MAC Sublayer Hierarchy

The silicon implementing this signal path is divided into strict functional domains. When hardware engineers refer to "the receiver," they almost always mean the analog circuitry at the chip's edge: the CDR, the slicer, and the DEMUX. This collection of analog circuits constitutes the Physical Medium Attachment (PMA) sublayer. Its responsibility is limited to getting bits off the wire and handing them to the digital core. The PMA has no knowledge of Ethernet frames, PCIe transactions, or any protocol semantics.

The digital core sits behind the PMA and processes the incoming parallel data through two additional sublayers. The Physical Coding Sublayer (PCS) performs word alignment, descrambling, and forward error correction (FEC). The Media Access Control (MAC) sublayer reconstructs the actual protocol data units (Ethernet frames, PCIe packets) and hands them to the system processor. While these sublayers process the received data, they are functionally and physically distinct from the analog receiver.

The significance of this boundary is practical: when an oscilloscope or BERT probes the link, it is analyzing the analog signal feeding into the PMA. If the eye diagram shows an unrecoverable closure, the PMA slicer will pass bit errors into the PCS, and the FEC must attempt to correct them mathematically. The eye diagram measures the margin available to the analog receiver, not the end-to-end error rate after FEC correction.

## Differential Signaling and Common-Mode Rejection

High-speed serial links universally use differential signaling. The transmitter drives two complementary traces (D+ and D-) with equal and opposite voltage swings. The receiver measures only the voltage difference between them.

The reason is noise immunity. An electromagnetic interference source couples equally into both traces of a tightly routed differential pair. If a noise spike adds +100 mV to both D+ and D-, the differential voltage is unchanged:

$$(D^+ + \text{noise}) - (D^- + \text{noise}) = D^+ - D^-$$

The noise cancels completely. This property, common-mode rejection, is the only mechanism that makes reliable data transmission at 112 Gbps across a piece of fiberglass physically possible.

Both oscilloscopes and BERTs exploit this property during measurement. When two SMA cables carry D+ and D- to a scope, the instrument performs the mathematical subtraction (Channel 1 minus Channel 2) and displays the resulting clean differential waveform as the eye diagram. A BERT performs the same subtraction internally before feeding the differential signal to both its CDR (for clock recovery) and its slicer (for bit decisions).

## The Analog Front-End vs. DSP Receiver Architectures

Two fundamentally different receiver architectures exist for processing the incoming analog signal, and the distinction determines both the equalization strategy and the test methodology.

**Architecture A: Analog front-end (AFE).** The degraded differential waveform passes through continuous-time linear equalization (CTLE), enters an analog summing node where the decision feedback equalizer (DFE) subtracts its echo estimate, and the slicer makes a binary decision directly on the analog voltage. No analog-to-digital converter exists in this signal path. This architecture is used in the majority of SerDes implementations through 56 Gbps NRZ and is also the architecture inside a BERT error detector. The equalization mechanics are covered in [CTLE and DFE](#).

**Architecture B: DSP-based receiver.** The analog waveform enters a high-speed ADC that measures the exact instantaneous voltage (for example, +212 mV) and converts it to a multi-bit digital word. A digital signal processor (DSP) then performs equalization, clock recovery, and bit-decision logic entirely in the digital domain. No analog slicer exists in this path. This architecture appears in ultra-high-speed implementations (112 Gbps PAM4 Ethernet and beyond) where the channel loss is so severe that analog equalization alone cannot open the eye sufficiently.

The critical distinction for test and measurement is that an oscilloscope operates on Architecture B principles (it uses an ADC to digitize the analog voltage and produces a visual representation of the waveform), while a BERT operates on Architecture A principles (it uses a slicer to make binary decisions and produces statistical error data). These are not interchangeable measurements. They answer fundamentally different questions about the link.

## Architecture

---

### Eye Diagram Construction: Three Instruments, Three Methods

The eye diagram is a composite image formed by overlaying many individual bit transitions on a common time axis. Every instrument that produces an eye diagram must solve two problems: capturing the analog voltage at a sufficient number of time-voltage coordinates, and establishing a timing reference that defines where each unit interval (UI) begins and ends. The three primary instruments solve these problems through radically different mechanisms.

**Real-time oscilloscope (RTO).** The RTO contains a massively fast ADC (for example, 256 billion samples per second on a modern Keysight UXR-Series) that digitizes the continuous analog voltage from the transmission line into an unbroken record stored in memory. Because the entire waveform is captured at once, no external clock or trigger is required to construct the eye. Instead, software algorithms implement a mathematical CDR: they analyze the stored waveform, calculate the data transition positions, determine the bit rate, and identify the clock frequency. The software then chops the continuous record into individual one-UI slices and mathematically overlays them on a common time axis. This folding operation produces the eye diagram entirely in post-processing.

The RTO display depends on how the software renders the stored data. In linear view, the captured record appears as a long trace of voltage over time. In persistence mode with the timebase set to one or two UI, the software overlays all slices and the characteristic eye shape emerges. If the timebase is zoomed out to show the entire PRBS pattern, the repeated overlay of the same pattern with slight noise and jitter variations produces a thick, glowing trace rather than a recognizable eye.

**Sampling oscilloscope (DCA).** A sampling scope builds the eye diagram dot by dot over millions of measurement cycles. Its ADC is far slower than an RTO's ADC (perhaps 1 million samples per second), but it achieves far greater analog bandwidth through an elegant trick. At the input sits an ultra-high-bandwidth analog track-and-hold circuit consisting of an electronic switch (typically a diode bridge) and a microscopic capacitor. When a hardware trigger fires, the switch closes for a few picoseconds, and the capacitor charges to the instantaneous voltage on the transmission line. The switch then snaps open, freezing the captured voltage. The slow ADC takes its time to measure the held voltage, and the scope plots exactly one dot at one (time, voltage) coordinate on the display.

The bandwidth of a sampling scope is not limited by the ADC speed. It is limited entirely by the aperture time of the analog switch. If the switch takes 10 picoseconds to open and close, and the signal being measured oscillates at 100 GHz, the voltage swings significantly during the capture window and the capacitor records a blurred average. Measuring a 100 GHz signal requires a sub-picosecond aperture, and building an analog switch with that speed is one of the most difficult challenges in test and measurement hardware.

A sampling scope cannot build an eye diagram without an external timing reference. High-speed serial links transmit no separate clock signal, so the DCA requires a hardware CDR (either built into the measurement module or provided as a standalone instrument) to extract the clock from the data stream. The data signal is split using pick-off tees (high-bandwidth microwave power splitters) or, in integrated DCA modules, tapped internally. One copy of the signal feeds the measurement channel. The other copy feeds the hardware CDR, which locks onto the data transitions and generates a physical clock signal routed to the DCA's trigger input. Each tick of that recovered clock fires the track-and-hold circuit once, capturing one dot.

**Bit Error Ratio Tester (BERT).** A BERT does not contain an ADC and does not measure analog voltage. Its "eye diagram" is constructed through a fundamentally statistical process. The BERT contains two subsystems: a pattern generator (PG) that transmits a known test pattern, and an error detector (ED) that receives data and checks it for errors. The ED uses a slicer identical in principle to the slicer inside a SerDes receiver chip: a purely analog comparator that outputs a single bit (1 or 0) based on whether the instantaneous voltage exceeds a programmable threshold.

To construct a BER contour (also called a statistical eye), the BERT sweeps its slicer across a two-dimensional grid of time delay and voltage threshold. At each grid coordinate, the BERT samples millions of bits and records the bit error ratio. At the center of the eye, where the voltage and timing margins are largest, the BER is zero. As the slicer moves toward the edges (higher voltage thresholds or earlier/later sampling instants), it begins intersecting the signal transitions and the BER increases. The BERT maps the BER at every coordinate and renders the result as a heat map. The dark, open region at the center is where errors are absent. The colored contour lines at the boundary represent specific BER thresholds (for example,  $10^{-6}$ ,  $10^{-9}$ ,  $10^{-12}$ ).

A scope eye diagram shows what the signal looks like. A BERT BER contour shows how much margin a real silicon receiver has before it starts making errors. The two measurements are complementary, not redundant.

## How the BERT Error Detector Achieves Pattern Synchronization

The ED exploits the deterministic nature of PRBS (pseudo-random binary sequence) test patterns. A PRBS pattern is generated by a linear feedback shift register (LFSR) driven by a mathematical polynomial. The sequence appears random but is entirely predictable once the LFSR state is known.

The ED contains its own internal PRBS generator implementing the same polynomial. When data begins flowing, the ED examines a short segment of the incoming bitstream and uses it to synchronize its internal generator's state. Once pattern synchronization is achieved, the ED can predict every subsequent bit indefinitely. From that point, operation reduces to a simple comparison: the physical slicer output (the received bit) is compared against the mathematically predicted bit. A mismatch increments the error counter.

The ED has no concept of what voltage is arriving. It does not know whether the signal swing is 50 mV or 500 mV. The engineer configures the nominal slicer threshold (0 V for standard differential NRZ), and the slicer reports only whether the voltage is above or below that threshold at the instant the recovered clock triggers. The BER contour sweep deliberately moves this threshold away from the optimal center to map the boundaries of the eye.

## TX Testing vs. RX Testing

The physical asymmetry between the transmitter and receiver pins creates two distinct testing paradigms that require different instruments.

**Transmitter testing** uses an oscilloscope. The chip under test generates a continuous test pattern (such as PRBS-31) from its TX pins. The signal travels through the channel to the measurement point, and the scope captures the eye diagram and measures jitter, rise time, and voltage swing. The scope is a passive listener that characterizes the quality of the transmitted signal.

**Receiver testing** uses a BERT. Data does not "come out" of the RX pins. The RX pins are strictly inputs, and once a signal enters the chip, it proceeds directly into the internal analog front-end and digital core with no external observation point. Testing a receiver therefore requires stressing it to the boundaries of its specification and determining whether it can still recover data correctly.

The BERT pattern generator acts as a controlled-impairment transmitter. It injects calibrated amounts of jitter, noise, and signal attenuation to degrade the eye diagram to the exact limits specified by the protocol compliance standard. This degraded signal is driven into the chip's RX pins. Two mechanisms verify whether the receiver survived:

**Loopback mode.** The chip is placed in a special test configuration where the received bits, after passing through the RX analog front-end and digital core, are routed internally to the TX block and retransmitted. The BERT error detector receives the loopback signal and compares it against the original transmitted pattern. Any mismatch indicates that the receiver made a bit error.

**Internal error counters.** The chip's digital protocol layer is configured to expect a specific PRBS pattern. When the analog receiver passes an incorrect bit into the digital core, the logic detects the discrepancy and increments an error counter accessible through a management interface (such as a JTAG or I2C debug port). The engineer reads the accumulated error count to determine the receiver's bit error ratio under stress.

The scope tests the transmitter by listening. The BERT tests the receiver by attacking it.

## Worked Examples

---

### Constructing an Eye Diagram from a Linear Capture

An RTO captures a continuous 10 Gbps NRZ differential signal for 1 microsecond. At 256 GSa/s, the memory contains 256,000 voltage samples covering 10,000 bit periods.

**Step 1: Software CDR.** The instrument software analyzes the stored waveform and identifies the data transitions (rising and falling edges). From the transition positions, the software calculates the average bit period: 100 ps (corresponding to 10 Gbps).

**Step 2: Folding.** The software partitions the 1 microsecond capture into 10,000 individual slices, each exactly one UI (100 ps) wide. Every slice contains approximately 25.6 voltage samples.

**Step 3: Overlay.** All 10,000 slices are superimposed on a common 0-to-100-ps time axis. The rising edges from all slices cluster together, the falling edges cluster together, and the flat regions between transitions overlap. The characteristic diamond-shaped opening of the eye emerges from the composite overlay. Noise and jitter cause slight variations in the transition positions and settled voltage levels, which appear as thickness in the overlaid traces. A wider spread indicates more jitter or noise; a tighter bundle indicates a cleaner signal.

### BERT Contour Sweep at a Single Grid Coordinate

A BERT error detector is configured to test a 28 Gbps NRZ link at one specific coordinate: a voltage threshold of +50 mV and a timing offset of +5 ps from the recovered clock edge.

**Step 1: Configuration.** The engineer programs the slicer threshold to +50 mV (above the nominal 0 V center) and adds a +5 ps delay to the sampling clock. The PRBS-31 pattern generator is running continuously.

**Step 2: Acquisition.** The ED samples 10 billion bits at this (time, voltage) coordinate and detects 3 bit errors. The BER at this coordinate is  $3 \times 10^{-10}$ .

**Step 3: Sweep.** The BERT increments the voltage threshold to +55 mV and repeats. Then +60 mV, +65 mV, and so on, recording the BER at each point. As the threshold approaches the signal's settled high level, the BER increases rapidly toward 0.5 (random guessing). The BERT then resets the voltage to +50 mV and sweeps the timing offset to +10 ps, +15 ps, and beyond. Over the full two-dimensional sweep, the accumulated BER values form a contour map of the eye.

**Interpreting the contour.** The open center of the eye (where BER is zero or immeasurably small) represents the margin available to the receiver. The contour line at  $10^{-12}$  shows the boundary where the receiver would experience one error per trillion bits. Protocol compliance specifications typically require the eye to remain open at specific BER thresholds, and the BER contour directly measures whether that requirement is met.

## Edge Cases

---

### The Hardware CDR Configuration Trap

When an engineer configures a CDR loop bandwidth and peaking on a DCA or BERT, that configuration applies to the test equipment's CDR, not to the silicon chip's internal CDR. The test equipment CDR serves only to provide a timing reference for the measurement instrument. Protocol compliance standards (such as PCI-SIG or IEEE 802.3) define a mathematical "Golden PLL" with specific loop bandwidth and peaking parameters. The test equipment CDR must be configured to match the Golden PLL so that the measurement reflects what a compliant receiver would see. Misconfiguring the test equipment CDR (for example, setting the loop bandwidth too wide so that the CDR tracks jitter that a real receiver would not track) produces an artificially optimistic eye measurement that does not represent the true link margin.

### The DCA Bandwidth Bottleneck

A sampling scope's bandwidth is constrained entirely by the aperture time of its track-and-hold switch, not by the speed of its ADC. A DCA with a 70 GHz bandwidth rating and a 1 MSa/s ADC can measure signals that a 20 GHz real-time scope with a 256 GSa/s ADC cannot resolve. The bandwidth limitation is analog, not digital. As data rates push toward 112 Gbaud PAM4, the required aperture times drop below 1 picosecond, and the physical construction of analog switches fast enough to achieve these apertures represents a fundamental engineering constraint on measurement capability.

## Signal Splitting and Impedance Considerations

Splitting the data signal using pick-off tees introduces a 6 dB power loss at each split point (half the signal energy goes to each output port). On marginal links where the eye is already near closure, this additional loss can alter the measurement. Integrated DCA modules that tap the signal internally minimize this effect by using high-impedance sensing rather than resistive power splitting. When using external pick-off tees, the engineer must account for the insertion loss and verify that the splitter's bandwidth exceeds the signal bandwidth.

## Chapter 17: CDR and PLL Loop Dynamics

---

High-speed serial links transmit no separate clock signal. PCIe, USB, Ethernet, and every modern protocol embed the clock directly into the data stream, forcing the receiver to reconstruct a local clock from the timing of the incoming voltage transitions. This reconstruction is the job of the Clock and Data Recovery (CDR) circuit, and its core mechanism is a Phase-Locked Loop (PLL): a feedback control system that continuously adjusts a local oscillator to track the incoming data edges.

The CDR is not simply a clock extractor. It is a filter with carefully chosen dynamics. Low-frequency timing drift (from temperature changes, supply voltage variation, or crystal oscillator tolerance) must be tracked so that the recovered clock moves with the data and the eye stays open. High-frequency timing noise (from random thermal events, crosstalk, or power supply transients) must be rejected so that the recovered clock holds steady and does not chase every stray edge displacement. The boundary between tracking and rejection is the loop bandwidth, and the shape of the frequency response near that boundary determines whether the loop amplifies jitter (peaking) or smoothly attenuates it. These dynamics are governed entirely by the physics of the analog components inside the loop filter: a charge pump, a resistor, a capacitor, and a voltage-controlled oscillator.

Two fundamentally different receiver architectures implement clock recovery. The traditional analog CDR uses a PLL with a voltage-controlled oscillator that physically speeds up and slows down to track the data. The DSP-based digital CDR, which dominates the 112G/224G era, runs a free-running ADC clock and uses mathematical interpolation to reconstruct the optimal sampling point from digitized samples. Both architectures solve the same problem (aligning a local sampling instant to the center of the incoming data eye), but they solve it through radically different physical mechanisms with distinct tradeoffs in noise performance, power consumption, and silicon area.

Understanding CDR loop dynamics is essential for two reasons. First, the loop bandwidth and peaking directly determine which jitter components close the eye and which are invisible to the receiver. Second, protocol compliance testing requires test equipment (oscilloscopes and BERTs) to emulate the exact CDR behavior specified by the standard's "Golden PLL," so that the measured eye diagram reflects what a real receiver would see.

### Core Concepts

---

#### Why Embedded Clocking Requires CDR

Running an external clock at gigahertz frequencies introduces a fundamental synchronization problem. A millimeter of trace length difference or a fraction of a degree of temperature change causes the clock phase and data phase to drift apart. This skew shifts the sampling instant toward the edge of the eye, degrading bit error performance. At multi-gigahertz data rates, maintaining phase alignment between a separate clock trace and a data trace across manufacturing variation, thermal gradients, and cable lengths is physically impractical.

Embedding the clock into the data eliminates the synchronization problem entirely. The transmitter encodes information as voltage transitions on a single differential pair. The receiver extracts timing information from those same transitions and generates a local clock that is, by construction, aligned to the data it must sample. Skew and thermal wander are eliminated because the clock is derived from the data itself.

This approach requires guaranteed transitions in the data stream. A long unbroken run of identical bits produces no voltage transitions, leaving the CDR with nothing to track. Line coding schemes (8b/10b, 128b/130b, scrambling) prevent long runs and maintain transition density, as covered in [Line Coding, FEC, and Protocol Framing](#).

## The Phase-Locked Loop as a Feedback Control System

The CDR is built around a phase-locked loop consisting of four functional blocks arranged in a closed feedback path.

**Phase Detector.** The phase detector compares the timing of each incoming data transition against the edge of the locally generated clock. The comparison produces a signal proportional to the time gap between the two edges. A common hardware implementation is the Bang-Bang (Alexander) Phase Detector, which uses D-flip-flops and a transistor switch to control a charge pump. When a data edge arrives, a flip-flop turns on the switch and connects a precision current source to the loop filter. When the local clock edge arrives, a second flip-flop turns the switch off. If the clock is far behind the data, the switch stays closed longer and a large pulse of current flows. If the two edges arrive simultaneously, the switch barely opens and almost zero current flows. The output is a current pulse whose duration encodes the magnitude and polarity of the phase error. A data edge arriving before the clock produces an "UP" pulse; a data edge arriving after the clock produces a "DOWN" pulse.

**Charge Pump and Loop Filter.** The UP or DOWN pulses from the phase detector turn on a current switch (the charge pump), which sends a burst of current into the loop filter. The loop filter converts these current pulses into a smoothly varying control voltage. The components of the loop filter (resistors and capacitors) determine whether the loop is first-order or second-order, and their values set the loop bandwidth, damping factor, and peaking behavior. The physics of the loop filter is the central subject of this guide.

**Voltage-Controlled Oscillator (VCO).** The VCO generates the local clock signal at a frequency determined by its input voltage. Inside the VCO, a specialized diode called a varactor acts as a voltage-variable capacitor. The varactor sits within an LC tank circuit whose resonant frequency is  $f = 1/(2\pi\sqrt{LC})$ . Changing the voltage applied to the varactor physically expands or contracts the depletion region inside the silicon, changing its capacitance. That capacitance change shifts the resonant frequency of the LC tank, which changes the VCO output frequency. Every VCO has a hardware specification called VCO gain ( $K_{\text{vco}}$ ), measured in hertz per volt. If a VCO has a gain of 100 MHz/V, applying 1 V to its input increases the output frequency by 100 MHz. The relationship is strictly linear over the operating range.

**Feedback Path.** The VCO's output clock is routed back to the phase detector, closing the loop. The phase detector compares this clock against the next incoming data transition, and the cycle repeats continuously.

## Frequency as a Consequence of Phase Alignment

A PLL has no concept of frequency. The phase detector only sees phase: the time difference between a data edge and a clock edge at one specific instant. Frequency matching is purely a mathematical consequence of the system finding equilibrium. From calculus, frequency is the derivative of phase with respect to time. If the phase difference between clock and data is constant (not growing, not shrinking), the frequencies must be identical. The loop does not measure frequency, does not compare frequencies, and has no frequency detector. It controls frequency indirectly by controlling phase, and the mathematical relationship between phase and frequency ensures that a stable phase lock implies a stable frequency lock.

This principle reveals the distinction between two quantities that the loop connects:

- **Phase Difference** is a static measurement: the time gap (in picoseconds) between a data edge and a clock edge at one specific instant.
- **Rate of Phase Change** is literally a frequency difference. If the phase difference is growing over time, the clock and data are running at two different frequencies. If the phase difference is constant, the frequencies match exactly.

The phase detector measures only the phase difference. It has no memory and performs no differentiation. Every time an edge arrives, it simply measures the instantaneous phase gap and outputs a voltage proportional to that distance. The loop uses this static phase difference to set the VCO frequency, which in turn dictates the rate of phase change. When the loop reaches equilibrium, the phase difference is constant and the rate of phase change is zero, meaning the frequencies have matched.

## Loop Bandwidth: The Boundary Between Tracking and Filtering

The CDR acts as a low-pass filter for jitter. The loop bandwidth is the cutoff frequency of that filter.

**Below the loop bandwidth.** Timing variations slower than the loop bandwidth are tracked by the recovered clock. If the data slowly wanders back and forth due to thermal drift or supply voltage variation, the recovered clock wanders with it. The relative timing between clock and data remains constant, the sampling instant stays centered in the eye, and these slow jitter components do not degrade the bit error ratio. The CDR "tracks them out."

**Above the loop bandwidth.** Timing variations faster than the loop bandwidth are too rapid for the feedback loop to follow. The recovered clock holds steady while the incoming data edges vibrate around it. These high-frequency jitter components shift the data transitions relative to the fixed sampling instant and close the eye. They appear directly in the measured jitter.

The loop bandwidth is not set arbitrarily. Making it too wide causes the clock to chase stray noise spikes (a cosmic ray hit, a crosstalk burst, a random thermal event that displaces a single edge by 100 picoseconds). An excessively wide bandwidth would cause the clock to violently track that noise, destroying the timing reference for subsequent bits. Making the loop bandwidth too narrow prevents the clock from tracking legitimate low-frequency drift, causing the sampling instant to slide away from the eye center. Protocol standards specify the loop bandwidth as a compromise between these two failure modes.

## Clock Architecture and Phase Interpolators

Transmitter and receiver typically share a common reference architecture. A stable, low-frequency crystal oscillator on the motherboard (commonly 100 MHz) provides the frequency reference. The transmitter multiplies this reference up to the line rate using a PLL (for example, 100 MHz multiplied by 160 to produce 16 GHz). The receiver multiplies its own 100 MHz crystal to the same approximate frequency using a separate PLL.

The CDR's task is then to take that locally generated clock and fine-tune its phase to align precisely with the incoming data transitions. Many modern implementations accomplish this fine-tuning through a Phase Interpolator rather than directly adjusting the VCO frequency. The Phase Interpolator takes the local clock and mathematically blends between adjacent phase taps to shift the sampling edge in precise fractional-UI increments, achieving the same alignment result through digital phase control rather than analog frequency manipulation.

## Type-I and Type-II Loop Classification

PLL loop dynamics literature describes loops as "first-order" or "second-order" based on the number of energy-storage elements in the loop filter, and as "Type-I" or "Type-II" based on the number of integrators in the open-loop transfer function. In the context of CDR design, a first-order loop with a resistor-only filter is a Type-I system (proportional control only), and a second-order loop with a resistor-capacitor filter is a Type-II system (proportional plus integral control). The Type-I and Type-II naming convention maps directly onto the first-order and second-order architectural distinction that governs the loop's steady-state behavior, its ability to eliminate phase error, and its susceptibility to overshoot and ringing.

## Architecture

---

### First-Order Loop: Proportional Control (Type-I)

The simplest loop filter is a single resistor. The charge pump outputs a current proportional to the phase error. Ohm's law converts that current to a voltage across the resistor ( $V = I \times R$ ). This voltage directly controls the VCO frequency. The control voltage is strictly proportional to the instantaneous phase error:  $V_{\text{control}} = K \times \text{Phase Error}$ . A larger phase error produces a higher voltage; a smaller phase error produces a lower voltage; zero phase error produces zero voltage.

A first-order loop is unconditionally stable. Its frequency response is a smooth low-pass curve with zero peaking. It tracks low-frequency jitter with unity gain and rolls off monotonically above the loop bandwidth.

### Static Phase Offset: The Catch-22

The fatal limitation of a Type-I loop is the steady-state phase error. The mechanism is a catch-22 rooted in the physics of the resistor.

Consider a transmitter crystal running at 101 MHz while the receiver's VCO naturally oscillates at 100 MHz when its control input sits at 0 V. The VCO needs a non-zero control voltage to speed up from its resting frequency to match the data. With a VCO gain of 100 MHz/V, matching the 1 MHz frequency offset requires exactly +10 mV. The resistor generates voltage only when current flows through it, and the charge pump generates current only when the phase detector measures a non-zero time gap between the data edge and the clock edge.

Here is what happens cycle by cycle:

1. The data is arriving at 101 MHz while the VCO clock runs at 100 MHz. The clock is slower, so each successive clock edge falls further behind the corresponding data edge. The phase error is actively growing.
2. As the phase error grows, the phase detector outputs wider UP pulses. The voltage across the resistor rises in proportion: at 5 degrees of phase error, the voltage is 1.1 V; at 10 degrees, it is 1.3 V; at 15 degrees, it is 1.5 V.
3. The rising voltage pushes the VCO frequency upward: 100.1 MHz, 100.5 MHz, 100.9 MHz.
4. At exactly 15 degrees of phase error and 1.5 V of control voltage, the VCO reaches 101 MHz. At this instant, the clock and data are running at the same speed.

The system now freezes. The clock edges stop falling further behind the data edges, so the phase error stops growing. The phase detector keeps measuring the same 15 degrees of separation edge after edge, and the resistor keeps producing the same 1.5 V. The VCO holds steady at 101 MHz.

This equilibrium is self-arresting. The loop cannot correct the 15-degree phase error because doing so would destroy the voltage required to maintain the frequency match. If the clock attempted to catch up to the data (reducing the phase error toward zero), the voltage across the resistor would drop. As the voltage dropped, the VCO frequency would fall below 101 MHz, the clock would start falling behind again, and the phase error would grow back to 15 degrees. The system is trapped: the phase error is stuck at a permanent, non-zero offset because that offset is the only mechanism available for generating the control voltage the VCO requires. This permanent offset is called **static phase offset**.

The hardware cannot escape this trap because the Type-I loop filter has no memory. The voltage exists only as long as the current flows, and the current flows only as long as there is an error. The loop uses phase difference to control frequency, but it cannot drive phase difference to zero while simultaneously maintaining the voltage that keeps the frequency matched.

## Second-Order Loop: Proportional Plus Integral Control (Type-II)

Adding a capacitor in series with the resistor transforms the loop filter into a proportional-integral (PI) controller. The total control voltage delivered to the VCO is the sum of two components:

$$V_{\text{total}} = V_{\text{resistor}} + V_{\text{capacitor}} = I \times R + \frac{1}{C} \int I dt$$

The resistor provides an instantaneous voltage proportional to the current (proportional control). The capacitor accumulates charge over time, producing a voltage that integrates the history of all past error currents (integral control).

The series arrangement is critical. In a parallel configuration, the capacitor would absorb and smooth out the fast voltage transients from the resistor, eliminating the proportional response. In series, the VCO receives the instant "kick" from the resistor (to react to sudden phase jumps) plus the slow, steady memory of the capacitor (to correct permanent frequency offsets). The resistor and capacitor values jointly determine the loop bandwidth and damping factor.

**Eliminating steady-state phase error.** When the transmitter runs at 101 MHz and the VCO rests at 100 MHz, the phase detector generates a persistent trickle of error current. This current charges the capacitor, and its voltage slowly rises: 1 mV, 5 mV, 9 mV. The VCO speeds up progressively. As the frequency gap closes, the phase error shrinks and the error current diminishes. At the moment the clock achieves perfect phase alignment, the charge pump stops outputting current. Unlike the resistor, which drops to zero voltage immediately, the capacitor holds its accumulated charge. It stores the +10 mV like a battery, continuously feeding it to the VCO. The loop has reached perfect equilibrium: the phase error is zero, the current is zero, and the VCO is held at exactly 101 MHz by the capacitor's stored charge. The integrator drives the steady-state phase error to absolute zero regardless of the frequency offset between transmitter and receiver crystals.

In real silicon, the capacitor slowly leaks charge through parasitic resistance. At gigabit data rates, however, data transitions arrive so frequently that the phase detector sends continuous micro-pulses of correction current to maintain the stored voltage at the correct level. The leakage is negligible in practice.

Virtually all high-speed standards (PCIe, USB, Ethernet) mandate second-order CDRs because every physical link involves two independent crystal oscillators with parts-per-million (PPM) frequency tolerances that cannot be bridged without integral control.

### **The Overshoot Mechanism: How Type-II Eliminates Phase Error**

The capacitor does not simply accumulate charge to the correct voltage and stop. The physical mechanism by which the Type-II loop drives phase error to zero requires a temporary frequency overshoot.

Phase cannot be changed without a difference in frequency. Phase is the integral of frequency over time, so the only way for the clock to catch up to the data is if the clock temporarily runs *faster* than the data. This is the fundamental constraint that separates the Type-II from the Type-I.

Scenario: Data is 101 MHz, VCO rests at 100 MHz.

1. **The Lag.** The VCO is slower than the data, so the data pulls ahead. A large phase error accumulates as the clock edges fall further and further behind.
2. **The Charge.** The phase detector sees this growing error and pumps current into the capacitor. The voltage rises, and the VCO frequency rises with it.
3. **The Crossing Point.** The VCO reaches 101.0 MHz. At this instant, the frequencies match. The phase error stops growing. But a massive phase error remains from the entire time the capacitor was charging up and the clock was running slower than the data.
4. **The Overshoot.** That large residual phase error still drives the phase detector to output UP pulses. The charge pump keeps pushing current into the capacitor. The voltage continues to rise, pushing the VCO frequency past 101.0 MHz to, say, 101.2 MHz.
5. **Closing the Gap.** The VCO is now running faster than the data (101.2 MHz vs. 101.0 MHz). Because the clock is faster, it begins to catch up to the data. The phase error shrinks.
6. **The Settling.** As the phase error shrinks toward zero, the phase detector stops pumping excess current. The control voltage settles back down. The loop reaches equilibrium at the exact moment the phase error hits zero and the VCO frequency returns precisely to 101.0 MHz. The capacitor holds the correct voltage indefinitely.

The Type-I loop physically cannot execute this overshoot. In a Type-I system, the control voltage is tied to the instantaneous phase error. To overshoot the frequency to 101.2 MHz, the phase error would have to grow to the value that generates enough voltage for 101.2 MHz. But the moment the clock speeds up and the phase error starts shrinking, the voltage instantly drops, immediately pulling the frequency back down before the phase can ever reach zero. The Type-I system lacks the component (the capacitor) that would allow it to hold excess voltage while the phase error is being corrected.

The capacitor is the physical memory device that makes the overshoot possible. It stores the baseline voltage required for 101 MHz while the phase detector temporarily pushes it higher to close the phase gap. Once the gap is closed, the excess charge drains and the voltage settles to the baseline.

### **Ringings: The Time-Domain Cost of Overshoot**

The overshoot does not stop cleanly on the first pass. The Type-II loop exhibits ringing (decaying oscillation) before it settles, and this is a direct consequence of second-order control theory.

When the phase error reaches zero, the capacitor still holds the excess voltage that pushed the VCO past 101.0 MHz. The clock is still running faster than the data, so the clock edge physically passes the data edge. The phase error immediately becomes non-zero in the opposite direction (the clock is now early instead of late). The phase detector switches from sending UP pulses to sending DOWN pulses. The DOWN pulses extract electrons from the capacitor, lowering the voltage. The frequency drops.

If the loop filter were a pure capacitor with no resistor, this overshoot and correction cycle would repeat indefinitely with constant amplitude, exactly like a frictionless pendulum swinging through the center point. The system would never settle.

The resistor in the loop filter acts as electrical damping. It provides a proportional voltage component that reacts instantaneously to the current error, counteracting the stored energy in the capacitor. Each successive overshoot is smaller than the previous one. The control voltage and phase error oscillate with decreasing amplitude until the system parks at exactly zero phase error and exactly the correct frequency. Hardware designers tune the ratio of the resistor and capacitor values to control the damping factor. A critically damped or slightly overdamped response flattens the oscillation and settles quickly. An underdamped response rings longer but may converge faster on the initial approach. Protocol compliance specifications constrain the maximum peaking (typically 1 dB or 2 dB), which is the frequency-domain manifestation of the time-domain ringing.

## Digital CDR: Free-Running Clock and Mathematical Phase Recovery

Ultra-high-speed DSP-based receivers (112G+ Ethernet, 224G development) replace the analog PLL with a fully digital clock recovery architecture. There is no voltage-controlled oscillator, no charge pump, and no analog feedback loop.

**The free-running clock.** On the receiver board, a local crystal oscillator drives the ADC at a fixed, stable frequency. This clock is not controlled by a feedback loop; it runs freely. The ADC takes snapshots of the incoming waveform strictly at this rate. The transmitter on the opposite end of the link has its own crystal oscillator, running at a frequency that differs by a few parts per million. This slight frequency mismatch causes the ADC sampling point to drift slowly across the data eye over time.

**Mueller-Muller phase detection.** The DSP examines the digitized amplitude of the signal during bit transitions to determine the sampling phase offset. The algorithm compares the voltage of the current sample against the voltage of the previous sample at a data transition. If the ADC is sampling at the dead center of the eye, the voltages on either side of a transition are symmetrical: a 0-to-1 transition produces a sample at  $-0.5$  V followed by a sample at  $+0.5$  V, with equal magnitudes. If the ADC clock is drifting late, the sample during the 0 is taken as the voltage is already rising toward the transition and might measure  $-0.2$  V, while the sample during the 1 is taken further up the flat top of the pulse and might measure  $+0.6$  V. The asymmetry between these two amplitudes tells the DSP exactly which direction the clock is drifting, and the magnitude of the imbalance is mathematically proportional to the time offset.

**Phase interpolation.** The DSP maintains a digital phase offset register tracking exactly how far off-center the ADC is currently sampling. A digital FIR interpolation filter uses this offset to compute the voltage that the waveform would have had at the exact center of the eye, even though the ADC never physically sampled that instant. The interpolator examines the sample taken slightly before the ideal point and the sample taken slightly after, then calculates the intermediate value using the known signal shape. The DSP continuously updates the phase register and applies this mathematical correction to every incoming sample.

The Mueller-Muller phase detector has a critical advantage over the Bang-Bang (Alexander) phase detector used in analog CDRs. The Bang-Bang detector is an edge-based detector: it examines only the polarity of the data transition and produces a binary early/late decision. The Mueller-Muller detector is a baud-rate detector: it operates on the amplitude of the signal at the sampling point itself, extracting phase information without requiring a separate sample at the data edge. This allows the ADC to operate at exactly the symbol rate (one sample per UI) rather than at twice the symbol rate, halving the ADC's required sampling speed and power consumption.

## Cycle Slipping

The free-running ADC clock and the transmitter clock are not phase-locked. Their slight frequency mismatch causes the phase offset to accumulate continuously. If the transmitter is slightly faster than the receiver's ADC clock, each successive sample drifts a little later relative to the data eye. Eventually, the accumulated phase offset reaches 360 degrees (one full bit period), and the DSP resets the phase counter to zero. This event is called a **cycle slip**.

**Dropped bits (data faster than the ADC clock).** The data bits arrive slightly faster than the ADC can sample them. Each sample drifts progressively later in the data eye. After thousands or millions of bits, a data bit slides entirely past the ADC between two consecutive samples. The receiver physically misses that bit. The phase offset register reaches 360 degrees and resets.

**Duplicated bits (ADC clock faster than the data).** The ADC samples arrive slightly faster than the data transitions. Each sample drifts progressively earlier in the data eye. Eventually, the ADC samples twice while the same data bit is present on the wire. The receiver records two copies of a single bit. The phase offset register reaches  $-360$  degrees and resets.

A cycle slip is not a catastrophic failure. The digital CDR immediately resumes normal tracking after the phase counter reset. The system does not lose synchronization or require retraining. The single dropped or duplicated bit is a localized error that Forward Error Correction (FEC) is designed to absorb. DSP-based SerDes architectures depend on robust FEC (Reed-Solomon codes at the physical layer) specifically because cycle slips are an expected, periodic consequence of running without a phase-locked oscillator. The FEC algorithm identifies the missing or duplicated bit in the received sequence and mathematically reconstructs the correct data, allowing the system to run indefinitely without a physical PLL.

## Peaking: The Cost of the Integrator

The capacitor that eliminates steady-state phase error introduces a time delay into the feedback path. A capacitor cannot change its voltage instantaneously; it requires time to accumulate charge. In AC circuit analysis, this charging time manifests as a phase shift (up to 90 degrees) between the current entering the capacitor and the voltage appearing across it.

This delayed correction interacts destructively with jitter at frequencies near the loop bandwidth.

**Low-frequency jitter.** The capacitor's charging delay is negligible relative to the slow jitter period. The VCO tracks the data accurately.

**High-frequency jitter.** The jitter oscillates so rapidly that the capacitor averages out to a flat DC voltage. The tracking stops, and the low-pass filter does its job.

**Jitter near the loop bandwidth.** At this crossover frequency, the jitter oscillates at the exact rate where the capacitor's delay is maximally destructive. By the time the capacitor charges up and delivers a correction voltage telling the VCO to slow down, the incoming jitter has already reversed direction and the VCO should be speeding up. The delayed correction adds to the existing error rather than subtracting from it. This constructive interference causes the VCO to swing further out of alignment than the original jitter. The feedback loop is chasing its tail: each correction arrives at the wrong moment and amplifies the oscillation rather than damping it.

On a Bode plot of the Jitter Transfer Function (JTF), this amplification appears as a bulge above 0 dB just before the rolloff frequency. A first-order filter shows a flat 0 dB line (unity gain, perfect 1:1 tracking) that smoothly curves downward at the cutoff. A second-order filter with peaking shows the same flat 0 dB line, but the curve bulges upward (for example, to +1.5 dB or +2 dB) just before the cutoff, then drops steeply. That physical bulge is the peaking.

The practical consequence is direct: if incoming data carries a jitter component at exactly the peaking frequency with an amplitude of 10 picoseconds, and the JTF gain at that frequency is +2 dB (a voltage ratio of approximately 1.26), the recovered clock oscillates by approximately 12.6 picoseconds. The clock is shaking more violently than the data. The relative timing displacement between clock and data at this frequency is worse than if the CDR were absent, closing the eye and spiking the bit error ratio.

## VCO Phase Adjustment Through Temporary Frequency Change

The VCO has only one control input, and that input controls frequency. It has no separate phase control. Yet the CDR must correct both frequency offsets and phase errors.

A change in frequency inherently produces a change in phase. Consider two identical square waves running at 10 GHz with their rising edges perfectly aligned. To shift the second wave's phase so that its edges arrive 5 picoseconds earlier, the VCO control voltage is temporarily increased, speeding the oscillator to 11 GHz for a brief interval. The faster oscillation causes the edges to accumulate a phase lead. Once the edges have walked exactly 5 picoseconds forward, the control voltage drops back, returning the VCO to 10 GHz. The two waves now run at the same frequency, but the second is permanently shifted 5 picoseconds ahead. By transiently manipulating frequency, the VCO achieves a permanent phase correction.

## The Jitter Transfer Function

The Jitter Transfer Function (JTF) quantifies the CDR's filtering behavior. It is not a voltage measurement. It is a ratio of time quantities:

$$\text{JTF}(f) = 20 \log_{10} \left( \frac{\text{Jitter Amplitude Out}}{\text{Jitter Amplitude In}} \right)$$

At 0 dB (a ratio of 1), the recovered clock vibrates by exactly the same amount as the incoming data. The clock and data move together, and the eye stays open. At -20 dB (a ratio of 0.1), the recovered clock vibrates by only one-tenth of the input jitter amplitude. The clock is stable while the data moves, and that relative motion closes the eye. At +2 dB (a ratio of approximately 1.26), the recovered clock vibrates more than the data, and the amplified relative motion closes the eye even further.

The critical distinction between clock frequency and jitter frequency causes persistent confusion. The clock frequency (for example, 16 GHz) describes how fast the bits arrive on the wire. The jitter frequency (for example, 5 MHz) describes how fast the phase error oscillates back and forth. Peaking occurs at the jitter frequency where the loop's reaction delay causes constructive interference, typically in the single-digit to tens-of-megahertz range. It has no direct relationship to the multi-gigahertz bit rate.

## Worked Examples

---

### Steady-State Phase Error in a Type-I Loop

A receiver's VCO has a natural resting frequency of 100 MHz at 0 V input and a VCO gain of 10 MHz/V. The incoming data arrives at 101 MHz. The phase detector gain produces a control voltage strictly proportional to the phase error.

**Step 1: Initial state.** The VCO is running at 100 MHz. The data is arriving at 101 MHz. The clock is slower, so the phase error begins to grow.

**Step 2: Voltage rises with phase error.** At 5 degrees of phase error, the phase detector produces 1.1 V, pushing the VCO to approximately 100.1 MHz. At 10 degrees, it produces 1.3 V. At 15 degrees, it produces 1.5 V. At exactly 1.5 V, the VCO reaches 101 MHz.

**Step 3: The freeze.** The clock and data are now running at the same speed. The clock edges stop falling further behind. The phase detector measures 15 degrees on the next edge, and the next, and the next. It outputs 1.5 V every time. The voltage "stops rising" simply because the phase detector keeps measuring the same 15-degree distance.

**Step 4: Equilibrium.** The system is locked with the clock edge permanently 15 degrees behind the data edge. This offset generates the exact 1.5 V the VCO needs to hold 101 MHz. The frequencies match, but the clock samples off-center in the eye.

**Step 5: Attempting higher gain.** Increasing the phase detector gain would reduce the steady-state error (the same 1.5 V could be produced by a smaller phase offset). Excessively high gain makes the feedback loop violently unstable: any microscopic phase perturbation produces an enormous voltage spike that drives the VCO frequency wildly out of control. Hardware designers accept a moderate steady-state error as the stability cost of a first-order loop, or they add a capacitor to eliminate the error entirely.

### Peaking Amplification at the Loop Bandwidth

A second-order CDR has a loop bandwidth of 10 MHz and 2 dB of peaking. The incoming data carries a periodic jitter component at 10 MHz with an amplitude of 10 ps peak-to-peak.

**Step 1: JTF gain at 10 MHz.** The peaking specification indicates +2 dB gain at the loop bandwidth frequency. Converting to a linear ratio:  $10^{2/20} \approx 1.26$ .

**Step 2: Recovered clock jitter.** The clock oscillates by  $10 \text{ ps} \times 1.26 = 12.6 \text{ ps}$  peak-to-peak. The clock is moving 2.6 ps more than the data at this frequency.

**Step 3: Eye closure.** The relative jitter between clock and data at 10 MHz is 2.6 ps peak-to-peak. This additional timing displacement reduces the eye opening and increases the probability of a slicer error. The effect is concentrated at the peaking frequency; jitter at lower or higher frequencies experiences unity gain or attenuation, respectively.

## Edge Cases

---

### CDRs in Test Equipment and the Golden PLL

The error detector (ED) inside a BERT faces the same clock recovery challenge as a silicon receiver: it receives a serial stream with no separate clock and must recover timing to sample the incoming bits for error comparison. The BERT therefore contains a high-performance CDR with precisely configurable loop bandwidth and peaking.

The configurability exists because test equipment must not be "too good." If a BERT's CDR tracks jitter perfectly and recovers data that a standard silicon receiver would fail on, the compliance test is invalid. Protocol standards (PCI-SIG for PCIe, IEEE 802.3 for Ethernet) define a mathematical reference receiver model called the Golden PLL, which specifies the exact loop bandwidth and maximum peaking. The BERT's CDR must be configured to match this specification so that its error detection reflects the performance a real receiver would achieve.

The same principle applies to oscilloscopes. When a real-time oscilloscope draws an eye diagram, its software CDR recovers the clock and folds the waveform. If the software CDR bandwidth is set too low, the scope will not track low-frequency jitter that a real receiver would track, and the eye will appear artificially closed. If the bandwidth is set too wide, the scope will track jitter that a real receiver cannot follow, and the eye will appear artificially open. The CDR bandwidth must match the Golden PLL specification defined by the protocol standard.

### Infinite Bandwidth and the Noise Rejection Tradeoff

Setting the CDR loop bandwidth to infinity would cause the clock to track every timing displacement, including single-event transients (crosstalk bursts, power supply glitches, cosmic ray hits). A stray noise spike that displaces one data edge by 100 ps would cause the recovered clock to jump 100 ps, destroying the timing reference for the next several bits. The finite loop bandwidth gives the flywheel inertia: the clock tracks the slow, real drift of the system but ignores fast, random noise. The loop bandwidth is the engineering compromise between drift tracking and noise rejection, and the protocol specification mandates its exact value.

### Capacitor Leakage in the Loop Filter

The integrator capacitor in a second-order loop filter theoretically holds its charge indefinitely, maintaining the control voltage that keeps the VCO at the correct frequency even when the phase error is zero. In real silicon, parasitic resistance across the capacitor causes slow charge leakage. At the data rates of modern serial links (multi-gigahertz), data transitions arrive so frequently that the phase detector delivers continuous micro-pulses of correction current faster than the leakage drains charge. The net effect of parasitic leakage is negligible in practical CDR implementations.

## References

---

### Related Guides:

- [CTLE and DFE](#): Receiver-side equalization architecture, analog vs. DSP DFE
- [SerDes Architecture and Eye Diagrams](#): System-level signal path, hardware vs. software CDR
- [Line Coding, FEC, and Protocol Framing](#): Transition density, FEC for cycle slip tolerance

## Chapter 18: Line Coding, FEC, and Protocol Framing

---

A serial link is more than an analog channel. Between the physical voltage transitions on the wire and the usable data that reaches the processor sits a set of encoding, error correction, and framing operations that collectively transform raw bits into a reliable digital transport. Line coding guarantees the electrical properties the analog channel requires: DC balance to prevent baseline wander across AC coupling capacitors, bounded run length to maintain the transition density that clock and data recovery circuits depend on, and a controlled frequency spectrum that keeps energy within the channel's passband. Forward error correction injects mathematical redundancy that allows the receiver to detect and repair bit errors without retransmission, relaxing the analog margin requirements that would otherwise be physically impossible to meet at extreme data rates. Protocol framing provides the structural markers that allow the receiver to identify where bytes begin, distinguish data from control information, and hand complete packets to the upper layers of the protocol stack.

These operations also determine the relationship between two quantities that high-speed engineers must never conflate: the transfer rate (the physical signaling speed on the wire, measured in GT/s) and the data rate (the usable throughput after encoding overhead is removed, measured in Gb/s). Every measurement setup, from oscilloscope trigger configuration to BERT pattern generation, depends on this distinction.

### Core Concepts

---

#### Transfer Rate, Data Rate, and Encoding Overhead

The transfer rate specifies the number of physical symbols transmitted per second, independent of what those symbols contain. The unit is transfers per second (T/s), typically expressed as gigatransfers per second (GT/s). The transfer rate is what the oscilloscope sees, what the transmission line carries, and what the Unit Interval (UI) is derived from:

$$\text{UI} = \frac{1}{\text{Transfer Rate}}$$

A 3.2 GT/s signal has a UI of exactly 312.5 picoseconds. The oscilloscope triggers and recovers the clock based on those transfers regardless of what the bits represent.

The data rate describes the usable throughput after encoding overhead is subtracted. PCIe Gen 1 runs at 2.5 GT/s, but the link uses 8b/10b encoding: for every 8 bits of actual data, 10 bits are transmitted on the wire. The raw transfer rate is 2.5 billion symbols per second, but the data throughput is only 2.0 Gb/s ( $2.5 \times 8/10$ ). The two extra bits per symbol carry no user data. They exist solely to maintain DC balance and transition density.

The distinction matters for signal integrity measurement. The analog channel and physical layer test equipment do not differentiate between payload bits and overhead bits. Every symbol must survive the channel with adequate margin. Signal integrity engineers measure in GT/s because they are characterizing the physical transport, not the protocol throughput.

## Line Coding for DC Balance and Transition Density

AC coupling capacitors sit in series on the data path in PCIe, Ethernet, USB, and virtually every modern serial link. A capacitor blocks DC and passes AC. If the data stream contains a long unbroken run of identical bits, the capacitor sees what appears to be a static voltage and begins to charge, causing the baseline voltage at the receiver to drift. This baseline wander shifts the signal away from the decision threshold and causes bit errors. The physics of AC coupling and baseline wander are covered in [CTLE and DFE](#).

Line coding prevents this failure by guaranteeing an approximately equal density of ones and zeros over time.

**8b/10b encoding.** Used by PCIe Gen 1 and Gen 2, Gigabit Ethernet, SATA, SAS, Fibre Channel, and USB 3.0. The encoder maps every 8-bit data byte to a 10-bit symbol selected from a lookup table. The table is constructed so that each 10-bit symbol contains either an equal number of ones and zeros (five of each) or at most one extra bit of either polarity. The encoder tracks a running disparity counter that records whether the cumulative transmitted bit stream has sent more ones or more zeros. If the running disparity is positive (more ones have been sent), the encoder selects the symbol variant that contains more zeros, and vice versa. This feedback mechanism keeps the long-term DC content of the signal at zero, holding the AC coupling capacitors at their neutral charge state.

The encoding also enforces a maximum run length of five identical consecutive bits. This bounded run length ensures that voltage transitions occur frequently enough for the CDR to maintain phase lock. The overhead cost is 20%: two out of every ten transmitted bits carry no user data.

**128b/130b encoding.** Used by PCIe Gen 3 and later. The encoder groups 128 data bits into a block and prepends a 2-bit sync header. The sync header is always either **01** or **10**, guaranteeing at least one transition at every block boundary. The 128-bit payload is scrambled using a linear feedback shift register (LFSR) that produces a pseudo-random bit pattern with statistically balanced ones and zeros. The overhead drops from 20% (8b/10b) to approximately 1.6% ( **2/130** ), a massive efficiency improvement. The tradeoff is that scrambling provides statistical DC balance rather than the deterministic per-symbol guarantee of 8b/10b. Over any sufficiently long window, the scrambled stream converges to equal density, but short-term deviations are possible.

**64b/66b encoding.** Used by 10G+ Ethernet (IEEE 802.3). Structurally identical in principle to 128b/130b: a 2-bit sync header followed by a scrambled payload block. The sync header values ( **01** for data blocks, **10** for control blocks) provide both framing and transition density at the block boundary.

## Embedded Clocks vs. Forwarded Clocks

High-speed serial links (PCIe, Ethernet, USB) transmit no separate clock signal. The clock is embedded in the data transitions. The receiver uses a CDR circuit to extract timing information from the voltage edges and generate a local sampling clock. The mechanics of CDR, including PLL loop bandwidth, damping, peaking, and jitter transfer, are covered in [CDR and PLL Loop Dynamics](#).

The absence of a physical clock wire eliminates the term "clock cycle" from the serial link vocabulary. Hardware engineers instead use the term symbol (or Unit Interval). For NRZ signaling, one symbol carries one bit. For PAM4, one symbol carries two bits. The relationship between symbol rate, Nyquist frequency, and multi-level modulation is covered in [PAM4 Signaling and Gray Coding](#).

If the baud rate is 32 GT/s (32 billion symbols per second), the Nyquist frequency is exactly half: 16 GHz, because the fastest possible toggle pattern (101010...) takes two symbols to complete one full wave. For NRZ at 32 GT/s, that 16 GHz channel carries 32 Gb/s. For PAM4 at 32 GT/s, the same 16 GHz channel carries 64 Gb/s because each symbol encodes two bits.

**DDR memory uses the opposite architecture.** DDR is a parallel bus with a dedicated physical clock wire (CK) running alongside the data wires (DQ). The clock wire transmits a continuous, perfect alternating pattern. The data wires carry standard NRZ data. "Double Data Rate" means the memory controller outputs a new NRZ data symbol on each rising edge and each falling edge of the clock, achieving two transfers per clock cycle:

$$1.6 \text{ GHz clock} \times 2 \text{ edges} = 3.2 \text{ GT/s}$$

The Nyquist frequency of the data wires is half the transfer rate: 1.6 GHz. The physical clock frequency is also 1.6 GHz. These two values align exactly because one full clock cycle (rising edge plus falling edge) triggers exactly two data transfers.

Total bandwidth for a DDR channel uses the bus width. A standard non-ECC DDR4 channel is 64 bits wide:

$$3.2 \text{ GT/s} \times 64 \text{ bits} = 204.8 \text{ Gb/s} \quad (25.6 \text{ GB/s})$$

The forwarded-clock architecture avoids the complexity of CDR entirely. The receiver does not need to extract timing from data transitions because it receives an explicit clock reference. The cost is the dedicated clock wire, the strict trace-length matching required to keep the clock and data signals phase-aligned across the PCB, and the physical limitation that skew control becomes impractical as frequencies increase. This is why the highest-speed interfaces (PCIe Gen 5/6, 100G+ Ethernet) are all serial with embedded clocking.

## Architecture

---

### Protocol Stack Processing

Once the analog front end has made its bit decisions and the DEMUX has spread the serial stream into parallel lanes at a slower clock speed, the data enters the digital protocol stack. Processing proceeds through three stages:

**Word alignment.** The parallel lanes arrive as an undifferentiated stream of bits with no inherent framing. Digital logic scans for predetermined alignment markers: comma symbols (specific reserved 10-bit patterns) in 8b/10b-encoded protocols, or sync headers ( 01 / 10 ) in 64b/66b and 128b/130b-encoded protocols. Once the alignment marker is found, the logic locks the byte or block boundaries and establishes the framing that all subsequent processing depends on.

**Descrambling and decoding.** For 8b/10b protocols, the decoder reverses the lookup table, converting each 10-bit symbol back to its 8-bit data value and stripping the encoding overhead. For 128b/130b and 64b/66b protocols, the descrambler reverses the LFSR transformation to recover the original data. The transmitter scrambles data to prevent long runs of identical bits that would compromise CDR lock and AC coupling stability. The receiver must precisely reverse this operation.

**Forward error correction.** At 100G+ data rates, the raw bit error rate after analog equalization frequently exceeds the target (for example,  $10^{-12}$ ). FEC provides a mathematical safety net that corrects residual errors, including burst errors caused by DFE error propagation or transient noise events.

The Physical Coding Sublayer (PCS) performs word alignment, descrambling, and FEC. The Media Access Control (MAC) sublayer reconstructs the actual protocol data units (Ethernet frames, PCIe packets) and hands them to the system processor. The complete PMA-PCS-MAC sublayer hierarchy and the distinction between analog eye measurement and post-FEC error rate are covered in [SerDes Architecture and Eye Diagrams](#).

## CRC: Error Detection by Polynomial Division

Cyclic Redundancy Check (CRC) provides error detection at the frame level. The mathematical mechanism exploits a property of modulo-2 polynomial arithmetic.

The transmitter treats the data payload as a binary polynomial and divides it by a standardized generator polynomial. The division uses modulo-2 arithmetic, where addition and subtraction are both identical to the XOR operation. The transmitter appends the remainder of this division (the CRC bits) to the end of the payload before transmission.

Appending the exact remainder to the original data forces the entire newly formed package to become a perfect multiple of the generator polynomial. The mathematical principle is identical to base-10 arithmetic: dividing 17 by 5 yields a remainder of 2; subtracting that remainder produces 15, which is perfectly divisible by 5. In modulo-2 arithmetic, subtraction and addition are the same operation (XOR), so appending the remainder is mathematically equivalent to subtracting it.

The receiver divides the entire received package (data plus appended CRC bits) by the same generator polynomial. Dividing a perfect multiple by its generator mathematically guarantees a remainder of exactly zero. A non-zero remainder proves that at least one bit flipped during transmission.

CRC detects errors but cannot locate or correct them. When the CRC check fails, the protocol must request retransmission (as in TCP/IP error recovery) or rely on FEC to have already corrected the error at a lower layer.

## FEC: Error Correction by Reed-Solomon Coding

Forward error correction extends the polynomial mathematics of CRC from detection to actual correction. The most common algorithm in high-speed protocols (Ethernet IEEE 802.3, PCIe Gen 6) is Reed-Solomon coding.

The transmitter groups incoming data bits into discrete blocks and treats each block as a single algebraic polynomial. It divides this polynomial by a standardized generator polynomial and appends the parity bits (the mathematical remainder) to the block. The structure parallels CRC, but the generator polynomial and field arithmetic are chosen so that the remainder carries enough information to locate and repair errors, not merely detect them.

The receiver divides the entire received block (data plus parity) by the same generator polynomial. A remainder of zero confirms error-free transmission. A non-zero remainder is called the syndrome. The syndrome is not merely a flag; its specific numerical value acts as a coordinate map that identifies the exact positions of the corrupted bits within the block and their correct values. The receiver hardware inverts the corrupted bits on the fly, eliminating the latency penalty of halting the link to request retransmission.

Every Reed-Solomon code has a strict mathematical limit on how many symbol errors it can correct per block. A parameter  $t$  defines this limit: the code can correct up to  $t$  symbol errors per codeword. If a burst of errors exceeds  $t$ , the block is uncorrectable.

**Interleaving.** Physical noise bursts on a serial link tend to corrupt consecutive symbols. A localized noise event might flip 10 adjacent bits, exceeding the correction capability of a single FEC block. Interleaving defeats this failure mode. The transmitter distributes consecutive data symbols sequentially across several independent FEC blocks before transmission. A localized burst then affects only a few scattered symbols in each block rather than concentrating all errors in one. Each block easily repairs its few distributed errors within the mathematical correction limit.

## Worked Examples

---

### PCIe Encoding Overhead Across Generations

PCIe Gen 1 runs at 2.5 GT/s with 8b/10b encoding. The data throughput is:

$$2.5 \text{ GT/s} \times \frac{8}{10} = 2.0 \text{ Gb/s per lane}$$

PCIe Gen 3 runs at 8.0 GT/s with 128b/130b encoding. The data throughput is:

$$8.0 \text{ GT/s} \times \frac{128}{130} \approx 7.877 \text{ Gb/s per lane}$$

The encoding overhead dropped from 20% to 1.5%, and the transfer rate increased by a factor of 3.2. The combined effect is a nearly four-fold increase in per-lane data throughput.

### UI Calculation for DDR4-3200

DDR4-3200 operates at 3.2 GT/s (3.2 billion transfers per second). The Unit Interval is:

$$\text{UI} = \frac{1}{3.2 \times 10^9} = 312.5 \text{ ps}$$

The physical clock frequency is 1.6 GHz (half the transfer rate, because each clock cycle triggers two transfers). An oscilloscope measuring jitter or drawing an eye diagram on the DQ data wire uses 312.5 ps as the UI for time-base alignment.

## CRC Verification

A transmitter sends a data payload treated as polynomial  $D(x)$  and divides by generator  $G(x)$ , obtaining remainder  $R(x)$ . The transmitted package is  $D(x) \cdot x^r + R(x)$ , where  $r$  is the degree of the generator (the number of CRC bits). The receiver divides the entire received package by  $G(x)$ :

$$\frac{D(x) \cdot x^r + R(x)}{G(x)}$$

The first term  $D(x) \cdot x^r / G(x)$  produces quotient  $Q(x)$  with remainder  $R(x)$ . Adding  $R(x)$  to  $R(x)$  in modulo-2 arithmetic yields zero (any value XORed with itself is zero). The receiver obtains a remainder of exactly zero, confirming no bit errors.

## Edge Cases

---

### Running Disparity and Special Symbols in 8b/10b


The 8b/10b code table contains two variants (positive and negative disparity) for most data symbols. The encoder selects the variant that steers the cumulative DC balance toward zero. Some symbols have only one variant and are inherently disparity-neutral. The code also reserves specific patterns as control symbols (K-codes) that can never appear in valid encoded data. These K-codes serve as unambiguous markers for comma detection, ordered set signaling, and link training. Their uniqueness within the code space is what makes reliable word alignment possible even in a continuous, unsynchronized bit stream.

### Scrambling vs. Deterministic Encoding

8b/10b encoding provides a deterministic, per-symbol guarantee of DC balance and bounded run length. Every transmitted symbol is individually constrained. 128b/130b and 64b/66b encoding rely on scrambling, which provides statistical balance: over any sufficiently long window, the density of ones and zeros converges to 50%, but short sequences within the stream can be temporarily unbalanced. The LFSR scrambler is designed so that pathological input patterns (such as all zeros or all ones) still produce output with adequate transition density. The 2-bit sync header at every block boundary provides a guaranteed transition regardless of the scrambled payload content, ensuring the CDR always receives periodic edges.

## FEC Correction Threshold and Residual Errors

An FEC code with correction capability  $t$  can repair up to  $t$  symbol errors per codeword. If the raw error rate is high enough that bursts occasionally exceed  $t$  (even after interleaving), those blocks are uncorrectable and produce post-FEC errors. Protocol specifications define target post-FEC bit error rates (typically  $10^{-15}$  or better) and select FEC parameters (  $t$  , block size, interleaving depth) to achieve that target given the expected pre-FEC raw error rate from the analog channel. The pre-FEC BER floor and the FEC correction threshold together determine the minimum analog eye opening that the equalization chain must achieve.



PART 6

# Coherent Optics

## Chapter 19: The Physics of the Electromagnetic Wave

---

Engineers must visualize light not as a stream of solid particles, but as an oscillating physical field moving through space. Bridging the mechanical world of electrical voltages with the quantum world of electromagnetic waves requires a precise geometric understanding of the physics.

### Visualizing the Electromagnetic Wave

---

Imagine a straight line pointing directly away from you. This central axis represents the exact direction the light is traveling down the fiber optic cable.

An electromagnetic wave consists of two linked components: an electric field and a magnetic field. These fields do not travel alongside the light. They are the light. The electric field oscillates entirely perpendicular to the forward direction of travel. The magnetic field oscillates simultaneously, but it remains perfectly perpendicular to both the direction of travel and the electric field.

Amplitude defines the physical height of this oscillation. It dictates the maximum strength of the electric field at the peak of the wave. Phase represents the timing of the wave. Freezing time reveals the phase as the exact position of the wave within its continuous cycle, whether at an absolute peak, a lowest trough, or an intermediate state.

### Understanding Polarization

---

Polarization is not an additional wave inserted into the system. Polarization simply defines the physical orientation of the electric field's oscillation relative to the environment.

An electric field oscillating strictly up and down along the vertical axis constitutes vertical polarization. Rotating that entire wave ninety degrees forces the electric field to oscillate left and right along the horizontal axis, constituting horizontal polarization.

These two spatial orientations remain perfectly perpendicular to each other. You can transmit a vertically polarized wave and a horizontally polarized wave down the exact same fiber optic core simultaneously. They travel in the identical forward direction and occupy the exact same physical space. Their electric fields never interact. This physical orthogonality allows optical systems to double their data capacity immediately by modulating independent data streams onto each distinct polarization state.

### Splitting the Laser into Polarizations

---

The physical separation of the light occurs immediately after the laser source, before any data enters the optical domain.

The continuous, unmodulated laser beam enters a specialized optical prism called a polarization beam splitter. This component physically divides the raw light into two distinct pathways within the transmitter hardware. The hardware designates one pathway for the vertical data channel and the second pathway for the horizontal data channel.

Each separate optical pathway routes into its own dedicated array of electro-optic modulators. The electrical data streams sculpt the phase and amplitude of the light in the vertical pathway completely independently from the horizontal pathway.

Completing the modulation process means both pathways now carry distinct encoded data. The hardware must physically rotate the orientation of the light in the second pathway to create the necessary orthogonal relationship. The light passes through a polarization rotator, physically twisting its electric field exactly ninety degrees relative to the first pathway.

The transmitter routes the vertical pathway and the newly rotated horizontal pathway into a polarization beam combiner. This final component merges the two separate beams back into a single column of light. This unified beam now contains both orthogonal polarizations carrying independent data, and it exits the transceiver directly into the fiber optic cable.

## The Alignment of Electric and Magnetic Fields

---

Transmitting two independent polarizations mandates the existence of two distinct electric fields and two distinct magnetic fields occupying the exact same physical space.

The vertically polarized wave consists of an electric field oscillating on the vertical axis and its paired magnetic field oscillating on the horizontal axis. The horizontally polarized wave consists of an electric field oscillating on the horizontal axis and its paired magnetic field oscillating on the vertical axis.

This geometric alignment means the vertical electric field occupies the exact same physical plane as the horizontal wave's magnetic field. The horizontal electric field occupies the exact same physical plane as the vertical wave's magnetic field.

Physics dictates that these overlapping fields will never interfere with each other. Electric fields only interact with other electric fields. Magnetic fields only interact with other magnetic fields. The vertical electric field remains completely blind to the magnetic field sharing its vertical axis. The entire dual-polarization system functions flawlessly because the two electric fields remain strictly perpendicular to each other, and the two magnetic fields remain strictly perpendicular to each other.

## Chapter 20: Electro-Optic Modulation and the Refractive Index

---

### Translating Electrons to Photons

---

Inside the transceiver, we must translate electrical signals moving through copper traces into optical signals moving through glass. We achieve this transduction using an electro-optic modulator.

The system begins with a highly stable laser emitting a continuous, uninterrupted beam of pure light. This continuous stream of photons acts as the carrier wave. The transceiver never rapidly turns the laser on and off to create data pulses. It feeds this constant optical stream directly into the modulator.

The modulator splits the incoming laser beam down two parallel microscopic pathways constructed from a specialized crystal material, such as lithium niobate.

### The Physics of the Refractive Index

---

Controlling the optical domain relies entirely on the molecular structure of the modulator crystal. Materials such as lithium niobate exhibit the electro-optic effect. The crystal lattice holds its atoms and electron clouds in a very specific, rigid geometry.

Applying an electrical voltage across this crystal generates a localized electric field. This electric field physically pulls and distorts the electron clouds within the crystal lattice. This microscopic structural tension alters the optical density of the material. The changing optical density forces the photons to interact more heavily with the distorted lattice, effectively reducing the speed of light within that specific section of the crystal.

### Velocity, Frequency, and Wavelength

---

The universal wave equation dictates that velocity equals frequency multiplied by wavelength. The energy of a photon is directly tied to its fundamental frequency. Conservation of energy requires the frequency to remain absolutely constant when light enters a new physical medium.

The locked frequency dictates that the drop in velocity inside the crystal forces a proportional drop in wavelength. The physical distance between the peaks of the wave compresses. The wave travels slower, the peaks pack closer together, and the number of peaks passing a specific point every second remains exactly the same.

### The Physics of Slowing Light

---

A laser does not emit white light. White light contains every color and wavelength, which a prism splits into a broad spectrum. A laser emits monochromatic light. It produces one single, incredibly precise wavelength operating at a single optical frequency. There is only one frequency present, eliminating the prism effect entirely.

This single-frequency light entering the crystal maintains an absolutely constant frequency. The number of cycles per second never changes. The physical wavelength simply compresses within the material.

Photons traveling through a vacuum encounter no resistance. Photons traveling through a crystal interact with the electromagnetic fields of the local atoms. Applying a voltage distorts those atomic fields, increasing the resistance. The wave propagates through this denser electrical environment at a slower velocity.

The wave immediately accelerates back to its original speed upon exiting the modulator. The wavelength stretches back to its original size. The permanent change lies entirely in the arrival time. Traveling at a slower velocity through the crystal means the wave took longer to traverse that specific physical distance. The wave exits the modulator slightly later than it would have without the applied voltage. The frequency remains identical, but the entire continuous wave is now permanently delayed by a fraction of a cycle. This permanent temporal delay constitutes the phase shift.

## Speed and Phase Translation

---

This localized change in the speed of the photons acts as the exact mechanism that creates a phase shift.

A wave traveling through the crystal tunnel at its normal speed exits at a precise moment in time. Slowing the wave down forces it to exit the tunnel slightly later. Exiting later means the peaks and troughs of the wave are physically delayed relative to a wave that traveled undisturbed. This temporal delay creates a geometric offset. The geometric offset defines the literal meaning of a phase shift.

The electrical data signal transitioning between different voltage levels billions of times per second causes the localized electric field to fluctuate perfectly in sync. The refractive index constantly shifts, the speed of the photons constantly changes, and the resulting optical phase shifts continuously to mirror the high-speed electrical data.

## Varying Phase and Amplitude

---

Applying a precise electrical voltage driven by the incoming data bits to one of the two crystal pathways slows down the light in that specific tunnel by a microscopic fraction of a second. This controlled delay shifts the phase of the wave in that pathway relative to the undisturbed wave traveling in the second pathway.

The two pathways merge back together at the exit of the modulator. The physics of optical interference immediately govern the output.

Delaying the first wave until its absolute peak perfectly aligns with the absolute trough of the second wave causes the two waves to cancel each other out completely. The resulting amplitude drops to zero. Aligning the peaks of both waves perfectly causes them to combine into a wave with maximum amplitude. Modulating the electrical voltage across multiple parallel pathways allows the transceiver to continuously sculpt the precise phase and amplitude of the outgoing photons before the light ever enters the transmission fiber.

## Chapter 21: The I/Q Modulator Architecture

---

### The Nested Modulator Array

---

Achieving complex data transmission requires controlling both phase and amplitude simultaneously to plot points on a two-dimensional constellation diagram like Quadrature Amplitude Modulation. A single basic modulator can only change the amplitude of a wave by causing interference between its own internal pathways.

Engineers construct a nested array called an I/Q modulator to gain full control over the two-dimensional space.

The hardware takes the light for a single polarization and splits it one more time into an In-phase pathway and a Quadrature pathway. The hardware applies a permanent, physical ninety-degree phase shift to the light entering the Quadrature pathway.

Each of these two pathways receives its own dedicated modulator. The first modulator controls the amplitude of the In-phase wave. The second modulator controls the amplitude of the Quadrature wave.

The electrical data stream feeds into both modulators independently. The hardware dynamically changes the amplitude of the In-phase wave and the amplitude of the ninety-degree shifted Quadrature wave based on the incoming voltage. The transmitter recombines these two distinct outputs to synthesize a brand new wave. The mathematical addition of these two orthogonal waves allows the system to instantly generate an optical signal with any specific phase and any specific amplitude required by the constellation grid. The nested array grants the transceiver deterministic control over the exact geometry of the exiting optical wave.

### The Modulator Architecture

---

The electrical data routing dictates the ultimate purpose of this geometric structure.

The system splits the initial laser into two polarization paths. It splits each of those paths into an In-Phase pathway and a Quadrature pathway. This results in four primary pathways, each containing its own Mach-Zehnder modulator. Inside each modulator, the light splits one final time into two microscopic arms to cause the amplitude interference, creating the exactly eight total internal beams.

The system does not feed the same electrical signal into all four modulators. The core objective of this complex architecture is maximizing data throughput. The transceiver takes the massive incoming data pipeline and divides it into four completely independent high-speed data streams. It feeds a unique electrical data stream into the Vertical In-Phase modulator, the Vertical Quadrature modulator, the Horizontal In-Phase modulator, and the Horizontal Quadrature modulator. This parallel architecture allows the system to transmit four distinct sets of data simultaneously using a single laser source.

## The Quadrature Phase Shift

---

The shift applied to the Quadrature wave is strictly a temporal delay along the exact same physical polarization plane. The hardware delays the Quadrature wave by exactly one quarter of a wavelength relative to the In-Phase wave.

Delaying a standard sine wave by ninety degrees transforms it into a cosine wave. These two waves exist on the exact same physical polarization axis, but their mathematical relationship becomes orthogonal. This mathematical orthogonality allows the digital signal processor at the receiver to isolate the amplitude of the sine wave and the amplitude of the cosine wave independently, even after the transmitter added them together into a single composite wave.

## The Illusion of Two Waves

---

The sum of the In-Phase wave and the Quadrature wave absolutely creates a single, unified wave.

They do not remain separate entities inside the fiber. They physically interfere with each other and merge completely the moment they exit their respective modulators.

Adding a sine wave and a cosine wave of the same frequency together always results in a single, brand new wave. This resulting wave possesses a specific final amplitude and a specific final phase.

This mathematical reality drives the entire architecture of the transmitter. Controlling the amplitude of the In-Phase wave and independently controlling the amplitude of the Quadrature wave allows the hardware to precisely dictate the amplitude and phase of the final combined wave. The two orthogonal waves act simply as the internal mathematical building blocks used to sculpt the final, singular photon stream entering the fiber optic cable.

## Physical versus Mathematical Orthogonality

---

Physical orthogonality prevents interference entirely. The vertical polarization wave and the horizontal polarization wave exist on perpendicular physical planes. They travel through the same glass core simultaneously without ever adding together or canceling each other out.

Mathematical orthogonality applies to waves traveling on the exact same physical plane. The In-Phase sine wave and the Quadrature cosine wave exist on the identical physical axis. They absolutely interfere with each other to form the single composite wave.

Their mathematical orthogonality means their specific ninety-degree temporal relationship allows the receiver to separate the original data. The digital signal processor at the receiver multiplies the single incoming composite wave by a locally generated pure sine wave to mathematically extract only the amplitude of the original In-Phase component. The processor multiplies that exact same composite wave by a locally generated pure cosine wave to mathematically extract only the amplitude of the original Quadrature component.

The waves physically combine into a single entity for the physical journey across the ocean. The mathematical properties of sine and cosine ensure the encoded data remains separable and distinct upon arrival.

## The Four Degrees of Freedom

---

The system possesses exactly four independent degrees of freedom: the amplitude of the vertical In-Phase wave, the amplitude of the vertical Quadrature wave, the amplitude of the horizontal In-Phase wave, and the amplitude of the horizontal Quadrature wave.

Engineers map raw digital data directly onto these four independent amplitude levels. Operating a system using Quadrature Amplitude Modulation allows the transmitter to define multiple distinct amplitude levels for each degree of freedom. Assigning four distinct amplitude levels to the vertical In-Phase channel allows that specific channel to represent two digital bits simultaneously. Replicating this strategy across all four degrees of freedom allows the transmitter to send eight total bits of data during every single optical clock cycle. This geometric multiplexing creates the massive terabit throughput associated with coherent optical networks.

## Translating Photons to Electrons

---

The digital processing architecture requires electrical inputs. The receiver must convert the optical signals back into electrical signals before any analog-to-digital conversion can occur.

The incoming fiber optic cable routes the light into a coherent receiver. The hardware reverses the entire transmission process. The receiver uses a polarization beam splitter to separate the vertical and horizontal planes. The hardware mixes each plane with a local oscillator laser to optically extract the In-Phase and Quadrature components.

The receiver directs these four resulting optical streams into four independent photodetectors. The photodetectors absorb the incoming photons and generate a proportional flow of electrons. A bright optical peak generates a strong electrical current. A dim optical trough generates a weak electrical current. The hardware successfully translates the varying optical phase and amplitude into four distinct, rapidly fluctuating electrical voltages.

## Digitization and Recombination

---

The analog nature of the signal ends immediately after the photodetectors.

The receiver feeds the four electrical voltage streams directly into four extremely high-speed Analog-to-Digital Converters. The converters sample the continuous analog voltage levels billions of times per second. They translate the physical height of the electrical wave into a precise binary number.

The four converters stream these binary numbers into a massive Digital Signal Processor. The processor mathematically analyzes the incoming numbers to reverse any physical distortions caused by the hundreds of kilometers of fiber optic glass.

The processor inspects the cleaned numerical values for all four degrees of freedom simultaneously. It compares the measured amplitude levels against the known geometric constellation map to determine exactly which specific zeros and ones the transmitter originally encoded. The processor extracts the independent data bits from the vertical and horizontal channels and reconstructs them in the correct sequential order. The hardware pushes this newly unified serial data stream out to the host network router, completing the transmission cycle.

## Chapter 22: Wideband Signal Analysis and EVM

---

### System Challenges and Optical Concerns

---

Sending complex phase-encoded light over hundreds of kilometers of glass introduces severe physical distortions that the wideband analyzer must identify and quantify.

**Chromatic Dispersion:** Silica glass possesses a refractive index that varies slightly depending on the exact wavelength of the light. An optical pulse actually contains a narrow band of frequencies. The higher frequency components travel through the glass at a slightly different speed than the lower frequency components. Over long distances, the optical pulse spreads out horizontally until it overlaps with adjacent pulses, destroying the integrity of the data sequence.

**Polarization Mode Dispersion:** Optical fiber is never perfectly circular or completely stress-free. Temperature changes, physical vibrations, and manufacturing imperfections cause the fiber core to become slightly elliptical. This physical distortion causes the two orthogonal polarization states to travel at different speeds and rotate unpredictably. The wideband analyzer tracks how violently these polarization states drift relative to each other.

**Laser Phase Noise:** Coherent detection relies on the absolute purity of the transmitter and receiver lasers. A real-world laser is never perfectly stable. It drifts continuously in frequency and phase. If the laser linewidth is too broad, the random phase jitter obscures the deliberate phase shifts encoded by the transmitter. The wideband analyzer quantifies this phase noise to ensure the internal digital signal processor can track and correct the laser drift before the link fails.

### The Analyzer as an Ideal Receiver

---

Engineers need a way to measure the exact health of the optical signal traveling through the fiber before a commercial receiver attempts to process it.

A Wideband Signal Analyzer designed for coherent optics acts as a perfect, idealized optical receiver. It physically connects to the fiber optic cable and ingests the raw photons directly. It does not look at the final electrical bits exiting a commercial receiver. It captures the raw optical wave to diagnose the physical layer physics.

The instrument contains its own pristine local oscillator lasers, its own flawless photodetectors, and its own extreme-precision Analog-to-Digital Converters. The analyzer ingests the four optical waveforms, converts them to electrical voltages, and digitizes them using hardware far superior to standard commercial transceivers.

### Measuring Error Vector Magnitude

---

Measuring the Error Vector Magnitude requires comparing the actual received signal against mathematical perfection.

The analyzer captures the incoming optical signal and digitizes the four degrees of freedom. The internal software reconstructs the two-dimensional constellation diagram. The software analyzes the exact geometric placement of every single received optical symbol.

The analyzer calculates the geometric distance between where a received symbol actually landed on the coordinate system and where it was theoretically supposed to land. The analyzer deduces the exact binary data the transmitter intended to send. The software calculates where that specific symbol should have landed on a perfect, mathematical grid. The analyzer draws a geometric vector between the perfect mathematical location and the actual physical location the optical symbol landed. The length of this vector represents the Error Vector Magnitude. A large magnitude indicates severe signal distortion.

## Isolating Physical Impairments

---

The primary value of the analyzer lies in identifying exactly why the Error Vector Magnitude is degrading.

The analyzer intercepts the optical signal before commercial digital signal processors attempt to erase the physical damage. The analyzer's software quantifies the exact amount of chromatic dispersion stretching the optical pulses. It measures the precise severity of the laser phase noise blurring the constellation points. It calculates the exact polarization drift twisting the vertical and horizontal planes. The analysis isolates specific physical impairments by identifying exactly how the symbols are scattered. It also measures the exact shape of the optical spectrum to ensure the signal does not bleed into adjacent wavelength channels within the fiber.

Engineers rely on this raw physical layer data. Designing advanced commercial receivers requires knowing exactly how much physical distortion the glass introduces. The wideband analyzer provides the absolute ground truth regarding the health of the photons.