




The Transformer Architecture

A Geometric Toy Example
from Scratch

By Lydia Pedersen

The Transformer Architecture: A Geometric Toy Example from Scratch

Copyright © 2026 Lydia Pedersen. All rights reserved.

No part of this publication may be reproduced, distributed, or transmitted in any form or by any means, including photocopying, recording, or other electronic or mechanical methods, without the prior written permission of the publisher, except in the case of brief quotations embodied in critical reviews and certain other noncommercial uses permitted by copyright law.

Table of Contents

Preface: The Big Picture & Tensor Notation

Chapter 1: Tokens, One-Hot Encodings, and the Embedding Matrix

Chapter 2: The Permutation Invariance Problem & Positional Encoding

Chapter 3: The Motivation for Q, K, and V (Asymmetric Similarity)

Chapter 4: The Attention Score and $\sqrt{d_k}$

Chapter 5: Causal Masking

Chapter 6: From Scores to Synthesis: Softmax and The Value Matrix

Chapter 7: The Cross-Head Mixer and The Projection Matrix

Chapter 8: The Residual Stream and the Central Memory Bus

Chapter 9: Taming the Stream: The Geometry of Layer Normalization

Chapter 10: The MLP as a Key-Value Memory Bank (Expansion)

Chapter 11: The MLP - Activation and Contraction

Chapter 12: Completing Layer 1 Residuals and Normalization

Chapter 13: Layer 2 Self-Attention

Chapter 14: Scoring Deep Context: Layer 2 Attention

Chapter 15: The Final Blend: Masking, Softmax, and Values in Layer 2

Chapter 16: Deepening the Representation: MLP and Residuals in Layer 2

Chapter 17: Mapping Back to Words

Chapter 18: Final Softmax and Predictions

Chapter 19: Cross-Entropy Loss

Chapter 20: The Beautiful Cancellation

Chapter 21: Backpropagating Through the Unembedding and Residual Stream

Chapter 22: Backpropagating Through Attention: The Softmax and the Mask

Chapter 23: Backpropagation Through Attention (Chapter 2: Routing to Q, K, and V)

Chapter 24: Updating the Embeddings and Conclusion