




The Transformer Architecture

A Geometric Toy Example
from Scratch

By Lydia Pedersen

The Transformer Architecture: A Geometric Toy Example from Scratch

Copyright © 2026 Lydia Pedersen. All rights reserved.

No part of this publication may be reproduced, distributed, or transmitted in any form or by any means, including photocopying, recording, or other electronic or mechanical methods, without the prior written permission of the publisher, except in the case of brief quotations embodied in critical reviews and certain other noncommercial uses permitted by copyright law.

Table of Contents

Preface: The Big Picture & Tensor Notation

Chapter 1: Tokens, One-Hot Encodings, and the Embedding Matrix

Chapter 2: The Permutation Invariance Problem & Positional Encoding

Chapter 3: The Motivation for Q, K, and V (Asymmetric Similarity)

Chapter 4: The Attention Score and $\sqrt{d_k}$

Chapter 5: Causal Masking

Chapter 6: From Scores to Synthesis: Softmax and The Value Matrix

Chapter 7: The Projection Matrix and The Cross-Head Mixer

Chapter 8: The Residual Stream and the Central Memory Bus

Chapter 9: Taming the Stream: The Geometry of Layer Normalization

Chapter 10: The Residual Stream and MLP Expansion

Chapter 11: Layer Normalization and The Multi-Layer Perceptron

Chapter 12: The MLP as a Key-Value Memory and Layer 1 Completion

Chapter 13: The MLP Activation and Layer 2 Self-Attention

Chapter 14: Final Layer 1 Residuals, Norms, and Layer 2 Attention

Chapter 15: Layer 2 Self-Attention

Chapter 16: Deepening the Representation: MLP and Residuals in Layer 2

Chapter 17: Mapping Back to Words

Chapter 18: Final Softmax and Predictions

Chapter 19: The Cross-Entropy Loss Function

Chapter 20: The Beautiful Cancellation

Chapter 21: Backpropagating Through the Unembedding and Residual Stream

Chapter 22: Backpropagation - Attention (Chapter 1: Softmax & Scores)

Chapter 23: Backpropagation - Attention (Chapter 2: Q, K, V)

Chapter 24: Updating the Embeddings and Conclusion